

Building Dynamic Tools with DynamoRIO on x86 and ARMv8

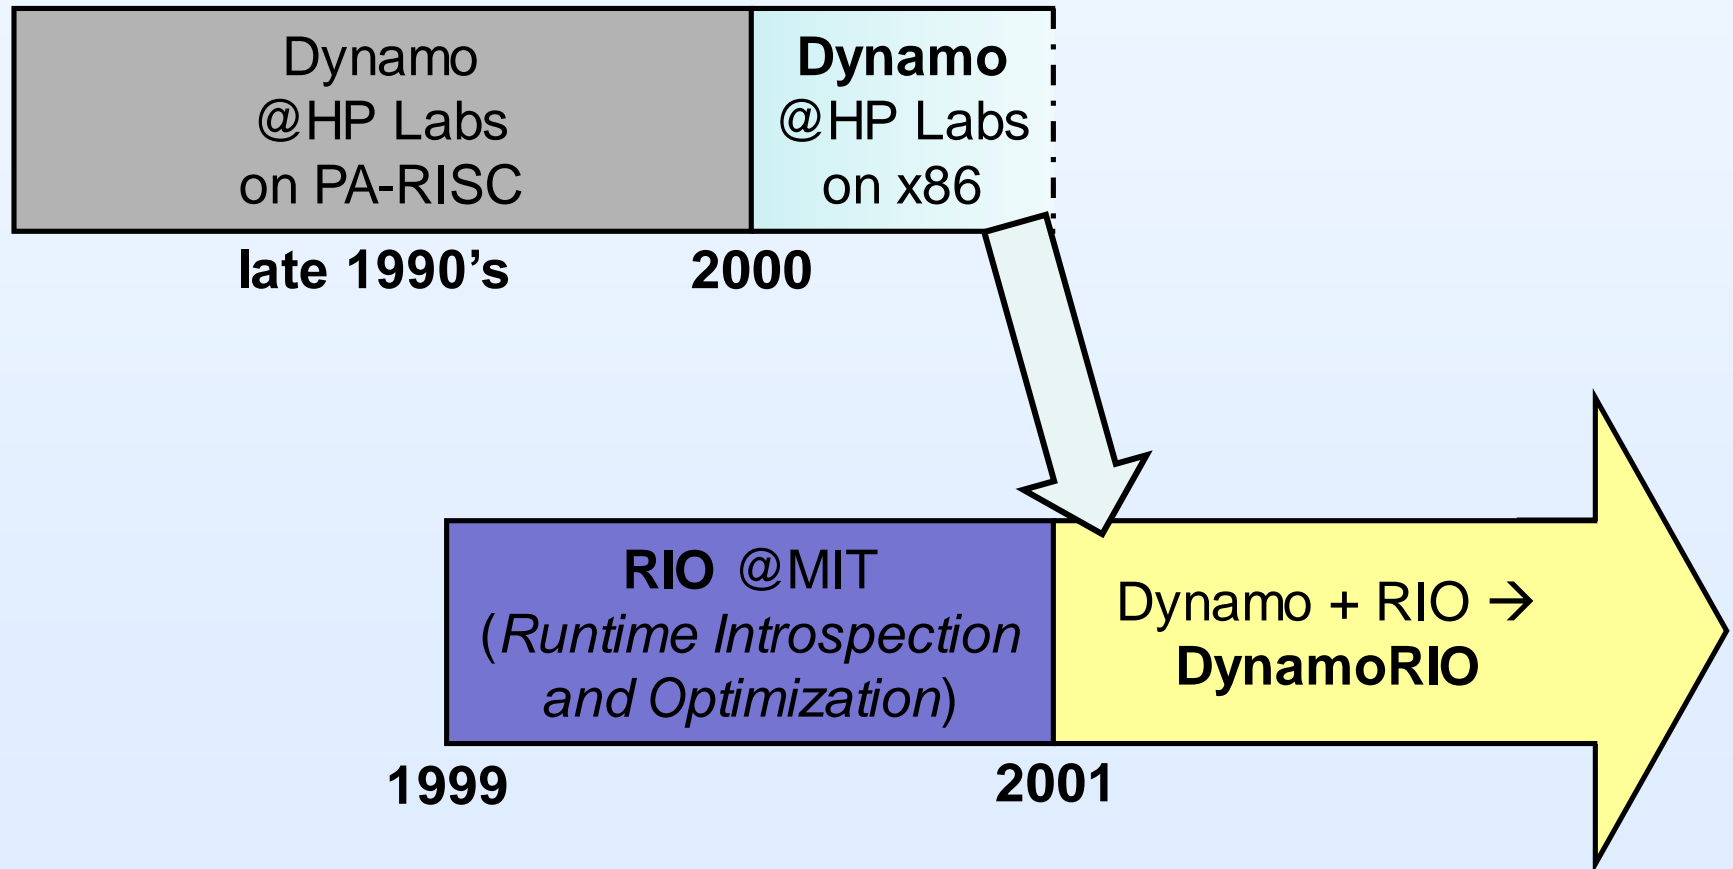
Derek Bruening
Google

Al Grant
Chris Adeniyi-Jones
Edmund Grimley Evans
Kevin Zhou
ARM

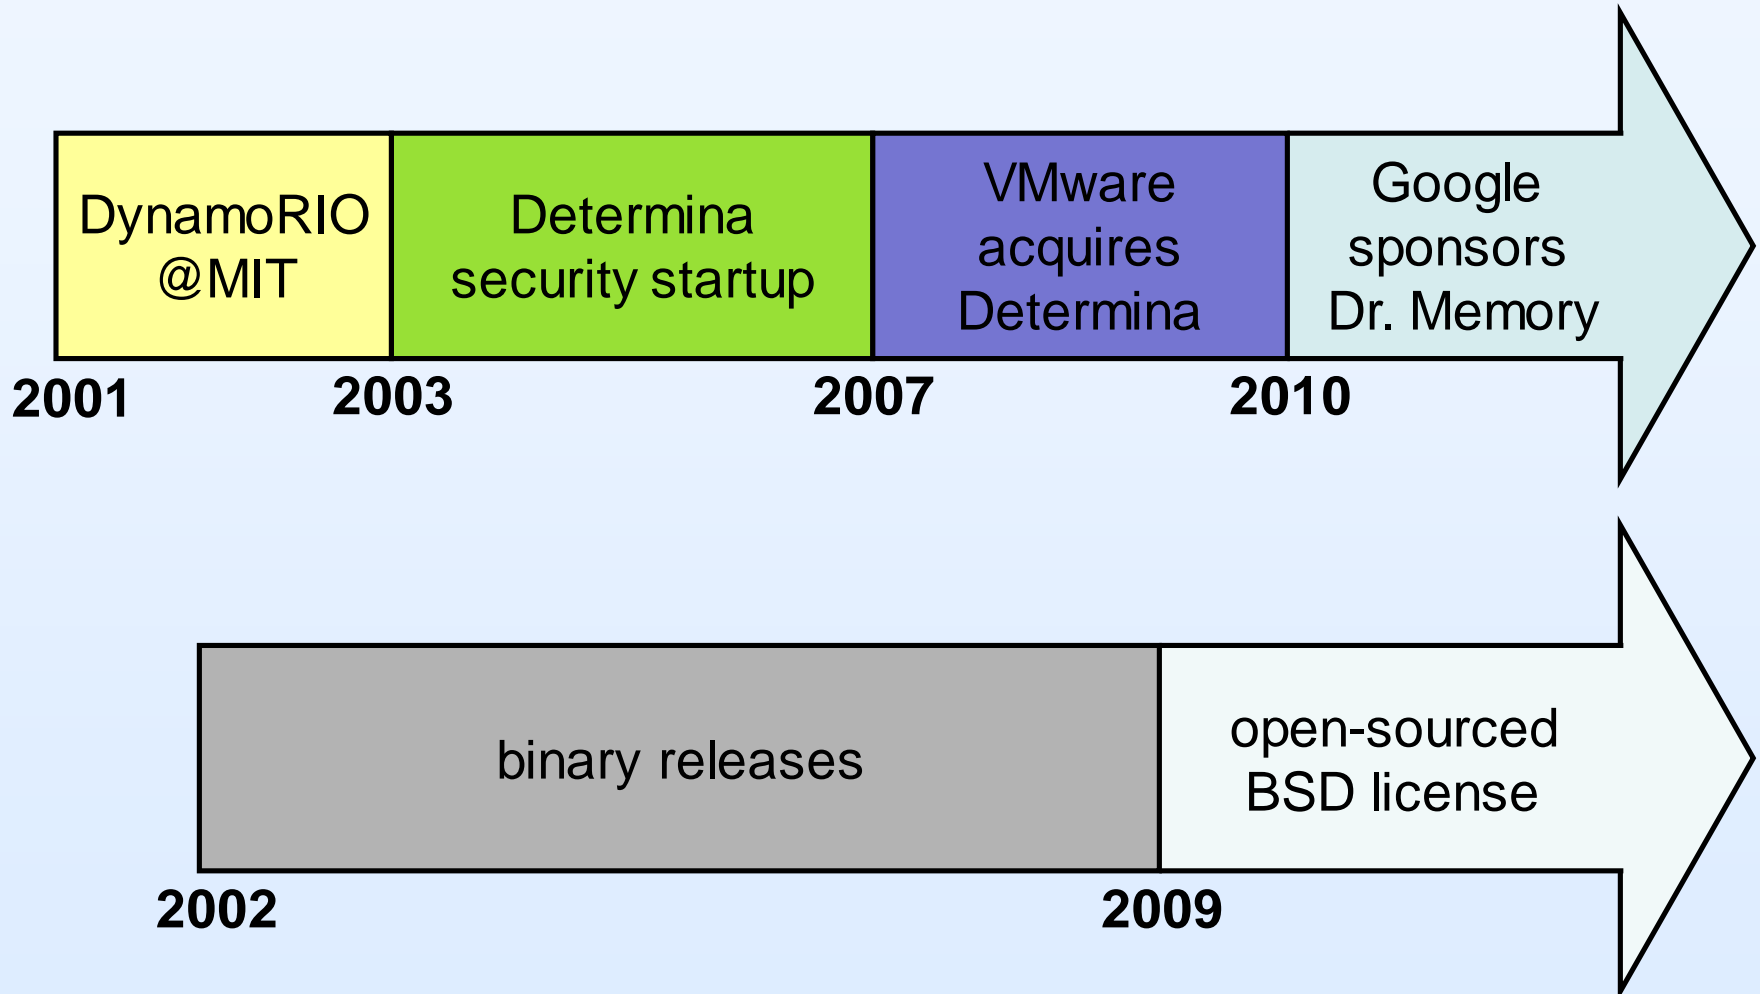
Tutorial Outline

- 8:30- 8:40 Welcome + DynamoRIO History
- 8:40- 9:10 Tool Demonstrations
- 9:10- 9:30 DynamoRIO System Overview
- 9:30- 9:45 DynamoRIO API Part 1
- 9:45-10:00 ARM Status + Demonstrations
- *10:00-10:30 Break*
- 10:30-10:55 DynamoRIO API Part 2
- 10:55-11:20 How to Create Your Own Tool
- 11:20-11:40 Tool Internals
- 11:40-11:45 Q & A

DynamoRIO



DynamoRIO History



Tool Demonstrations

8:30- 8:40	Welcome + DynamoRIO History
8:40- 9:10	Tool Demonstrations
9:10- 9:30	DynamoRIO System Overview
9:30- 9:45	DynamoRIO API Part 1
9:45-10:00	ARM Status + Demonstrations
10:00-10:30	<i>Break</i>
10:30-10:55	DynamoRIO API Part 2
10:55-11:20	How to Create Your Own Tool
11:20-11:40	Tool Internals
11:40-11:45	Q & A

Memory Trace Tool

- Instruments every load and store
- Writes the trace of memory accesses to a file
- Usage:
 - `drrun -c samples/bin64/libmemtrace_x86_text.so`
 `-- <app>`

Memory Trace Tool Results

```
% bin64/drrun -c samples/bin64/libmemtrace_x86_text.so --  
~/multi_threaded_pi_estimator
```

Format: <instr address>,<(r)ead/(w)rite>,<data size>,<data address>

0x00007fb31a5462d3,w,8,0x00007ffc4febfb18

0x00007fb31a549a74,w,8,0x00007ffc4febfb08

0x00007fb31a549a8f,r,8,0x00007fb31a767e70

0x00007fb3199a54fc,r,4,0x00007ffc4febfa38

0x00007fb3199a5526,r,1,0x0000000000400e98

0x00007fb31a549a96,w,8,0x00007fb31a767c98

0x00007fb31a549aa7,r,8,0x00007fb31a768000

0x00007fb3199af4a0,r,8,0x00007ffc4febfa98

0x0000000000400db3,r,8,0x00007ffc4febfa40

0x0000000000400db4,r,8,0x00007ffc4febfa48

Dr. CacheSim

- Multi-process, multi-thread online cache simulator
- Each thread collects and sends its own memory access trace to a single simulator
- Simple static scheduling
- Extensible
- Usage:
 - `drrun -t drcachesim -- <app>`

Dr. CacheSim Results

```
% bin64/drrun -t drcachesim -- ./multi_threaded_pi_estimator
Estimation of pi is 3.142425985001098
---- <application exited with code 0> ----
Core #0 (1 thread(s))
  L1I stats:
    Hits:                243,293
    Misses:               1,077
    Miss rate:           0.44%
  L1D stats:
    Hits:                67,738
    Misses:               3,471
    Miss rate:           4.87%
Core #1 (1 thread(s))
  L1I stats:
    Hits:                8,589
    Misses:               95
    Miss rate:           1.09%
  L1D stats:
    Hits:                3,302
    Misses:               211
    Miss rate:           6.01%
Core #2 (1 thread(s))
  L1I stats:
    Hits:                1,906
    Misses:               94
    Miss rate:           4.70%
  L1D stats:
    Hits:                749
    Misses:               121
    Miss rate:           13.91%
Core #3 (0 thread(s))
  LL stats:
    Hits:                1,186
    Misses:               3,883
    Local miss rate:     76.60%
    Child hits:          325,577
    Total miss rate:     1.17%
```

Dr. Memory: Memory Debugger

- Detects reads of uninitialized memory
- Detects heap errors
 - Out-of-bounds accesses (underflow, overflow)
 - Access to freed memory
 - Invalid frees
 - Memory leaks
- Detects other accesses to invalid memory
 - Stack tracking
- Detects Windows handle leaks
- Operates at runtime on unmodified Windows, Linux, MacOS, and Android binaries

Dr. Memory Results

The screenshot shows the Microsoft Visual Studio IDE with the 'Buggy' project open. The code editor displays the following C++ code:

```
};

int main()
{
    int *p = new int;
    if (*p != 10) /* bug: uninitialized */
        std::cout << "hi" << std::endl;
    delete p;

    hasdtr *x = new hasdtr[7];
    x[6].y = x[7].y; /* bug: overflow */
    // ... (commented out) ...
}
```

The 'Output' window at the bottom shows the results of the Dr. Memory analysis, configured to show output from 'Dr. Memory'. The output text is as follows:

```
Dr. Memory version 1.9.16867
Running ""c:\users\bruening\documents\visual studio 2015\Projects\Buggy\Debug\Buggy.exe""

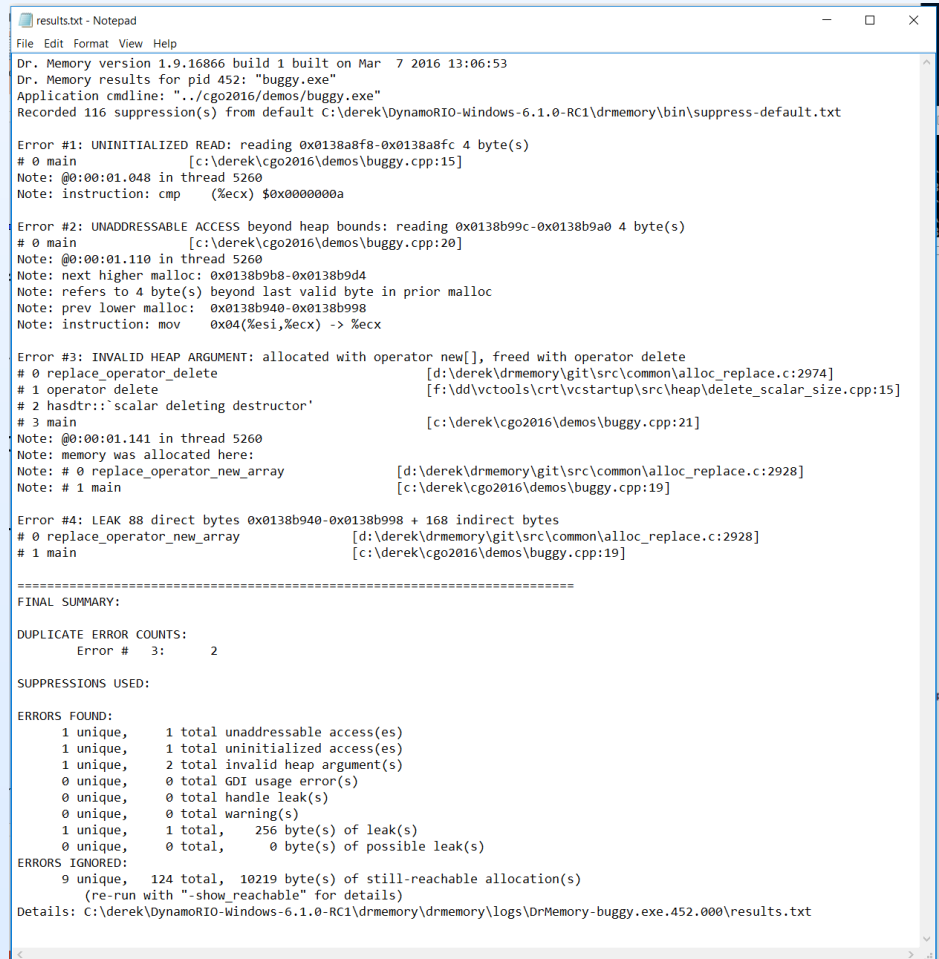
Error #1: UNINITIALIZED READ: reading 4 byte(s)
main
c:\users\bruening\documents\visual studio 2015\projects\buggy\buggy\buggy.cpp(15):
hi

Error #2: UNADDRESSABLE ACCESS beyond heap bounds: reading 4 byte(s)
main
c:\users\bruening\documents\visual studio 2015\projects\buggy\buggy\buggy.cpp(20):
Note: refers to 4 byte(s) beyond last valid byte in prior malloc
```

The status bar at the bottom indicates the current line is 15 in the file `c:\users\bruening\documents\visual studio 2015\projects\buggy\buggy\buggy.cpp(15):`.

Dr. Memory Command Line Results

- To run:
`drmemory -t drmemory -- <app>`
- On Windows, notepad shows results



```
results.txt - Notepad
File Edit Format View Help

Dr. Memory version 1.9.16866 build 1 built on Mar  7 2016 13:06:53
Dr. Memory results for pid 452: "buggy.exe"
Application cmdline: "..\cgo2016\demos\buggy.exe"
Recorded 116 suppression(s) from default C:\derek\DynamoRIO-Windows-6.1.0-RC1\drmemory\bin\suppress-default.txt

Error #1: UNINITIALIZED READ: reading 0x0138a8f8-0x0138a8fc 4 byte(s)
# 0 main [c:\derek\cgo2016\demos\buggy.cpp:15]
Note: @0:00:01.048 in thread 5260
Note: instruction: cmp (%ecx) $0x0000000a

Error #2: UNADDRESSABLE ACCESS beyond heap bounds: reading 0x0138b99c-0x0138b9a0 4 byte(s)
# 0 main [c:\derek\cgo2016\demos\buggy.cpp:20]
Note: @0:00:01.110 in thread 5260
Note: next higher malloc: 0x0138b9b8-0x0138b9d4
Note: refers to 4 byte(s) beyond last valid byte in prior malloc
Note: prev lower malloc: 0x0138b940-0x0138b998
Note: instruction: mov 0x04(%esi,%ecx) -> %ecx

Error #3: INVALID HEAP ARGUMENT: allocated with operator new[], freed with operator delete
# 0 replace_operator_delete [d:\derek\drmemory\git\src\common\alloc_replace.c:2974]
# 1 operator_delete [f:\dd\vctools\crt\vcstartup\src\heap\delete_scalar_size.cpp:15]
# 2 hasdtr::'scalar deleting destructor'
# 3 main [c:\derek\cgo2016\demos\buggy.cpp:21]
Note: @0:00:01.141 in thread 5260
Note: memory was allocated here:
Note: # 0 replace_operator_new_array [d:\derek\drmemory\git\src\common\alloc_replace.c:2928]
Note: # 1 main [c:\derek\cgo2016\demos\buggy.cpp:19]

Error #4: LEAK 88 direct bytes 0x0138b940-0x0138b998 + 168 indirect bytes
# 0 replace_operator_new_array [d:\derek\drmemory\git\src\common\alloc_replace.c:2928]
# 1 main [c:\derek\cgo2016\demos\buggy.cpp:19]

=====
FINAL SUMMARY:

DUPLICATE ERROR COUNTS:
Error # 3: 2

SUPPRESSIONS USED:

ERRORS FOUND:
1 unique, 1 total unaddressable access(es)
1 unique, 1 total uninitialized access(es)
1 unique, 2 total invalid heap argument(s)
0 unique, 0 total GDI usage error(s)
0 unique, 0 total handle leak(s)
0 unique, 0 total warning(s)
1 unique, 1 total, 256 byte(s) of leak(s)
0 unique, 0 total, 0 byte(s) of possible leak(s)

ERRORS IGNORED:
9 unique, 124 total, 10219 byte(s) of still-reachable allocation(s)
(re-run with "-show_reachable" for details)
Details: C:\derek\DynamoRIO-Windows-6.1.0-RC1\drmemory\drmemory\logs\DrMemory-buggy.exe.452.000\results.txt
```

Dr. Fuzz

- In-process function-level fuzzing
- Repeatedly runs a target function, varying the inputs
- Integrated with Dr. Memory for finding errors sooner
 - Also available in a standalone library
- Basic usage:
 - `drmemory -fuzz_function <func> -- <app>`
- Includes typical fuzz testing features:
 - Code coverage-guided fuzzing
 - Input corpus
 - Input dictionary
 - Custom mutators

Dr. Fuzz Results

```
% drrun -t drmemory -batch -fuzz_function some_function -fuzz_num_iters 10 -- ./fuzzme.exe
~~Dr.M~~ Dr. Memory version 1.11.0
~~Dr.M~~ Running "./fuzzme.exe"
in some_function
hello!
in some_function
~~Dr.M~~
~~Dr.M~~ Error #1: UNADDRESSABLE ACCESS beyond heap bounds: reading 0x00f58efc-0x00f58f00 4 byte(s)
~~Dr.M~~ # 0 some_function [c:\derek\cgo2017\demos\fuzzme.c:8]
~~Dr.M~~ # 1 main [c:\derek\cgo2017\demos\fuzzme.c:18]
~~Dr.M~~ Note: @0:00:00.609 in thread 4956
~~Dr.M~~ Note: next higher malloc: 0x00f58f20-0x00f59f20
~~Dr.M~~ Note: refers to 0 byte(s) beyond last valid byte in prior malloc
~~Dr.M~~ Note: prev lower malloc: 0x00f58ef0-0x00f58efc
~~Dr.M~~ Note: instruction: cmp (%eax,%edx,4) $0x00000000
in some_function
in some_function
hello!
in some_function
in some_function
in some_function
in some_function
in some_function
in some_function
in some_function
~~Dr.M~~
~~Dr.M~~ ERRORS FOUND:
~~Dr.M~~ 1 unique, 8 total unaddressable access(es)
```

Opcode Mix Tool

- Monitors every instruction executed
- Counts the dynamic executions of each ISA opcode for the target application workload
- Usage:
 - `drrun -c samples/bin32/opcodes.dll -- <app>`

Opcode Mix Tool Results

```
% bin64/drrun -c samples/bin64/libopcodes.so -- gzip lib64/libdrinjectlib.a
```

```
Client opcodes is running
```

```
Top 15 opcode execution counts in 64-bit AMD64 mode:
```

```
3932460 : shl
4607686 : shr
6736566 : cmovnb
7110994 : test
13250403 : mov
15827664 : xor
17002535 : lea
20301339 : add
23794160 : jnb
24993299 : jz
26312571 : movsxd
28064475 : sub
29210955 : and
37652399 : jnz
69217406 : mov
71688400 : movzx
79491451 : cmp
```


Dr. CPUSim: Legacy CPU Checker

- A tool for testing that applications will run correctly on legacy processors
- Looks for instructions that are illegal on the specified CPU model
- Fools `cpuid` checks
- Usage:
 - `drrun -t drcpusim -cpu PentiumPro -- <app>`

Dr. CPUSim Results

```
% bin32/drrun -t drcpusim -cpu Pentium -continue -- hello.exe  
<Invalid Pentium instruction "movd" @ hello.exe+0x85c3.  
Continuing.>  
<Invalid Pentium instruction "pshufd" @ hello.exe+0x85c7.  
Continuing.>  
<Invalid Pentium instruction "movups" @ hello.exe+0x85ce.  
Continuing.>  
<Invalid Pentium instruction "movdqa" @ hello.exe+0x85f0.  
Continuing.>
```

Dr. Strace: System Call Tracer

- Monitors all system calls executed by a target application
- Prints a trace of system calls and their arguments
- Operates on Windows, Linux, and Android
- Usage:
 - `drrun -t drstrace -- <app>`

Dr. Strace Results

NtGdiGetTextFaceW

```
arg 0: 0x740122ad (type=HANDLE, size=0x4)
arg 1: 0x20 (type=int, size=0x4)
arg 2: 0x001fcd10 (type=<struct>*, size=0x4)
arg 3: 0x0 (type=bool, size=0x4)
```

succeeded =>

```
arg 2: <NYI> (type=<struct>*, size=0x4)
retval: 0x9 (type=int, size=0x4)
```

NtOpenKeyEx

```
arg 0: 0x001fcd0c (type=HANDLE*, size=0x4)
arg 1: 0x109 (type=unsigned int, size=0x4)
arg 2: len=0x18, root=0x3c, name=150/152 "SOFTWARE\Microsoft\Windows
      NT\CurrentVersion\LanguagePack\SurrogateFallback", att=0x40, sd=0x00000000,
      sqos=0x00000000 (type=OBJECT_ATTRIBUTES*, size=0x4)
arg 3: REG_OPTION_RESERVED or REG_OPTION_NON_VOLATILE (type=named constant,
      value=0x0, size=0x4)
```

succeeded =>

```
arg 0: 0x001fcd0c => 0x134 (type=HANDLE*, size=0x4)
retval: 0x0 (type=NTSTATUS, size=0x4)
```

NtQueryKey.KeyCachedInformation

```
arg 0: 0x134 (type=HANDLE, size=0x4)
arg 1: 0x4 (type=named constant, size=0x4)
arg 2: 0x001fcb5c (type=<struct>*, size=0x4)
arg 3: 0xb0 (type=unsigned int, size=0x4)
arg 4: 0x001fca34 (type=unsigned int*, size=0x4)
```

succeeded =>

```
arg 2: _KEY_CACHED_INFORMATION {_LARGE_INTEGER {0x1ca043f05a7a595}, int=0x0, int=0x4,
      int=0x1a, int=0x1, int=0xc, int=0x18, int=0x22} (type=<struct>*, size=0x4)
arg 4: 0x001fca34 => 0x28 (type=unsigned int*, size=0x4)
retval: 0x0 (type=NTSTATUS, size=0x4)
```

Dr. Ltrace: Library Tracer

- Intercepts calls to exported library functions
- Prints a trace of each function called and some argument information
- Usage:
 - `drrun -t drltrace -- <app>`

Dr. Ltrace Results

- Snippet from running an app that prints to the screen:

```
~~~~ ntdll.dll!RtlEnterCriticalSection(0x006fc758, 0x0018fe50)
~~~~ KERNEL32.dll!WriteFile(0x00000268, 0x0018e9b0)
~~~~ KERNELBASE.dll!WriteFile(0x00000268, 0x0018e9b0)
Hello, world!
~~~~ ntdll.dll!RtlLeaveCriticalSection(0x006fc758, 0x0018fe50)
```

Dr. Cov

- Collects code coverage information
- Records which basic blocks have been executed
- Operates on unmodified binaries
- Results are post-processed into the standard lcov format
- Usage:
 - `drrun -t drcov -- <app>`
 - `tools/bin32/drcov2lcov -input drcov*.log`
 - Now pass the resulting data file to lcov's genhtml script

Dr. Cov Results

LCOV - coverage.info - c

file:///C:/derek/cgo2016/demos/drcov/demos/cov.cp

LCOV - code coverage report

Current view: [top level](#) - [demos](#) - [cov.cpp](#) (source / [functions](#))

Test: [coverage.info](#)

Date: [2016-03-10 13:53:55](#)

	Hit	Total	Coverage
Lines:	8	12	66.7 %
Functions:	0	0	-

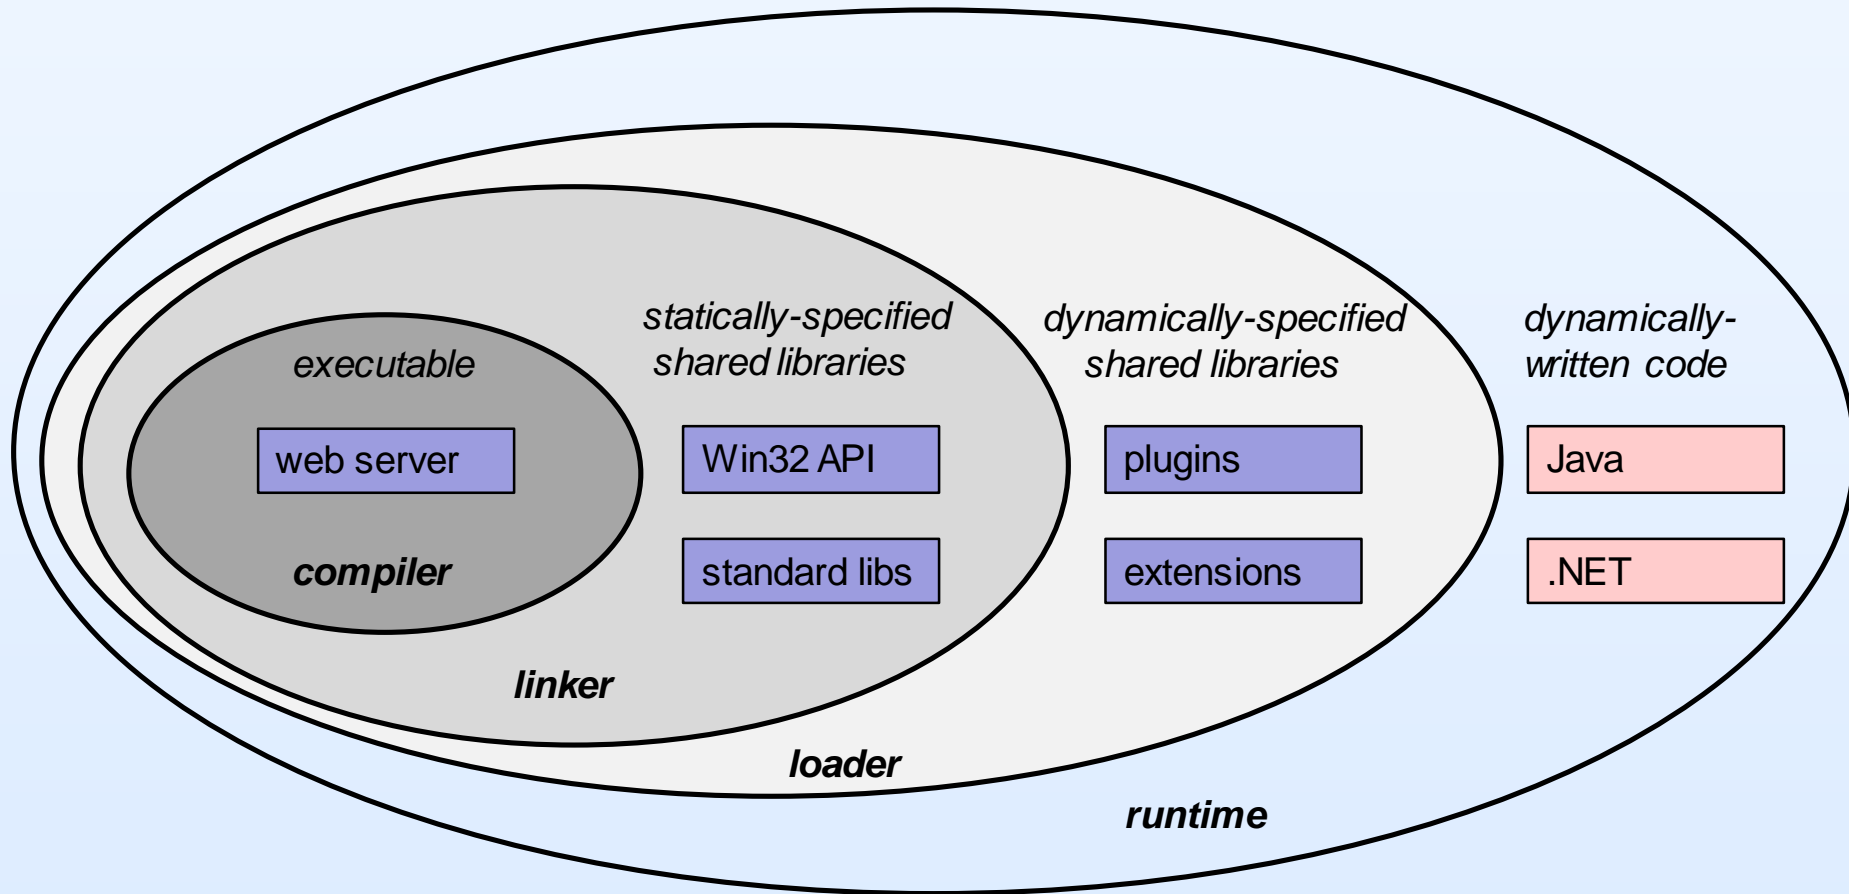
Line data	Source code
1	: #include <iostream>
2	1 : int foo(int x) {
3	1 : return x + 1;
4	1 : }
5	0 : int bar(int y) {
6	0 : return y - 1;
7	0 : }
8	: int main(int argc, const char *argv[])
9	1 : {
10	1 : std::cout << "hello" << foo(argc) << std::endl;
11	1 : if (argc > 2)
12	0 : std::cout << "lots of args" << std::endl;
13	1 : return 0;
14	1 : }

Generated by: [LCOV version 1.11](#)

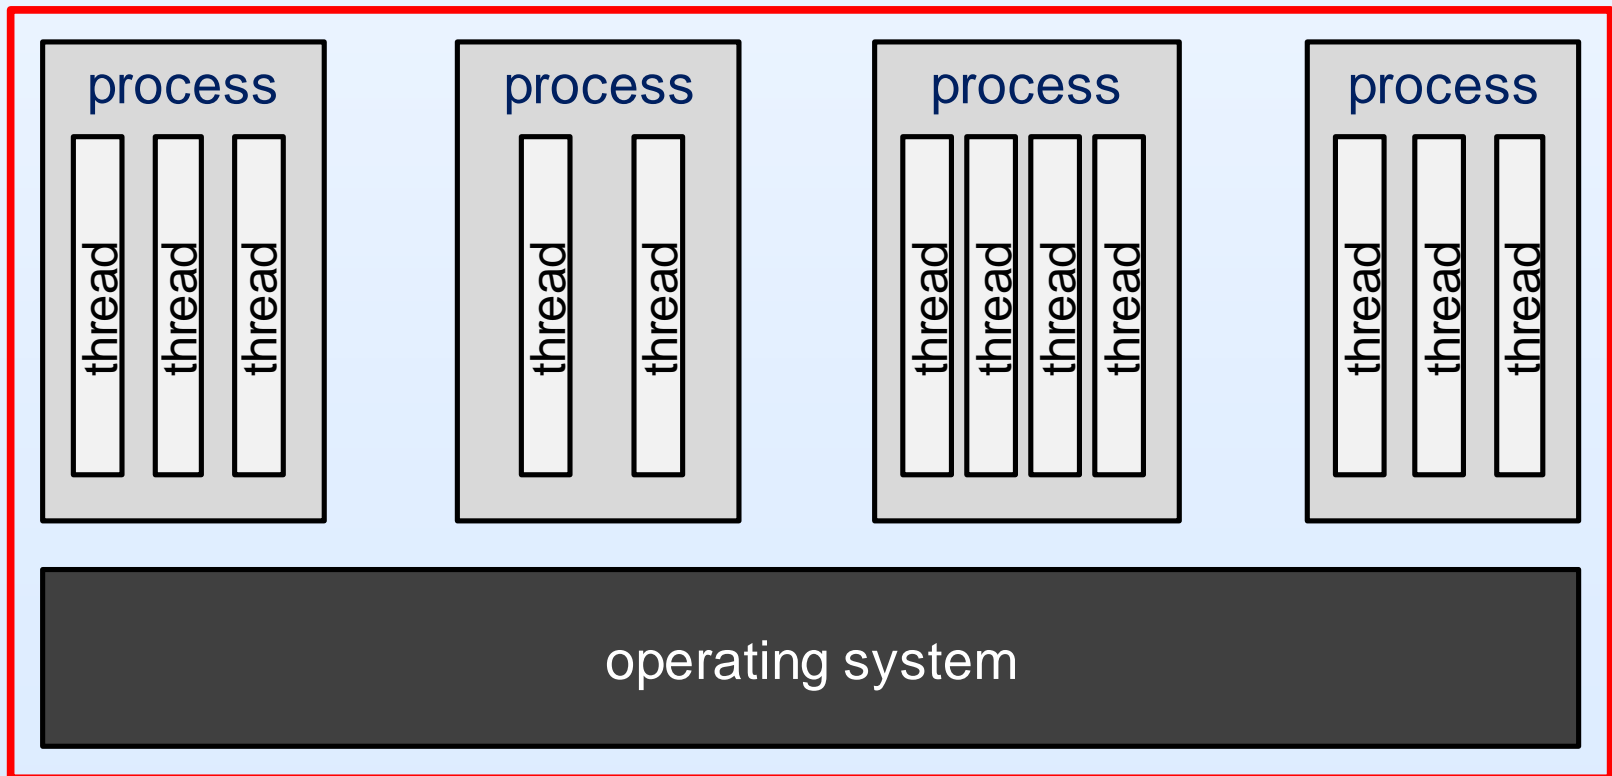
DynamoRIO Overview

8:30- 8:40	Welcome + DynamoRIO History
8:40- 9:10	Tool Demonstrations
9:10- 9:30	DynamoRIO System Overview
9:30- 9:45	DynamoRIO API Part 1
9:45-10:00	ARM Status + Demonstrations
10:00-10:30	<i>Break</i>
10:30-10:55	DynamoRIO API Part 2
10:55-11:20	How to Create Your Own Tool
11:20-11:40	Tool Internals
11:40-11:45	Q & A

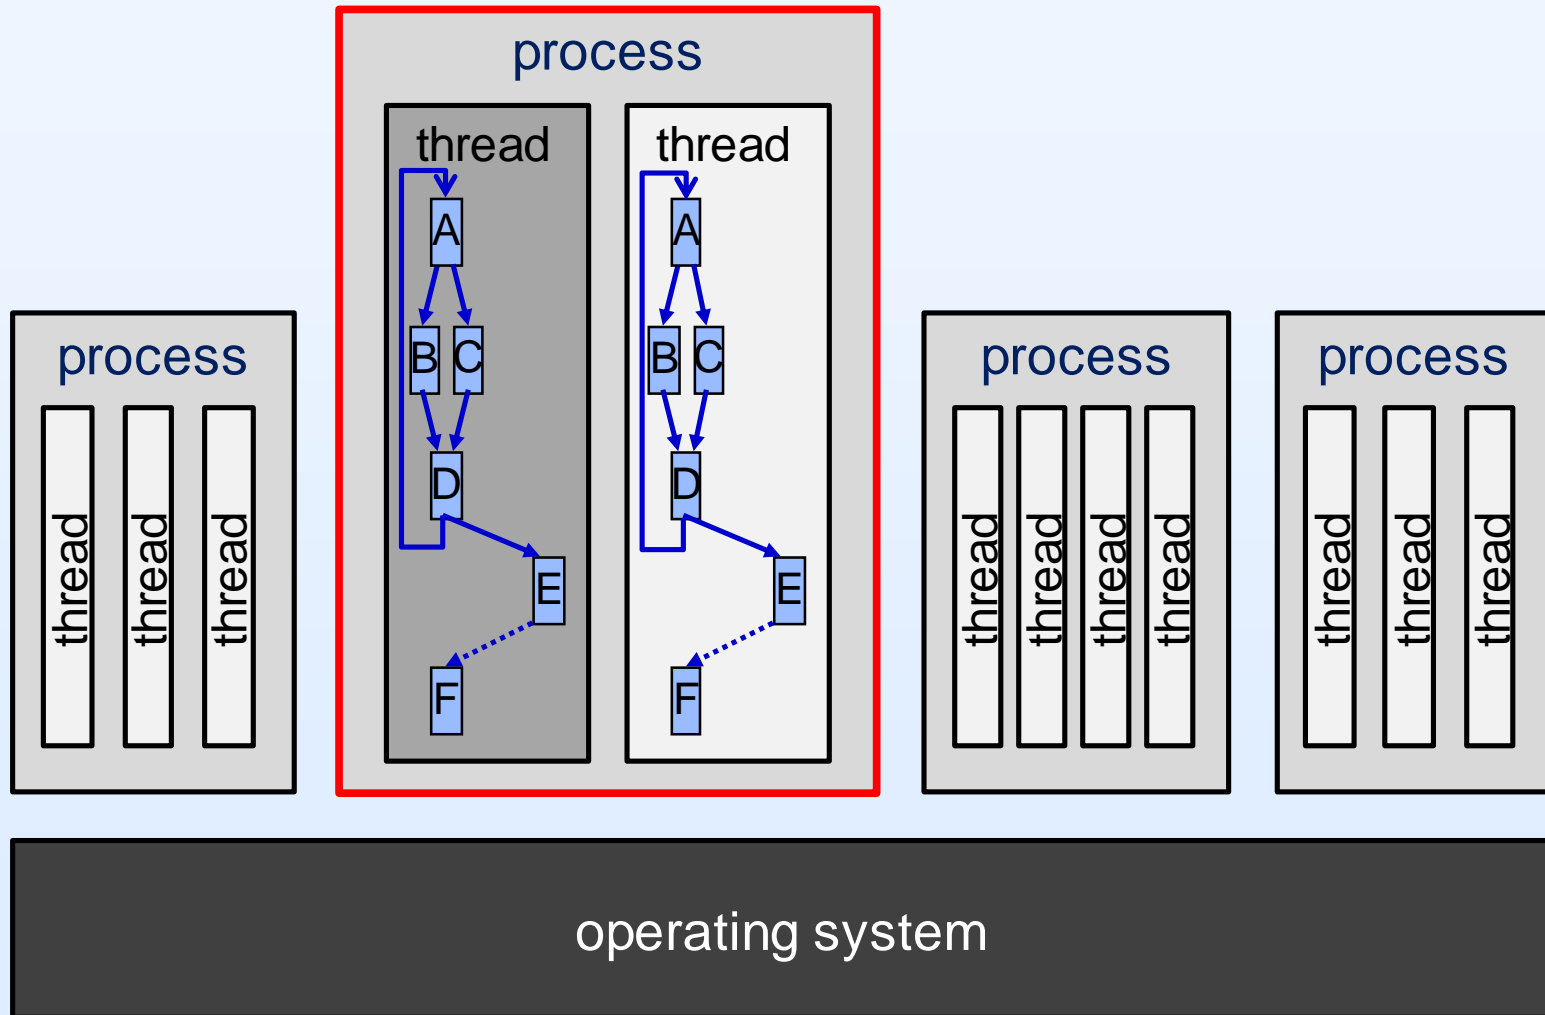
Reach of Toolchain Control Points



System Virtualization



Process Virtualization



Design Goals

- Efficient
 - Near-native performance
- Transparent
 - Match native behavior
- Comprehensive
 - Control every instruction, in any application
- Customizable
 - Adapt to satisfy disparate tool needs

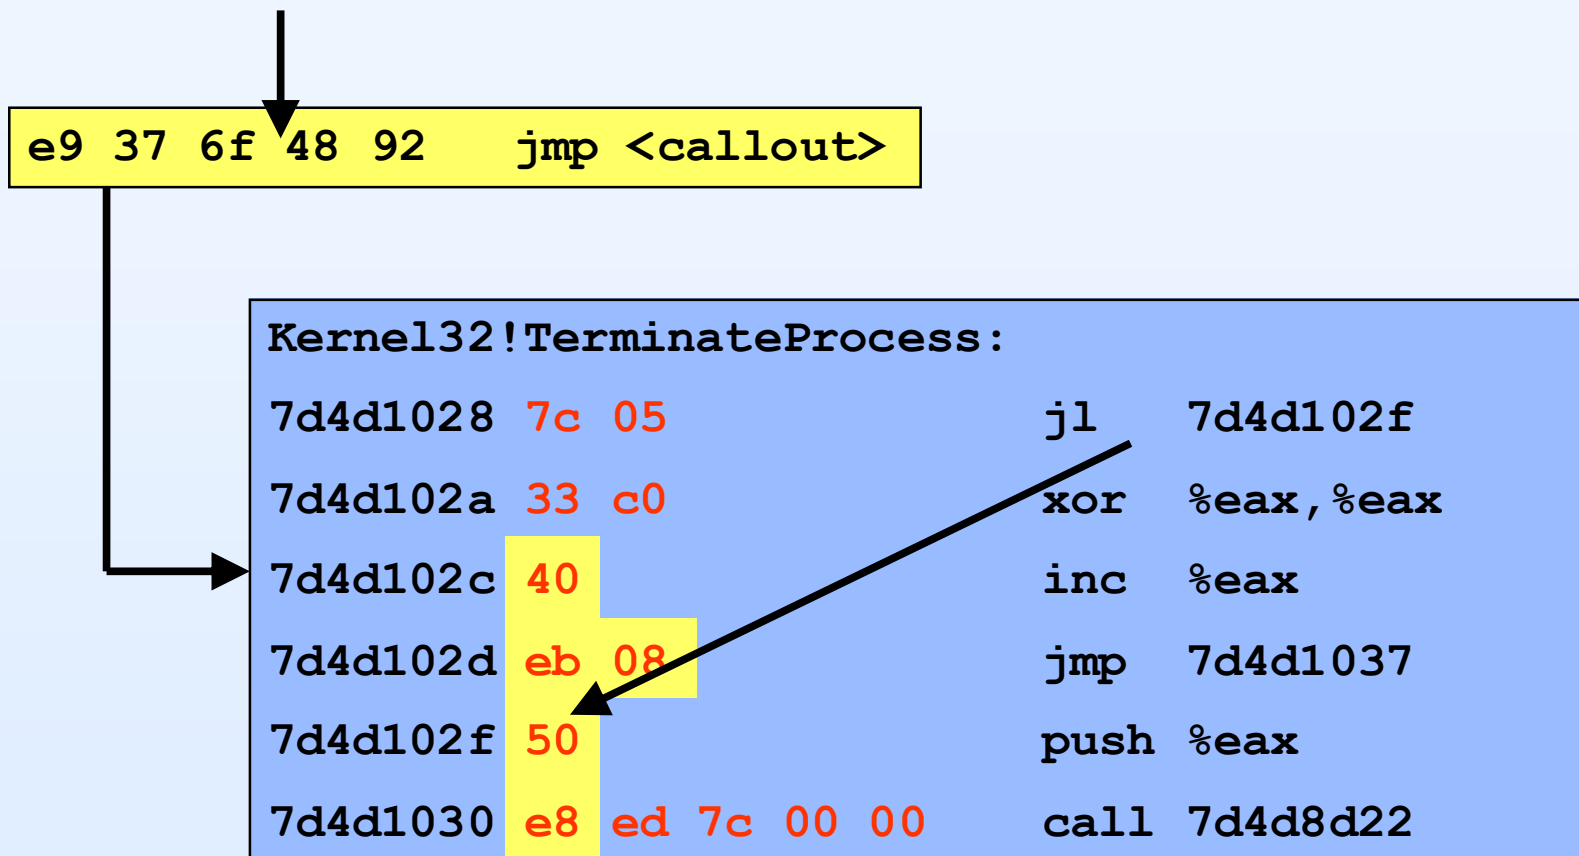
Challenges of Real-World Apps

- Multiple threads
 - Synchronization
- Application introspection
 - Reading of return address
- Transparency corner cases are the norm
 - Example: access beyond top of stack
- Scalability
 - Must adapt to varying code sizes, thread counts, etc.
- Dynamically generated code
 - Performance challenges

Overview Outline

- Efficient
 - Software code cache overview
 - Thread-shared code cache
- Transparent
- Comprehensive
- Customizable

Direct Code Modification



Debugger Trap Too Expensive

cc int3 (breakpoint)

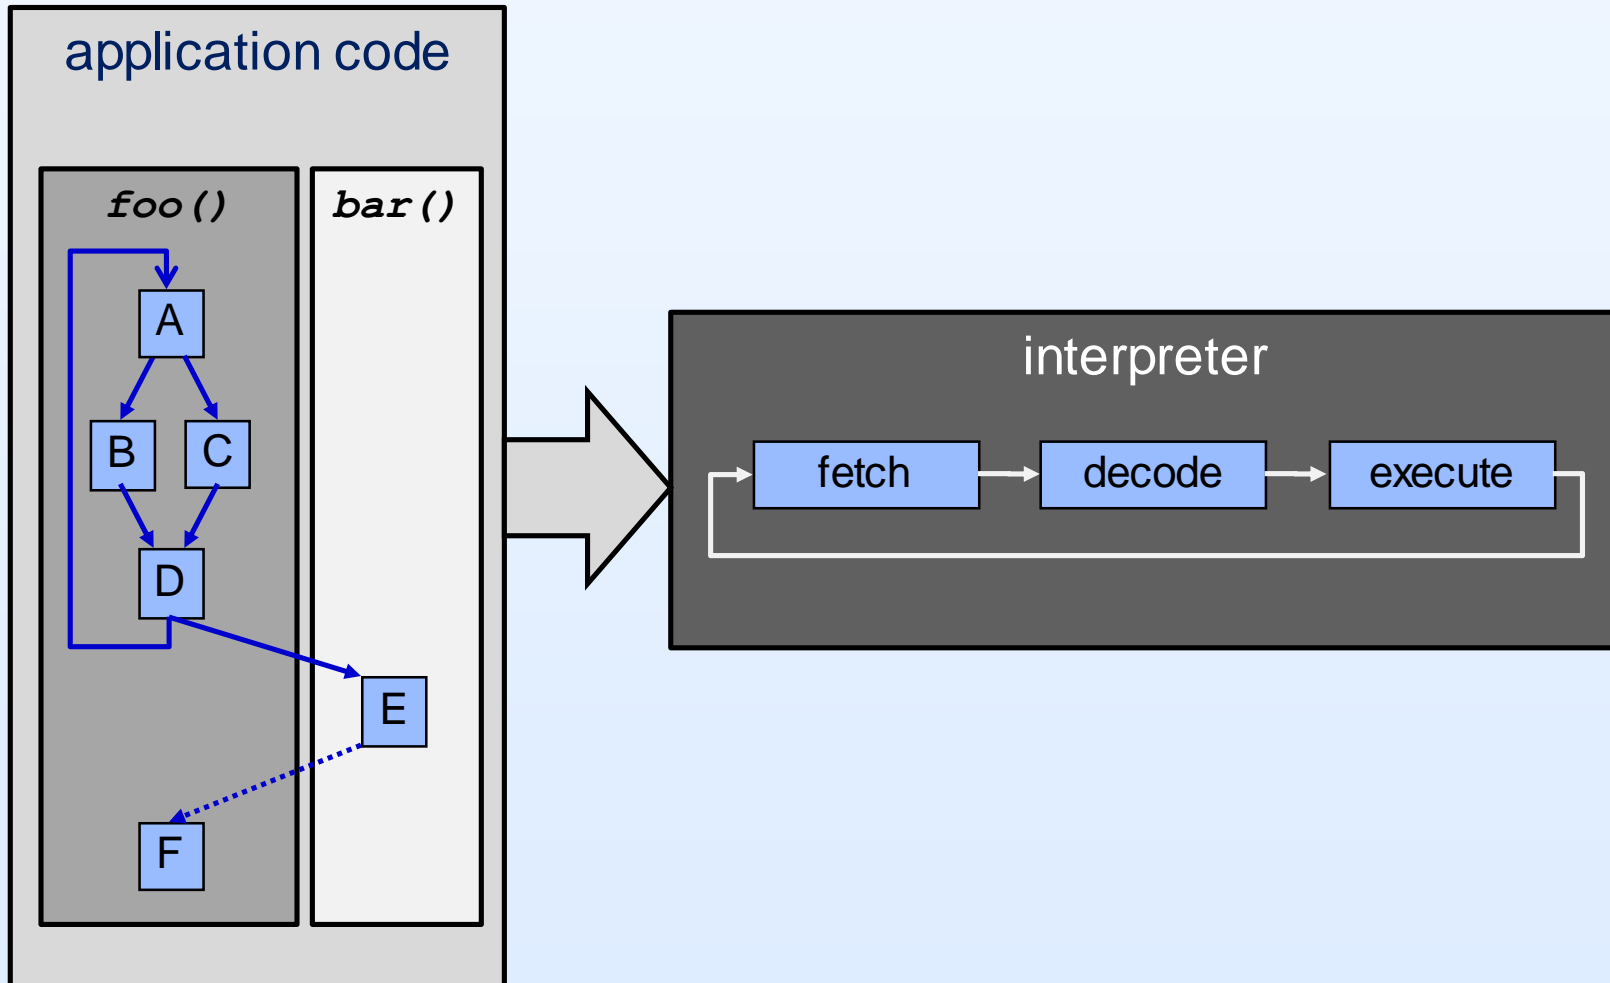
Kernel32!TerminateProcess:

7d4d1028	7c 05	j1	7d4d102f
7d4d102a	33 c0	xor	%eax,%eax
7d4d102c	40	inc	%eax
7d4d102d	eb 08	jmp	7d4d1037
7d4d102f	50	push	%eax
7d4d1030	e8 ed 7c 00 00	call	7d4d8d22

We Need Indirection

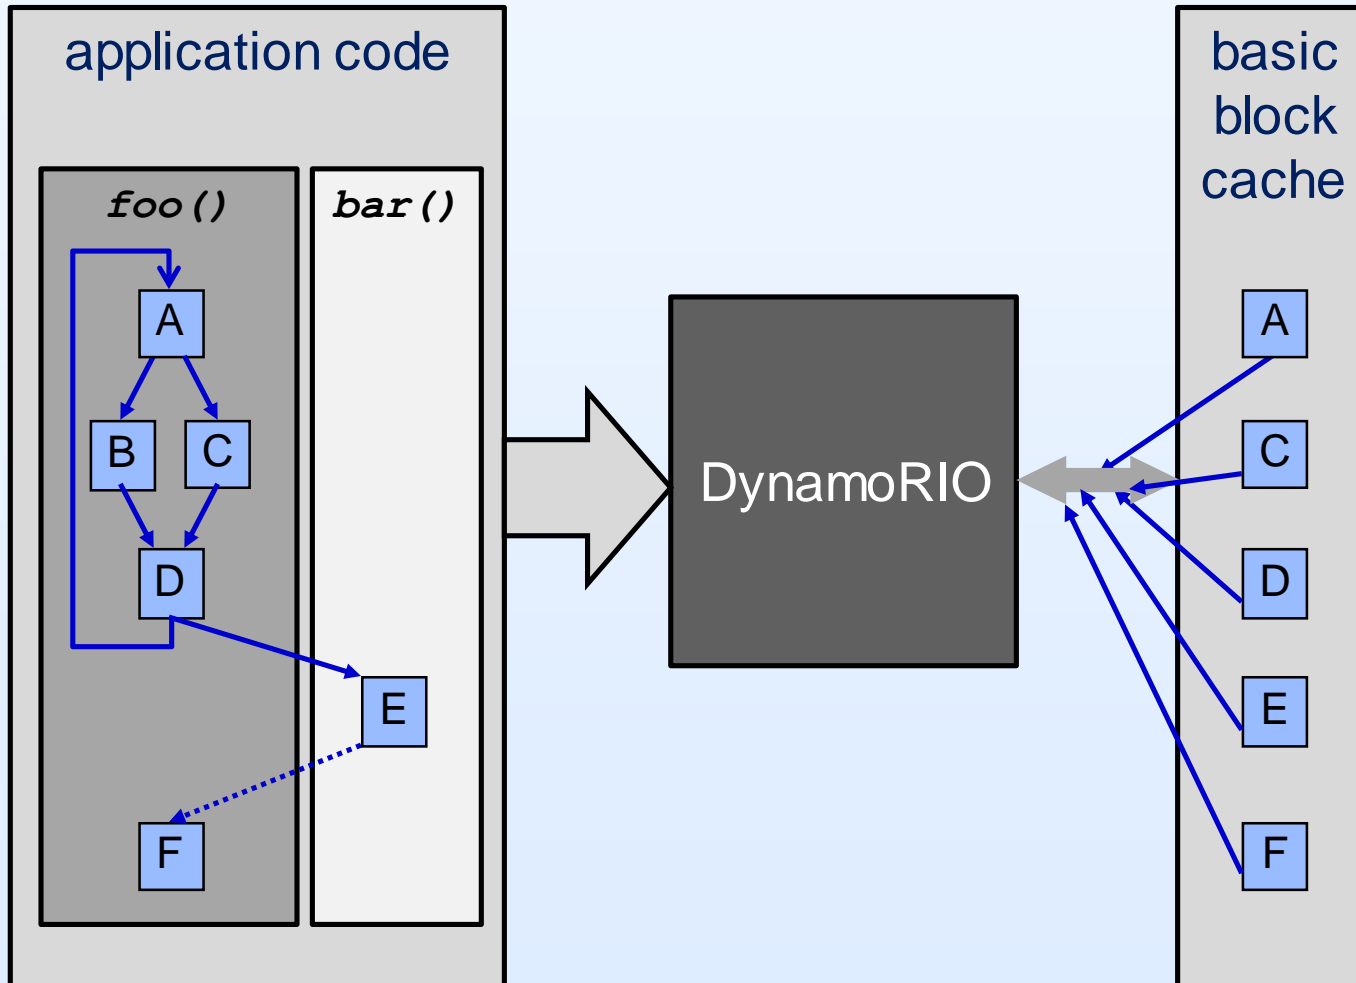
- Avoid transparency and granularity limitations of directly modifying application code
- Allow arbitrary modifications at unrestricted points in code stream
- Allow systematic, fine-grained modifications to code stream
- Guarantee that all code is observed

Basic Interpreter



Slowdown: 300x


Improvement #1: Basic Block Cache



Slowdown: ~~300x~~ 25x

Example Basic Block Fragment

```
add  %eax, %ecx  
cmp  $4, %eax  
jle  $0x40106f
```

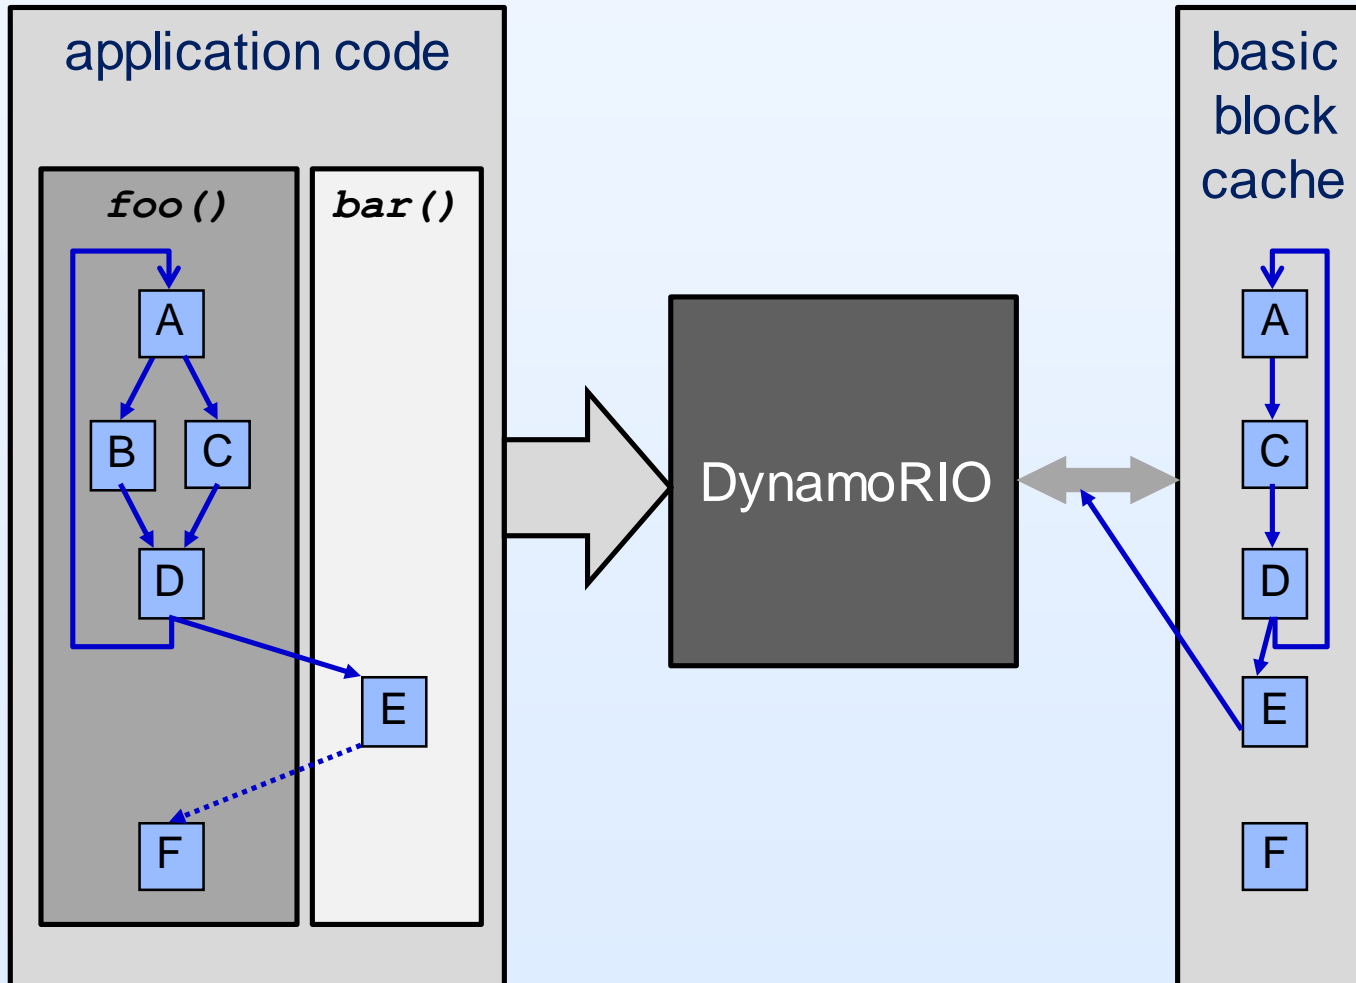


```
frag7: add  %eax, %ecx  
       cmp  $4, %eax  
       jle  <stub0>  
       jmp  <stub1>  
  
stub0: mov  %eax, eax-slot  
       mov  &dstub0, %eax  
       jmp  context_switch  
  
stub1: mov  %eax, eax-slot  
       mov  &dstub1, %eax  
       jmp  context_switch
```

dstub0
target: 0x40106f

dstub1
target: fall-thru


Improvement #2: Linking Direct Branches



Slowdown: ~~300x~~ ~~25x~~ 3x

Direct Linking

```
add  %eax, %ecx  
cmp  $4, %eax  
jle  $0x40106f
```

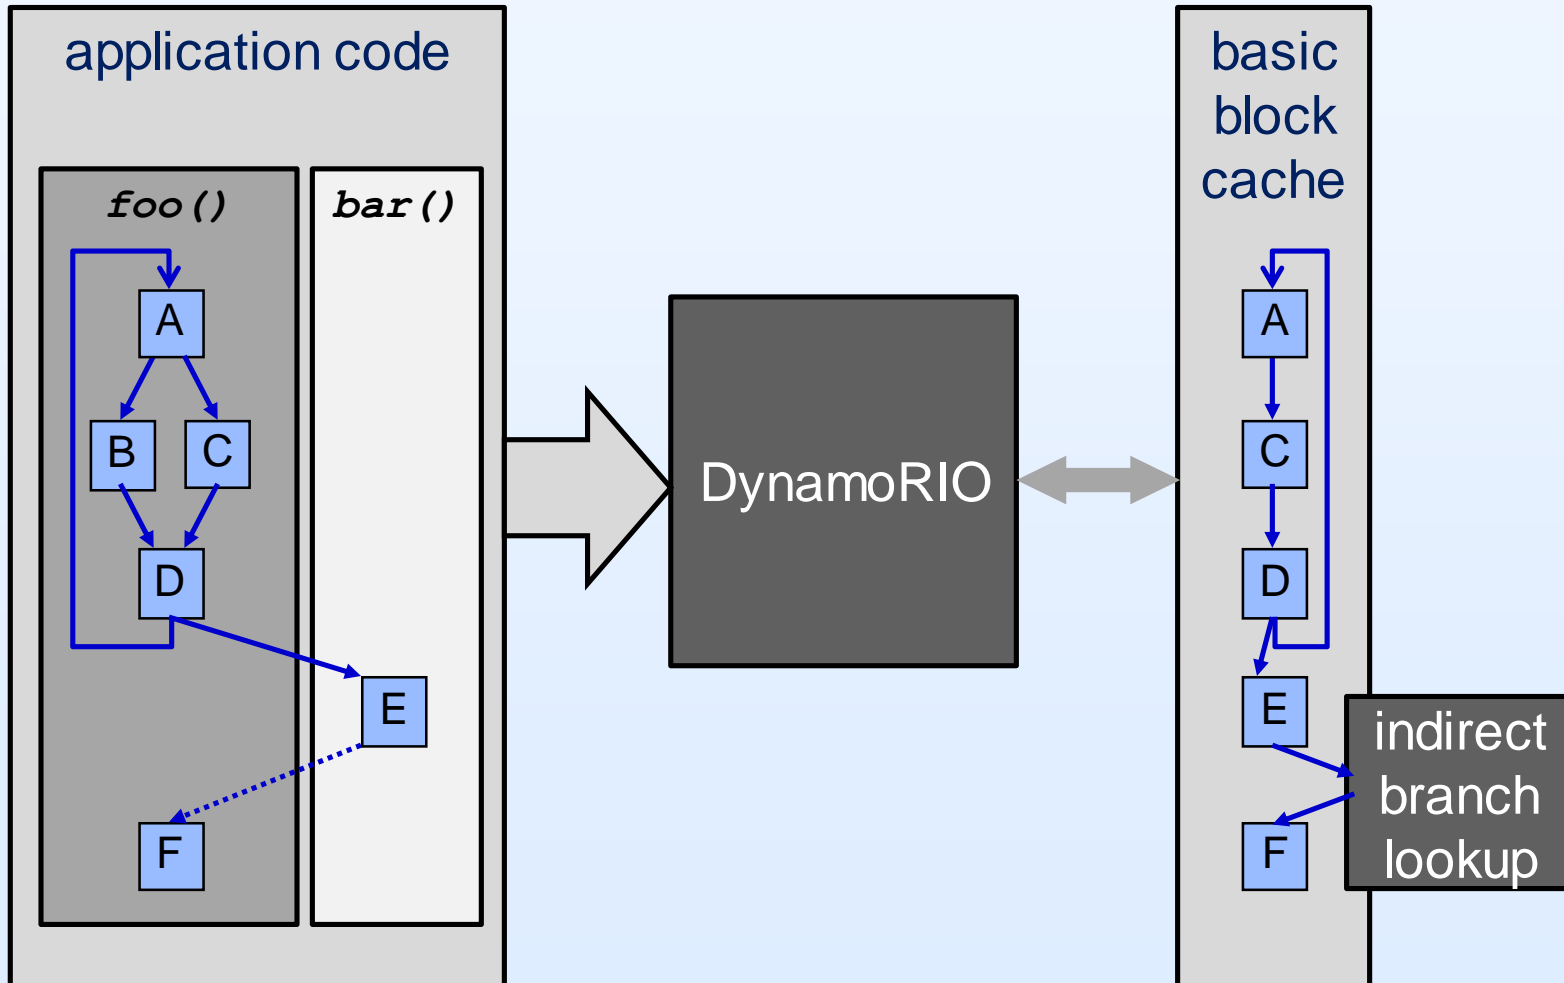


```
frag7: add  %eax, %ecx  
       cmp  $4, %eax  
       jle  <frag8>  
       jmp  <stub1>  
  
stub0: mov  %eax, eax-slot  
       mov  &dstub0, %eax  
       jmp  context_switch  
  
stub1: mov  %eax, eax-slot  
       mov  &dstub1, %eax  
       jmp  context_switch
```

dstub0
target: 0x40106f

dstub1
target: fall-thru


Improvement #3: Linking Indirect Branches



Slowdown: ~~300x~~ ~~25x~~ ~~3x~~ 1.2x

Indirect Branch Transformation

`ret`




```
frag8: mov %ecx, ecx-slot  
       pop %ecx  
       jmp <ib_lookup>
```


```
ib_lookup: ...  
          ...  
          ...
```

Improvement #4: Trace Building

basic block cache






trace cache






- Traces reduce branching, improve layout and locality, and facilitate optimizations across blocks
 - We avoid indirect branch lookup
- Next Executing Tail (NET) trace building scheme [Duesterwald 2000]

Incremental NET Trace Building

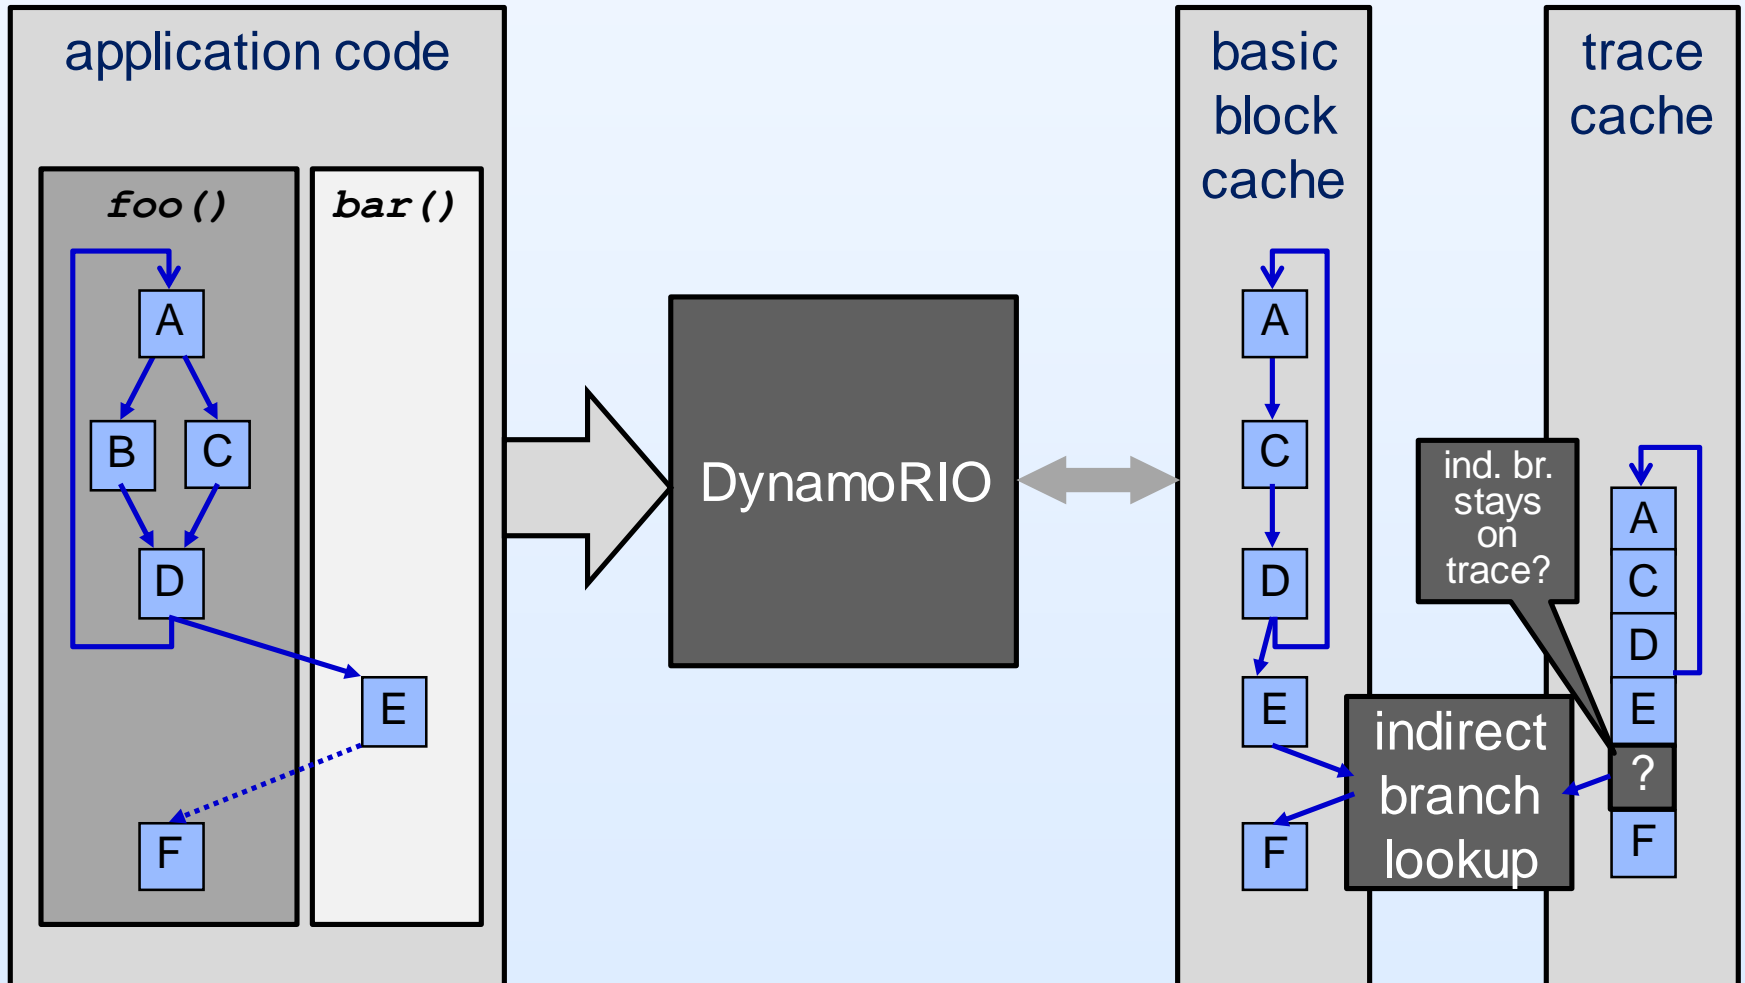
basic block cache



trace cache




Improvement #4: Trace Building




Slowdown: ~~300x~~ ~~25x~~ ~~3x~~ ~~1.2x~~ 1.1x

Base Performance



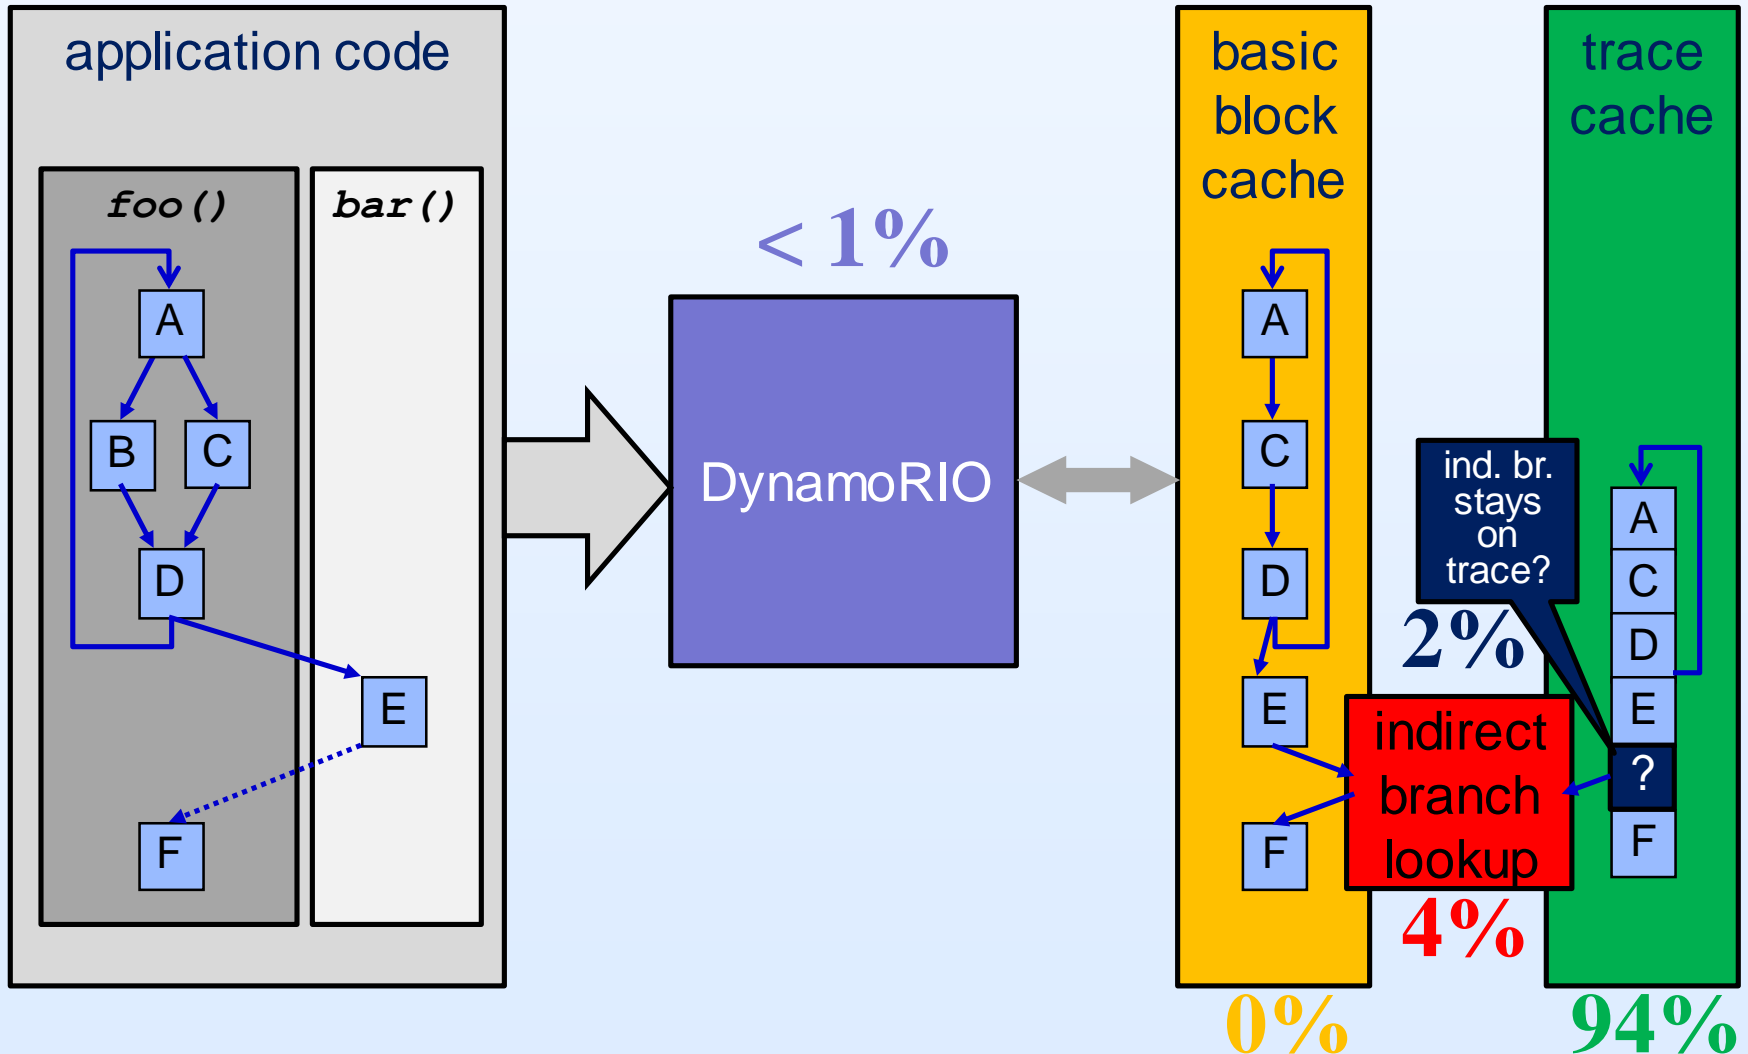
Base Performance: SPEC 2006




Sources of Overhead

- Extra instructions
 - Indirect branch target comparisons
 - Indirect branch hashtable lookups
- Extra data cache pressure
 - Indirect branch hashtable
- Branch mispredictions
 - Calls and returns turn into jumps on some architectures
- Application code modification

Time Breakdown for SPEC CPU INT



Avoiding Intermediate Layers



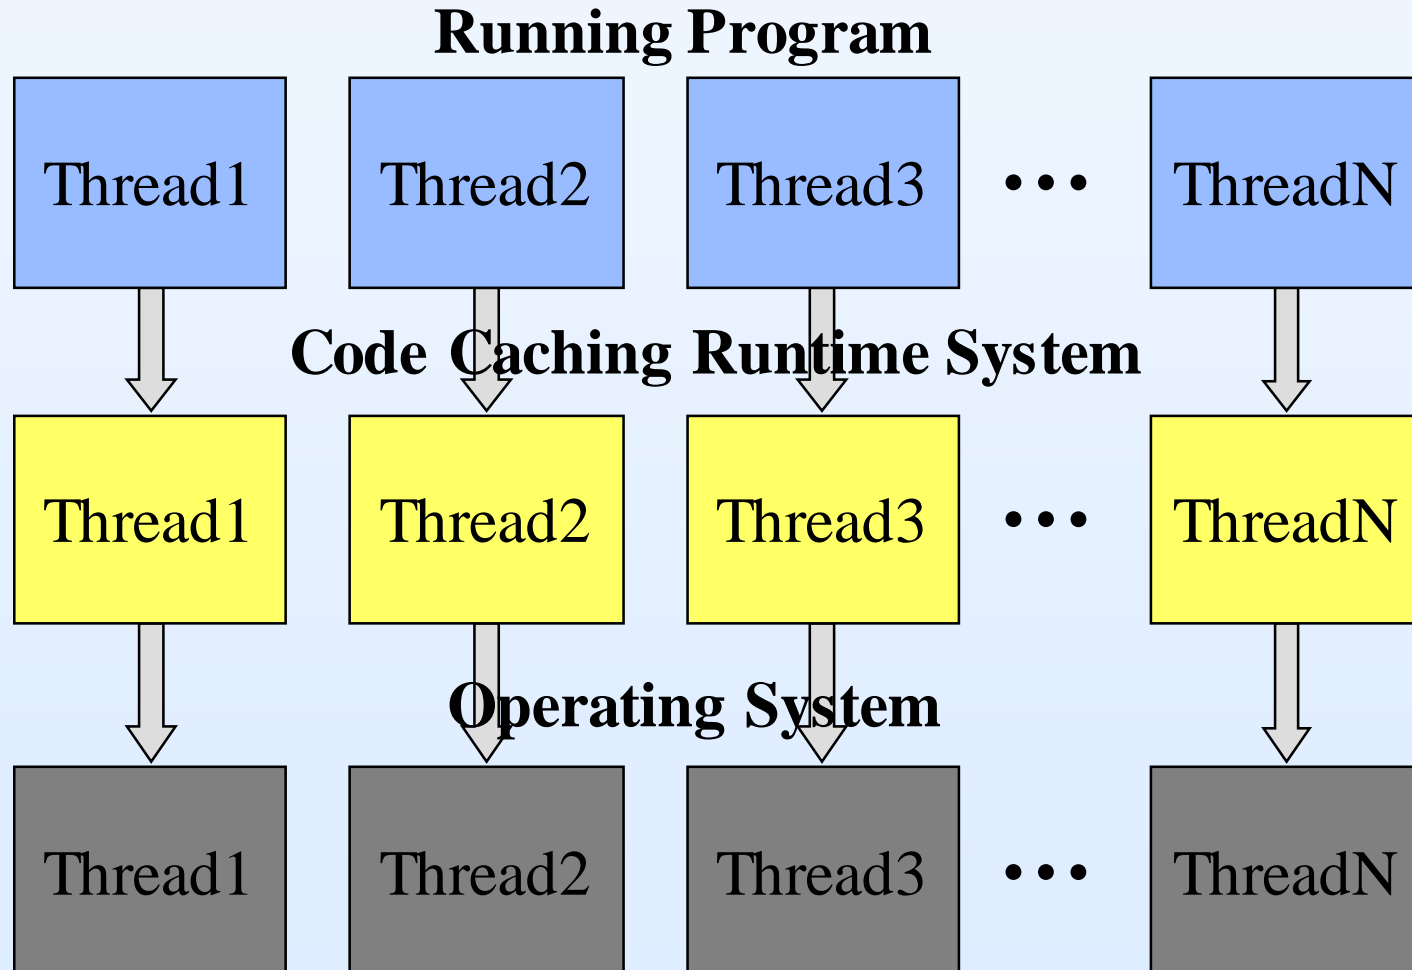
No Intermediate Layers

- IR mirrors underlying ISA
 - Preserves optimized application code
 - Intermediate layers incur significant performance impact
 - QEMU (user-mode) **6x** slower than DR, Valgrind **4x** slower than DR
 - *This is the key to good performance*
- Still have an abstraction layer
 - Block or trace = list of instructions
 - Instruction = lists of source and destination operands
 - Tool code often still cross-platform
 - “Does this instruction read memory?”

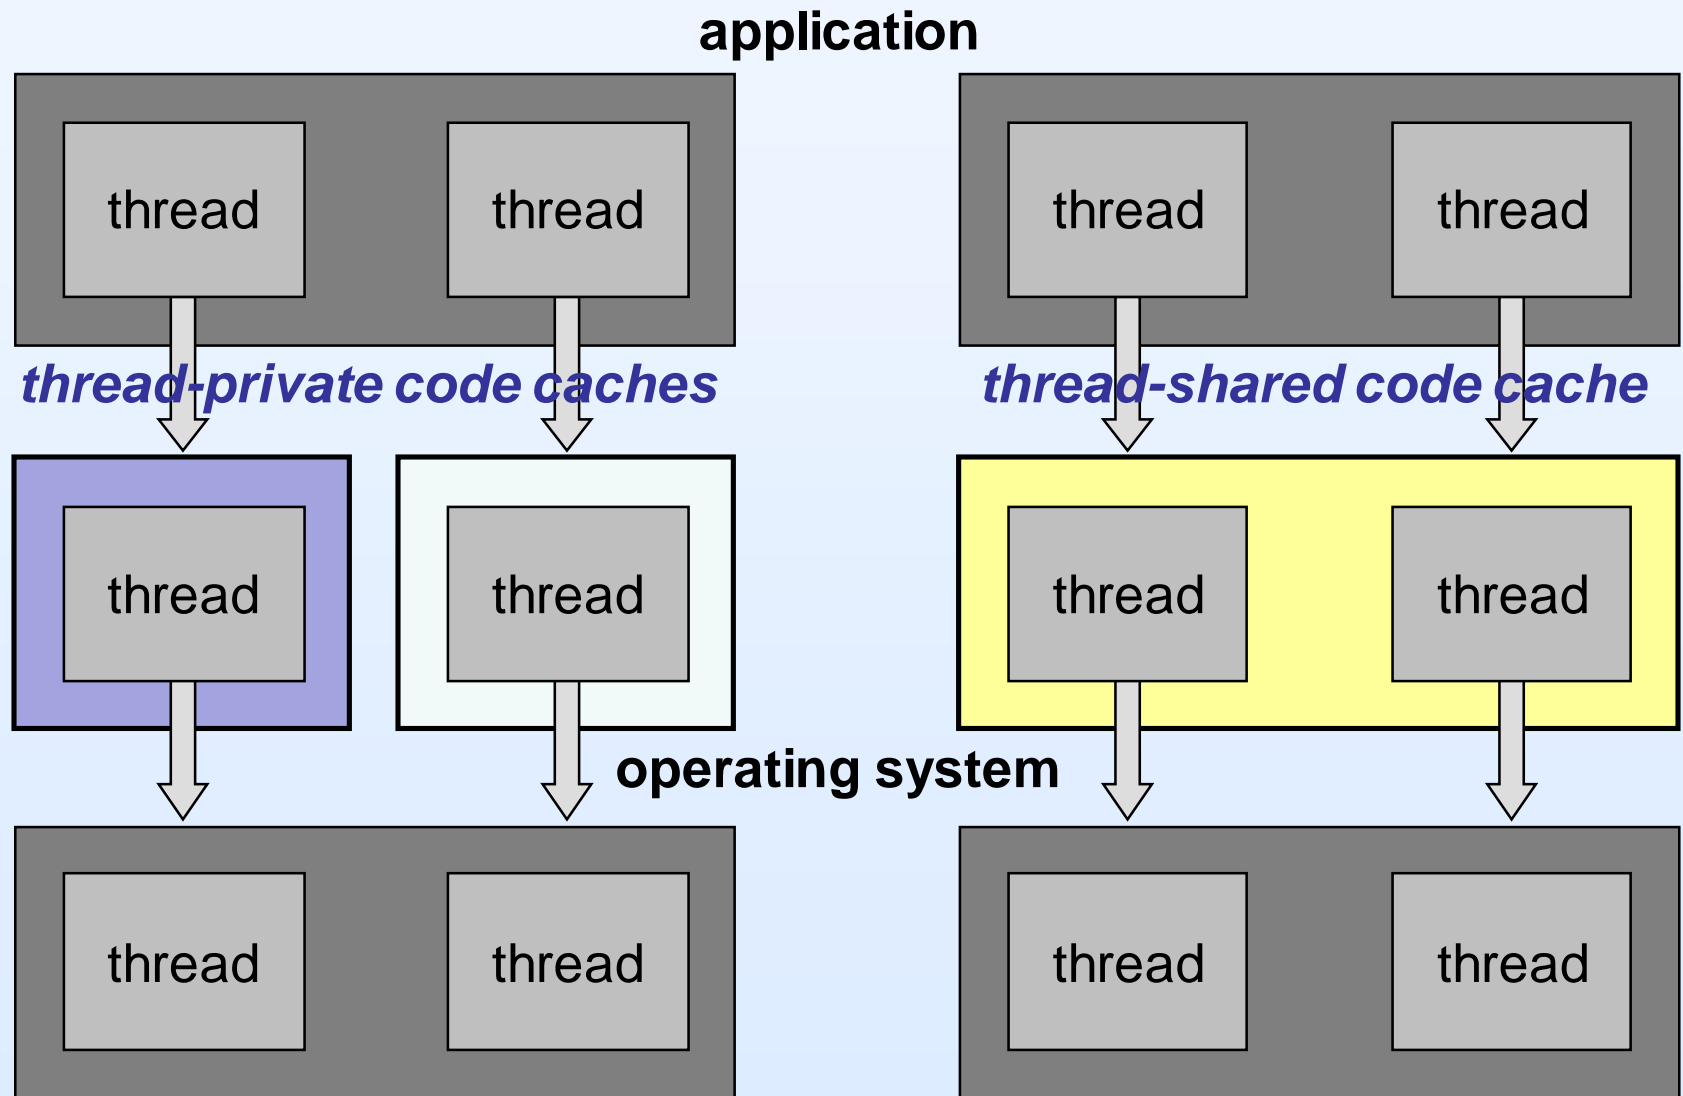
Overview Outline

- Efficient
 - Software code cache overview
 - Thread-shared code cache
- Transparent
- Comprehensive
- Customizable

Threading Model



Code Cache Threading Models




Thread-Private versus Thread-Shared

- Thread-private
 - Less synchronization needed
 - Absolute addressing for thread-local storage
 - Thread-specific optimization and instrumentation
- Thread-shared
 - Scales to many-threaded apps

Overview Outline

- Efficient
- Transparent
 - Transparency principles
 - Cache consistency
 - Synchronization
- Comprehensive
- Customizable

Unavoidably Intrusive



Transparency

- Do not want to interfere with the semantics of the program
- Dangerous to make any assumptions about:
 - Register usage
 - Calling conventions
 - Stack layout
 - Memory/heap usage
 - I/O and other system call use

Painful, But Necessary

- Difficult and costly to handle corner cases
- Many applications will not notice...
- ...but some will!
 - Microsoft Office: Visual Basic generated code, stack convention violations
 - COM, Star Office, MMC: trampolines
 - Adobe Premiere: self-modifying code
 - VirtualDub: UPX-packed executable
 - etc.

Transparency Principles

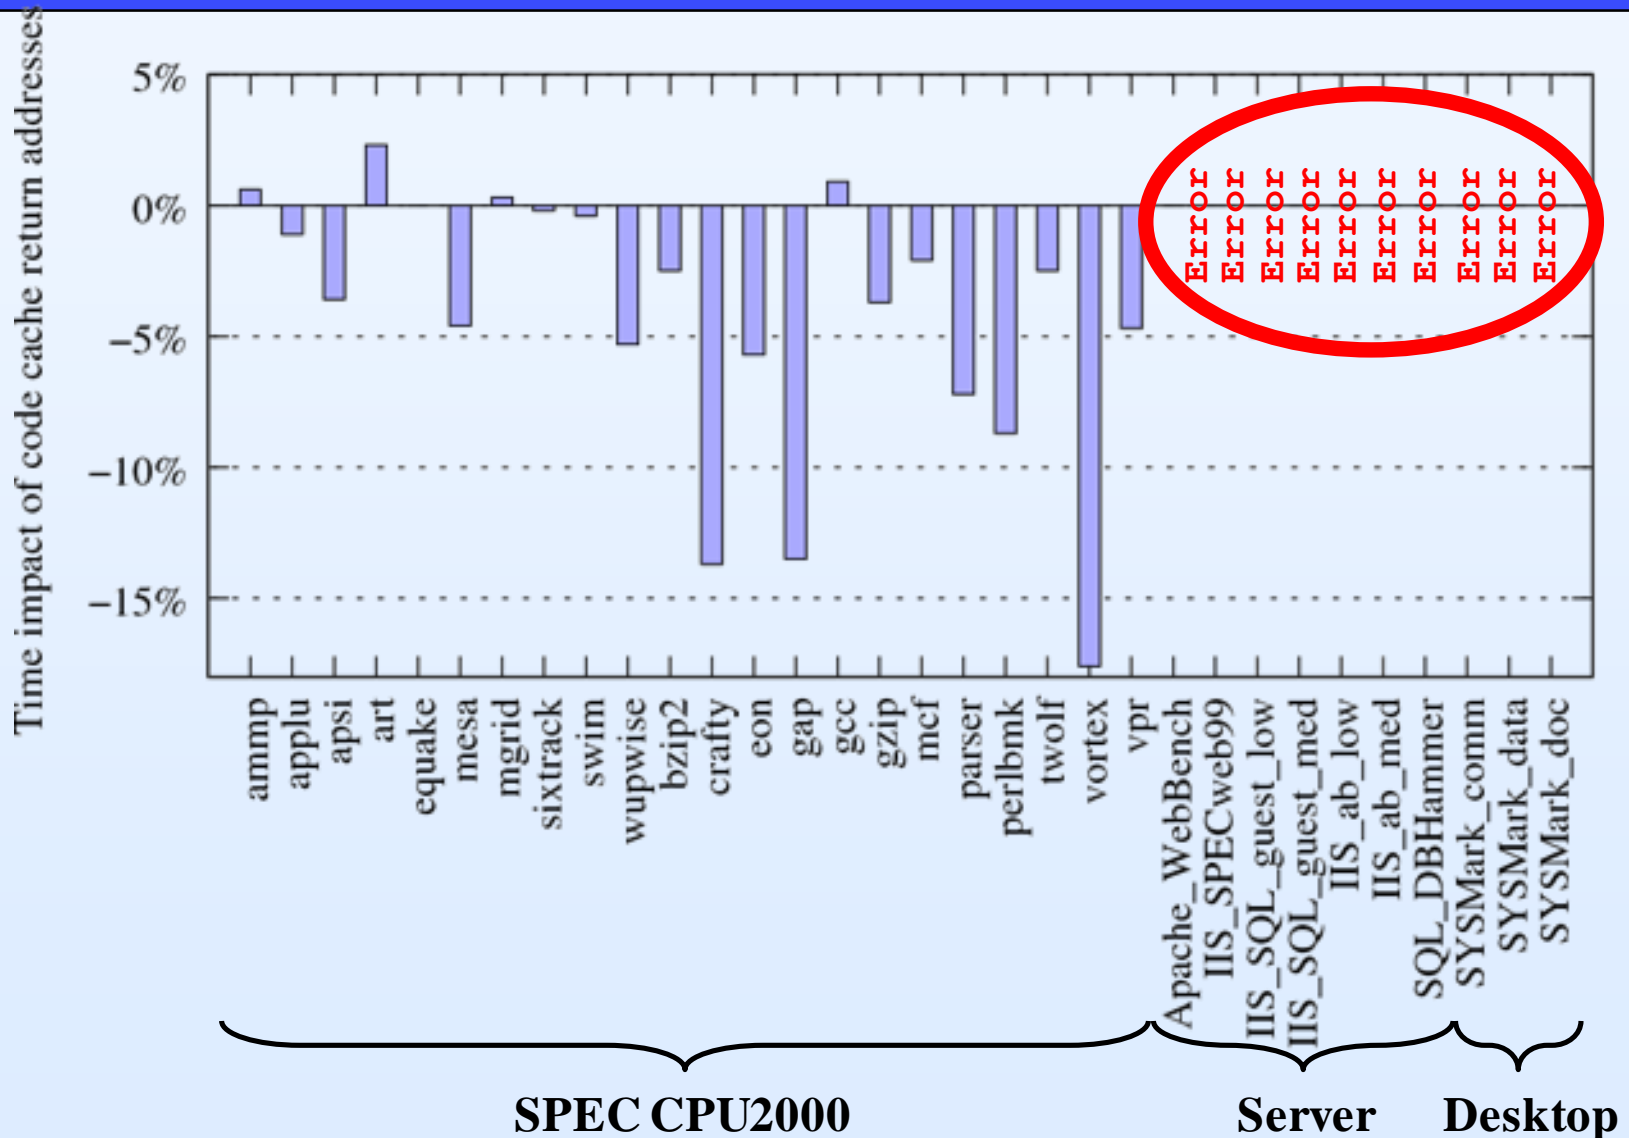
- **Principle 1: *As few changes as possible***
 - Set a high bar for value before changing the native environment
- **Principle 2: *Hide necessary changes***
 - Whatever is valuable enough to change must be hidden
 - Changes that cannot be hidden should not be made
- **Principle 3: *Separate resources***
 - Avoid intra-process resource conflicts

Bruening et al. "Transparent Dynamic Instrumentation" VEE'12

Principle 1: *As few changes as possible*

- Application code
- Executable on disk
- Stored addresses
- Threads
- Application data
 - Including the stack!

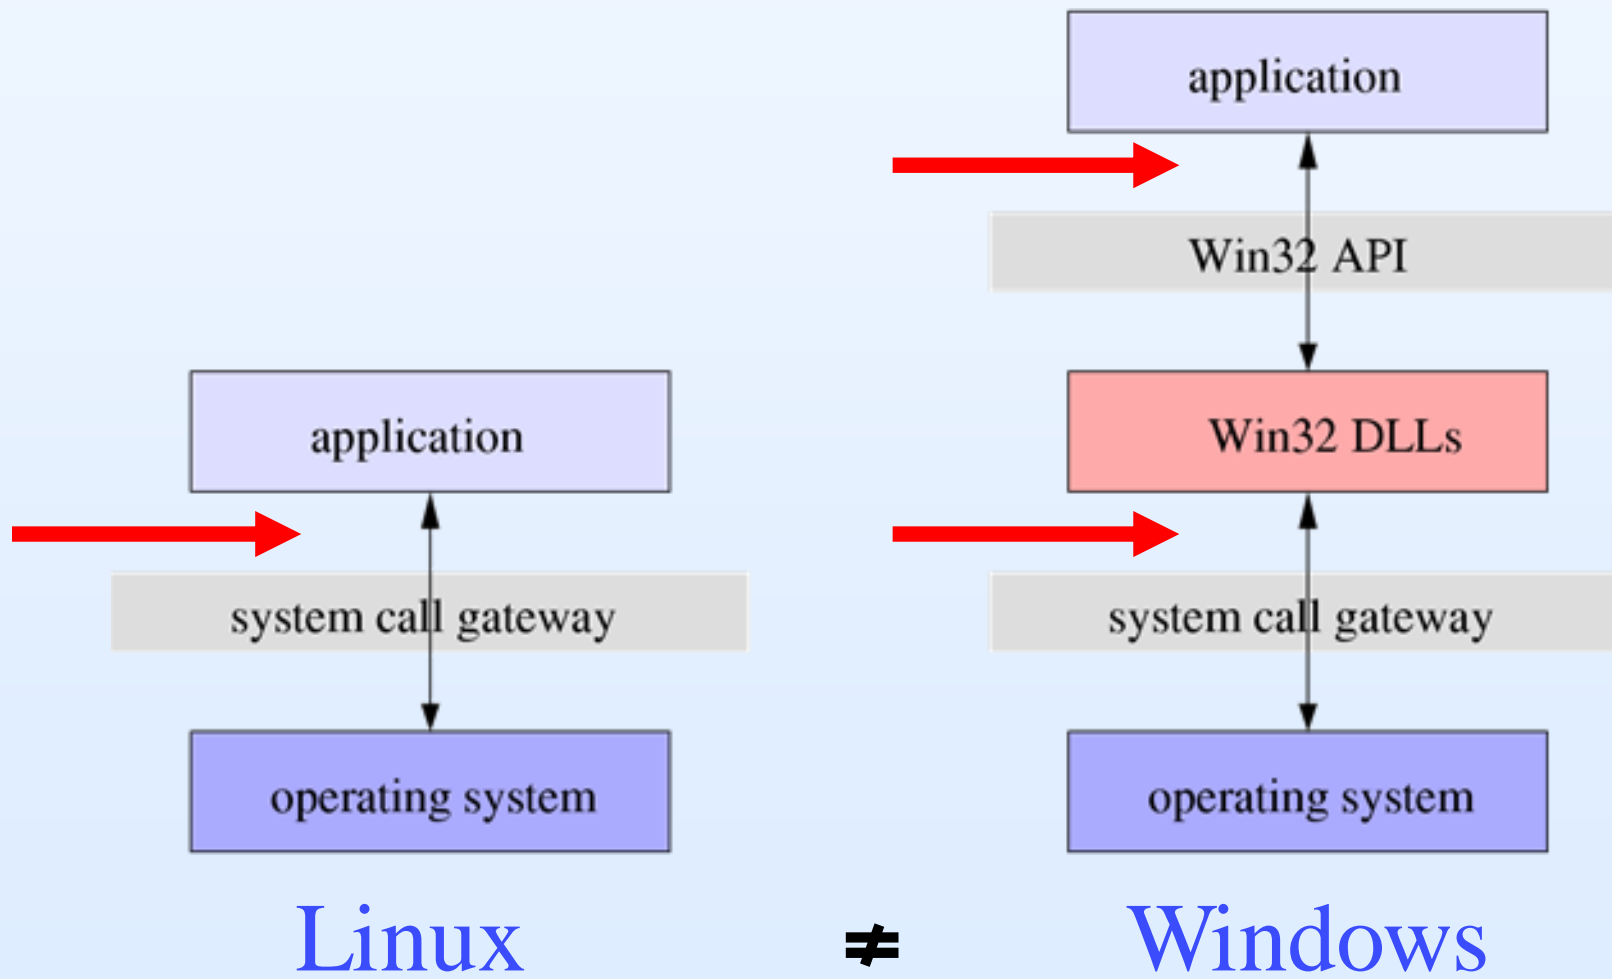
Return Address Transparency



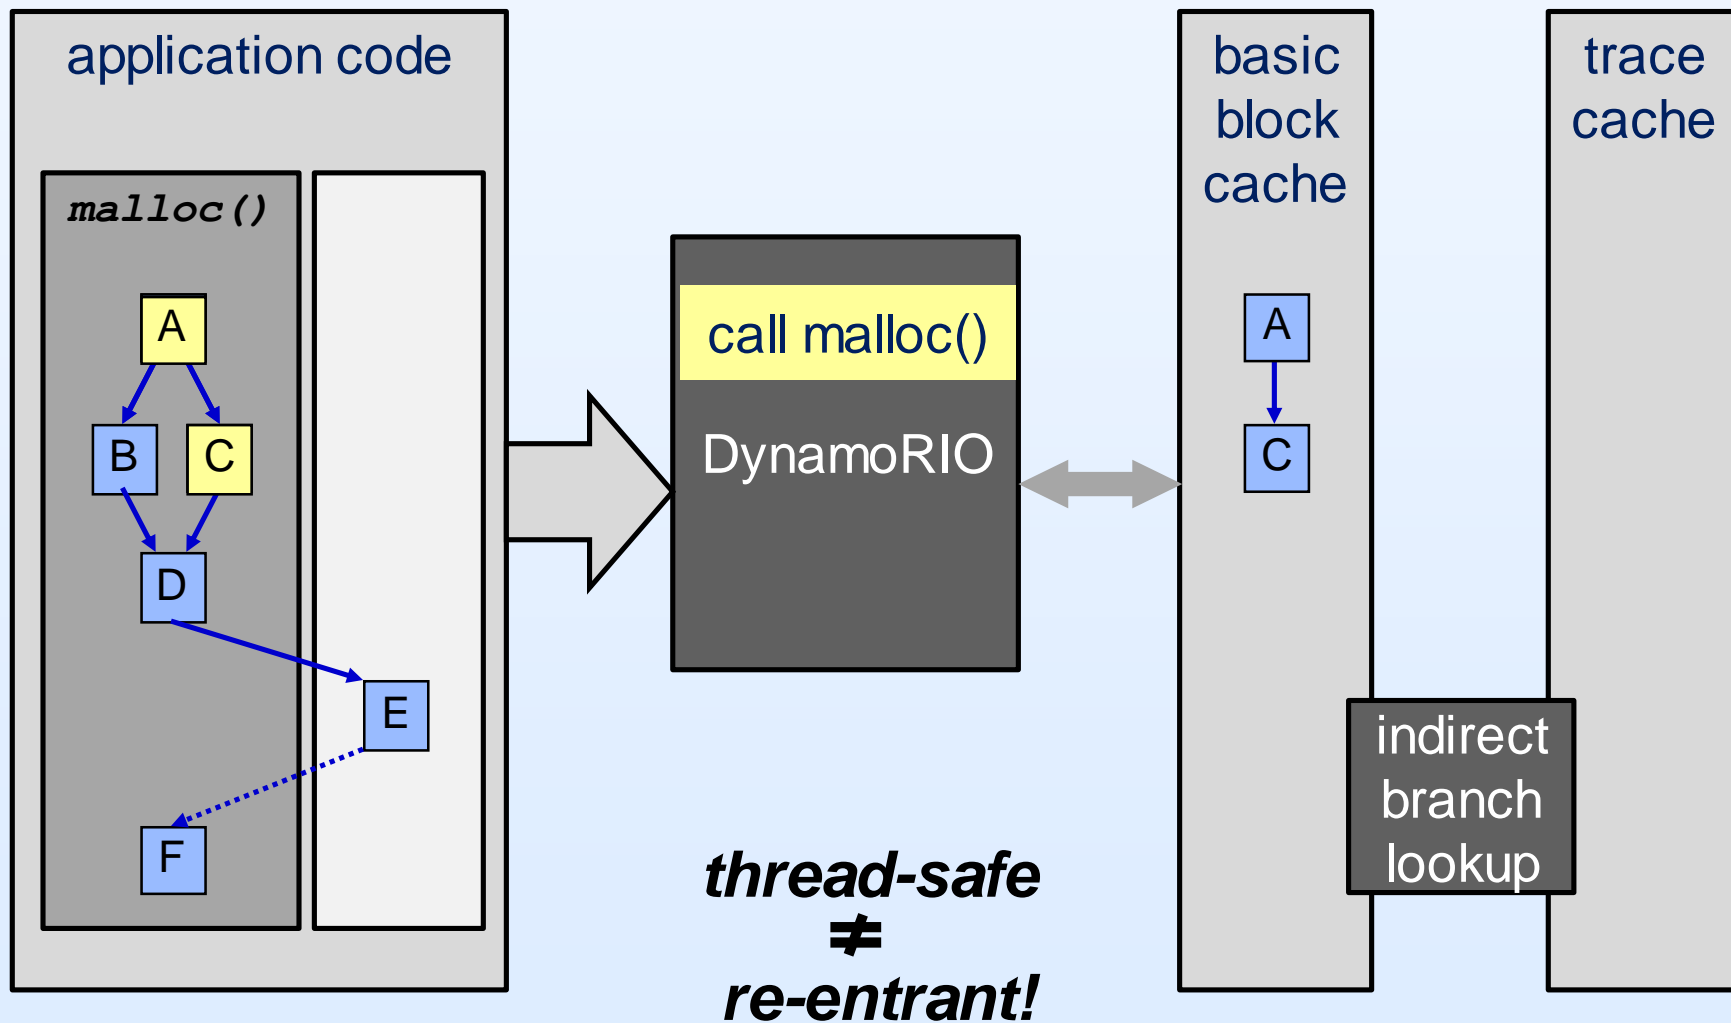
Principle 2: *Hide necessary changes*

- Application addresses
- Address space
- Error transparency
- Code cache consistency

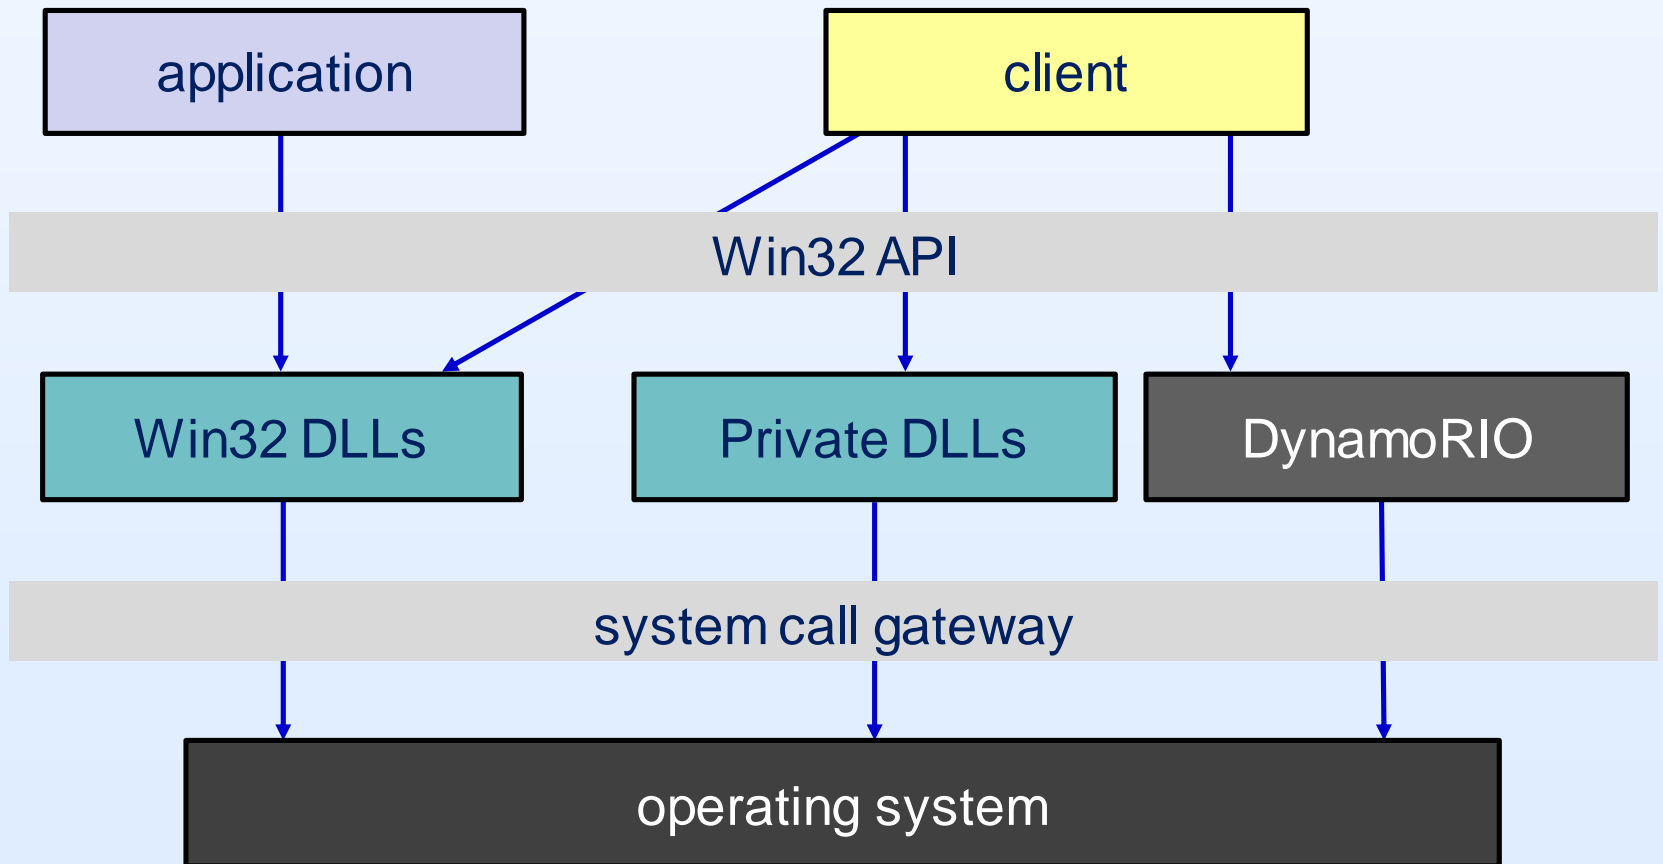
Principle 3: *Separate resources*



Arbitrary Interleaving



Private Libraries



Transparency Landscape


	Principle 1: <i>As few changes as possible</i>	Principle 2: <i>Hide necessary changes</i>	Principle 3: <i>Separate resources</i>
Code	application code, stored addresses	machine context, cache consistency	
Data	stack, heap, registers, condition flags		separate stack, heap, context, i/o
Concurrency	threads, memory ordering		disjoint locks
Other		preserve errors	

Overview Outline


- Efficient
- Transparent
 - Transparency principles
 - Cache consistency
 - Synchronization
- Comprehensive
- Customizable

Code Change Mechanisms

RISC



x86



How Often Does Code Change?

- Not just modification of code!
- Removal of code
 - Shared library unloading
- Replacement of code
 - JIT region re-use
 - Trampoline on stack

Code Change Events

	Memory Unmappings	Generated Code Regions	Modified Code Regions
SPECFP	112	0	0
SPECINT	29	0	0
SPECJVM	7	3373	4591
Excel	144	21	20
Photoshop	1168	40	0
Powerpoint	367	28	33
Word	345	20	6

Adaptive Consistency Algorithm

- Use page protection by default
 - Most code regions are always read-only
- Subdivide written-to regions to reduce flushing cost of write-execute cycle
 - Large read-only regions, small written-to regions
- Switch to instrumentation if write-execute cycle repeats too often (or on same page)
 - Switch back to page protection if writes decrease

Bruening et al. "Maintaining Consistency and Bounding Capacity of Software Code Caches" CGO'05

Overview Outline

- Efficient
- Transparent
 - Transparency principles
 - Cache consistency
 - Synchronization
- Comprehensive
- Customizable


Synchronization Transparency

- Application thread management should not interfere with the runtime system, and vice versa
 - Cannot allow the app to suspend a thread holding a runtime system lock
 - Runtime system cannot use app locks
- Disjoint locks
 - App thread suspension requires safe spots where no runtime system locks are held
 - Time spent in the code cache can be unbounded
 - Our invariant: no runtime system lock can be held while executing in the code cache

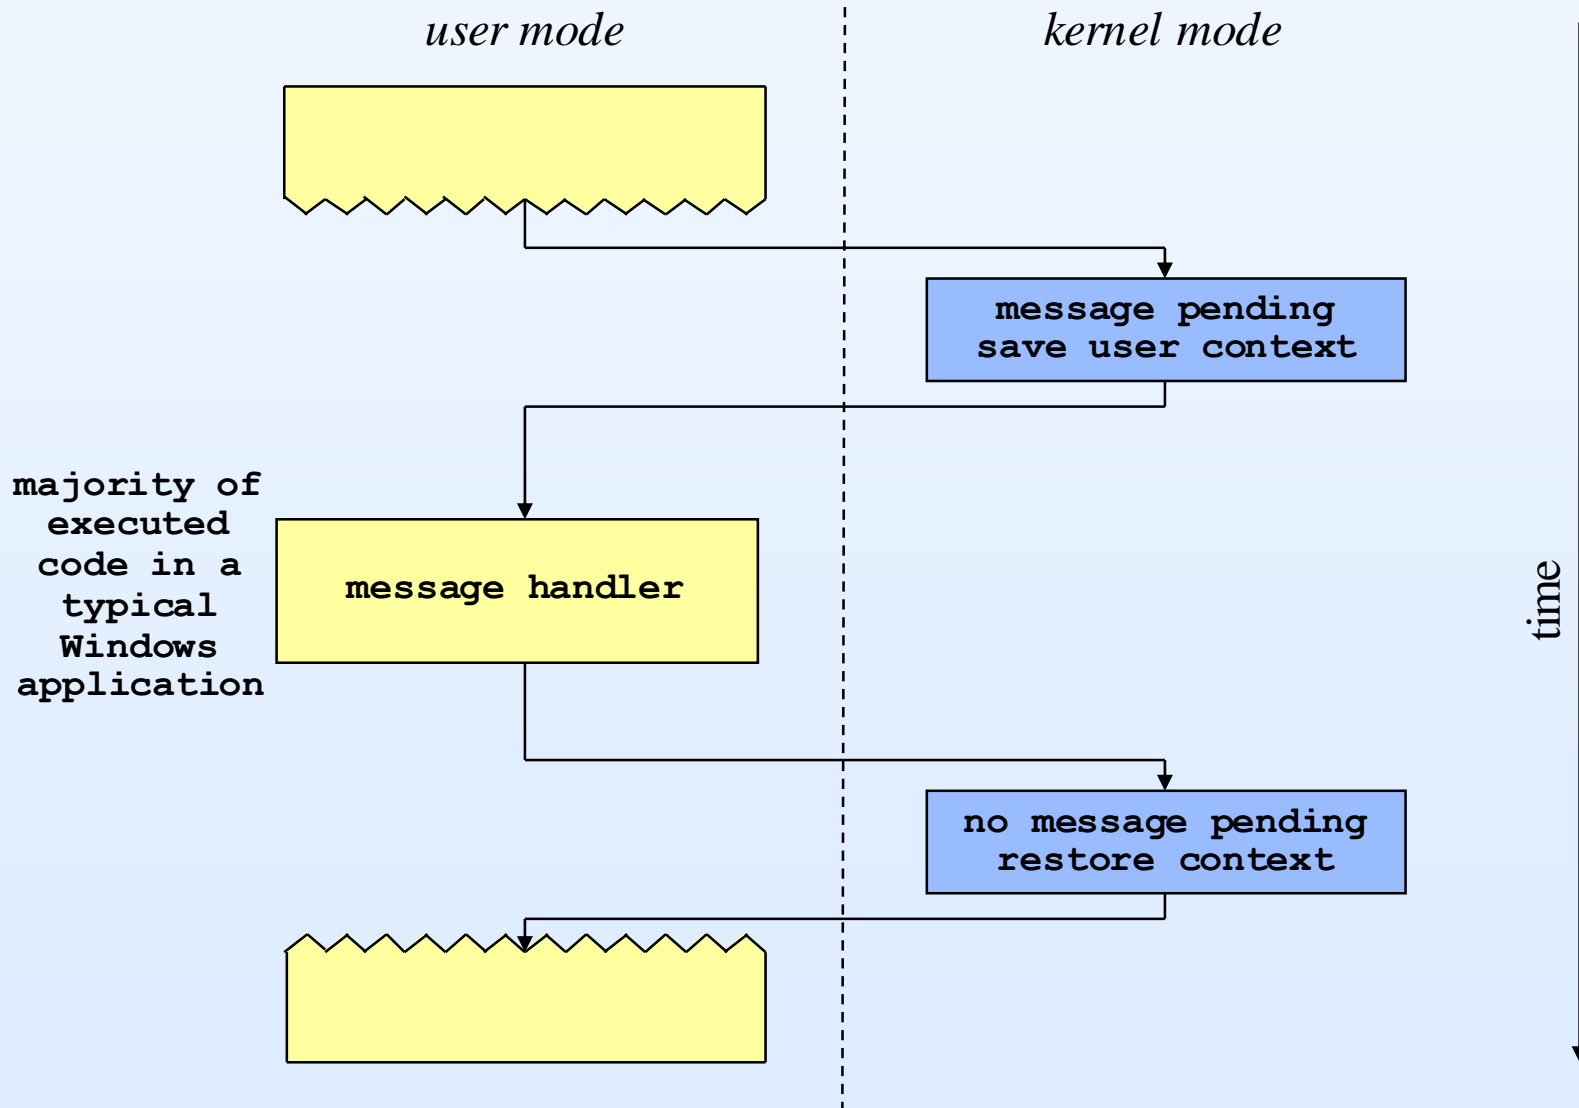
Overview Outline

- Efficient
- Transparent
- Comprehensive
- Customizable

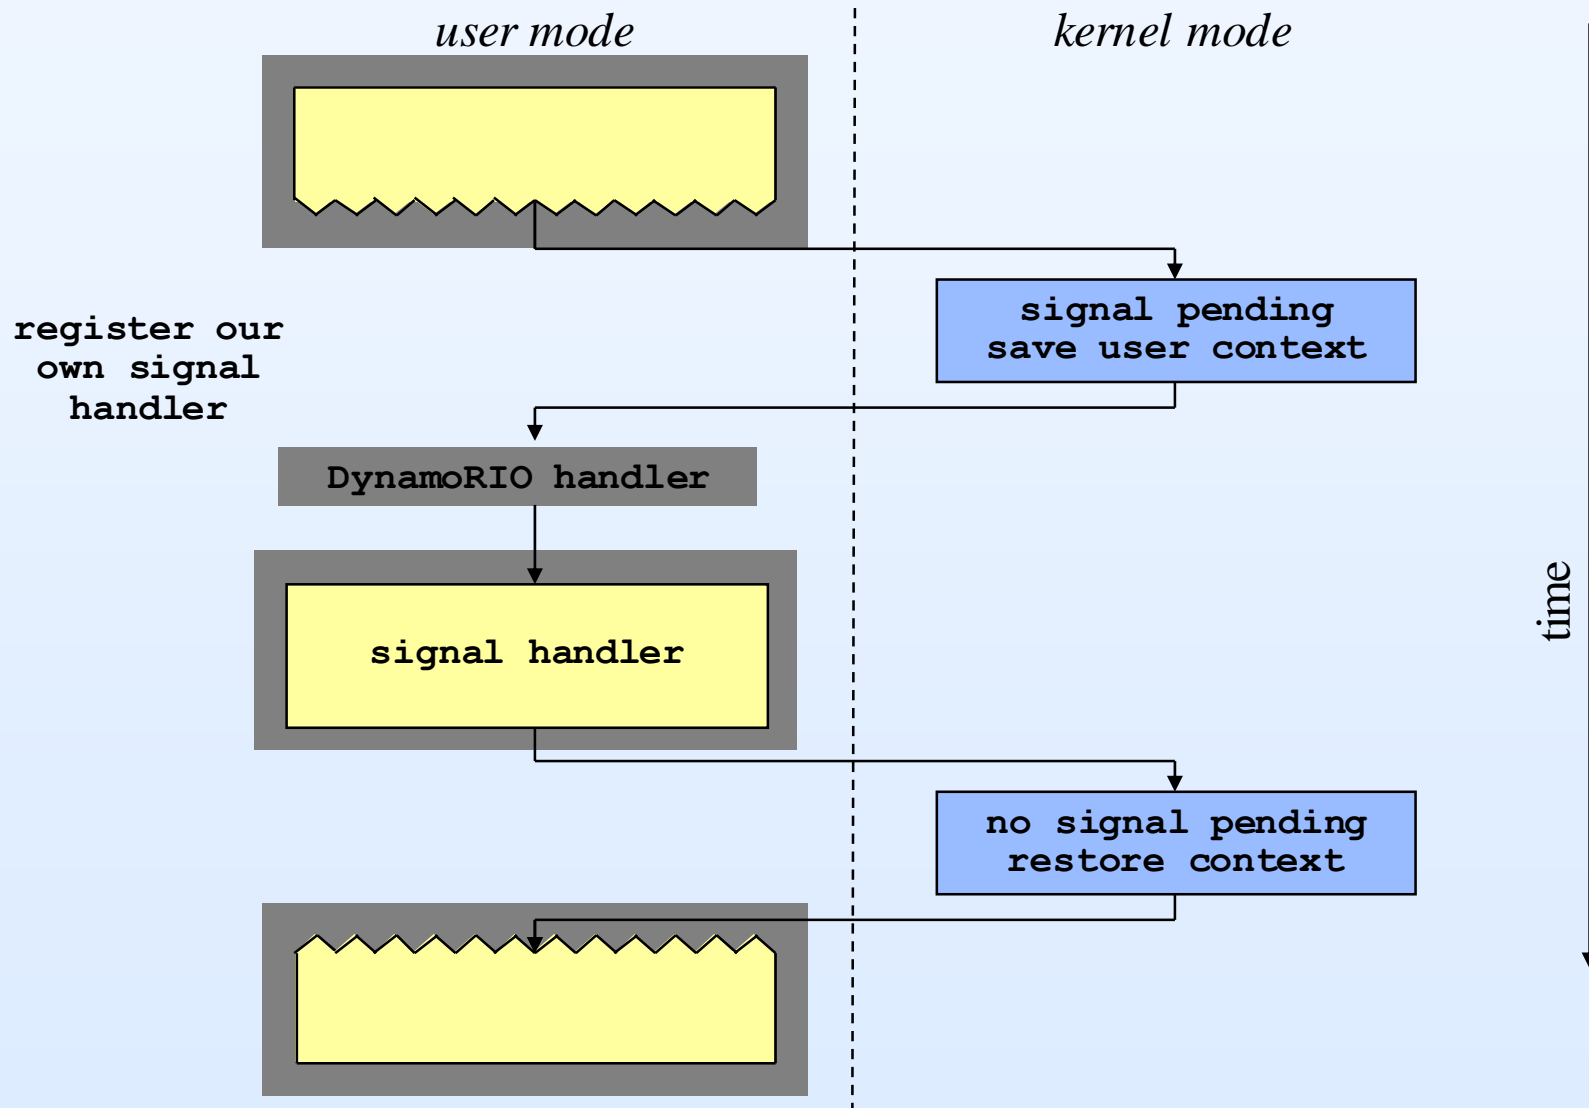
Above the Operating System



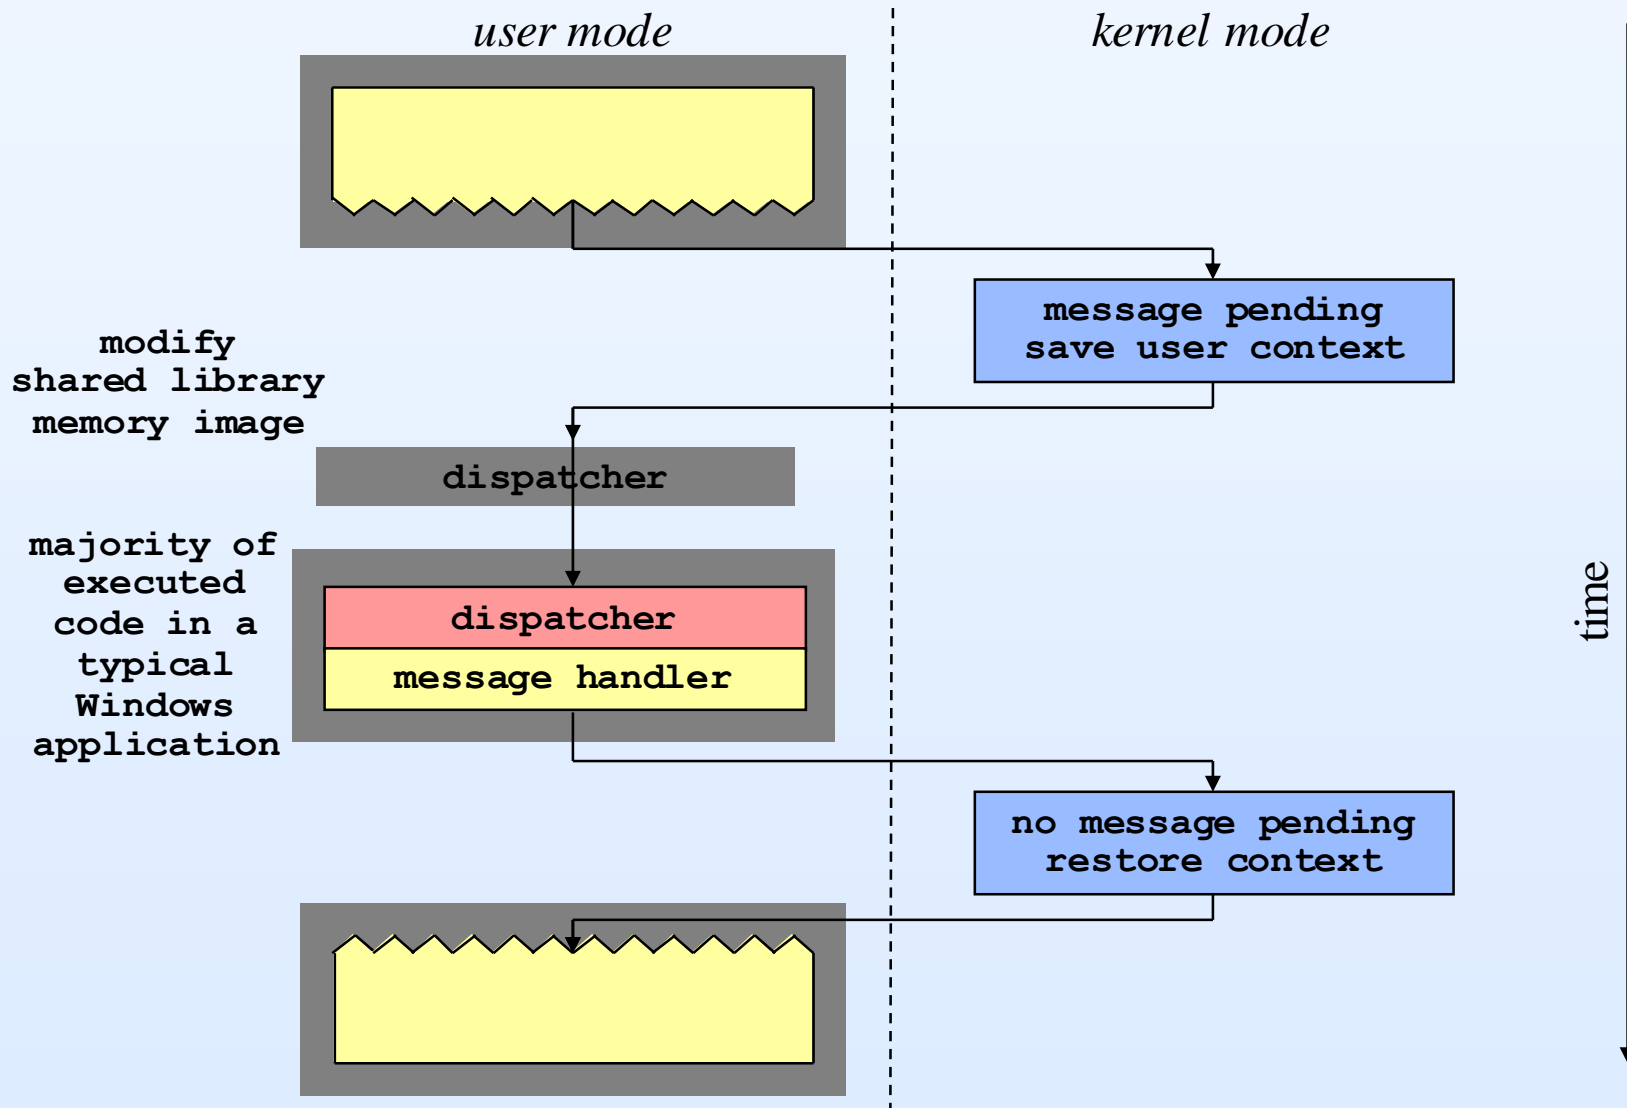
Kernel-Mediated Control Transfers



Intercepting Linux Signals



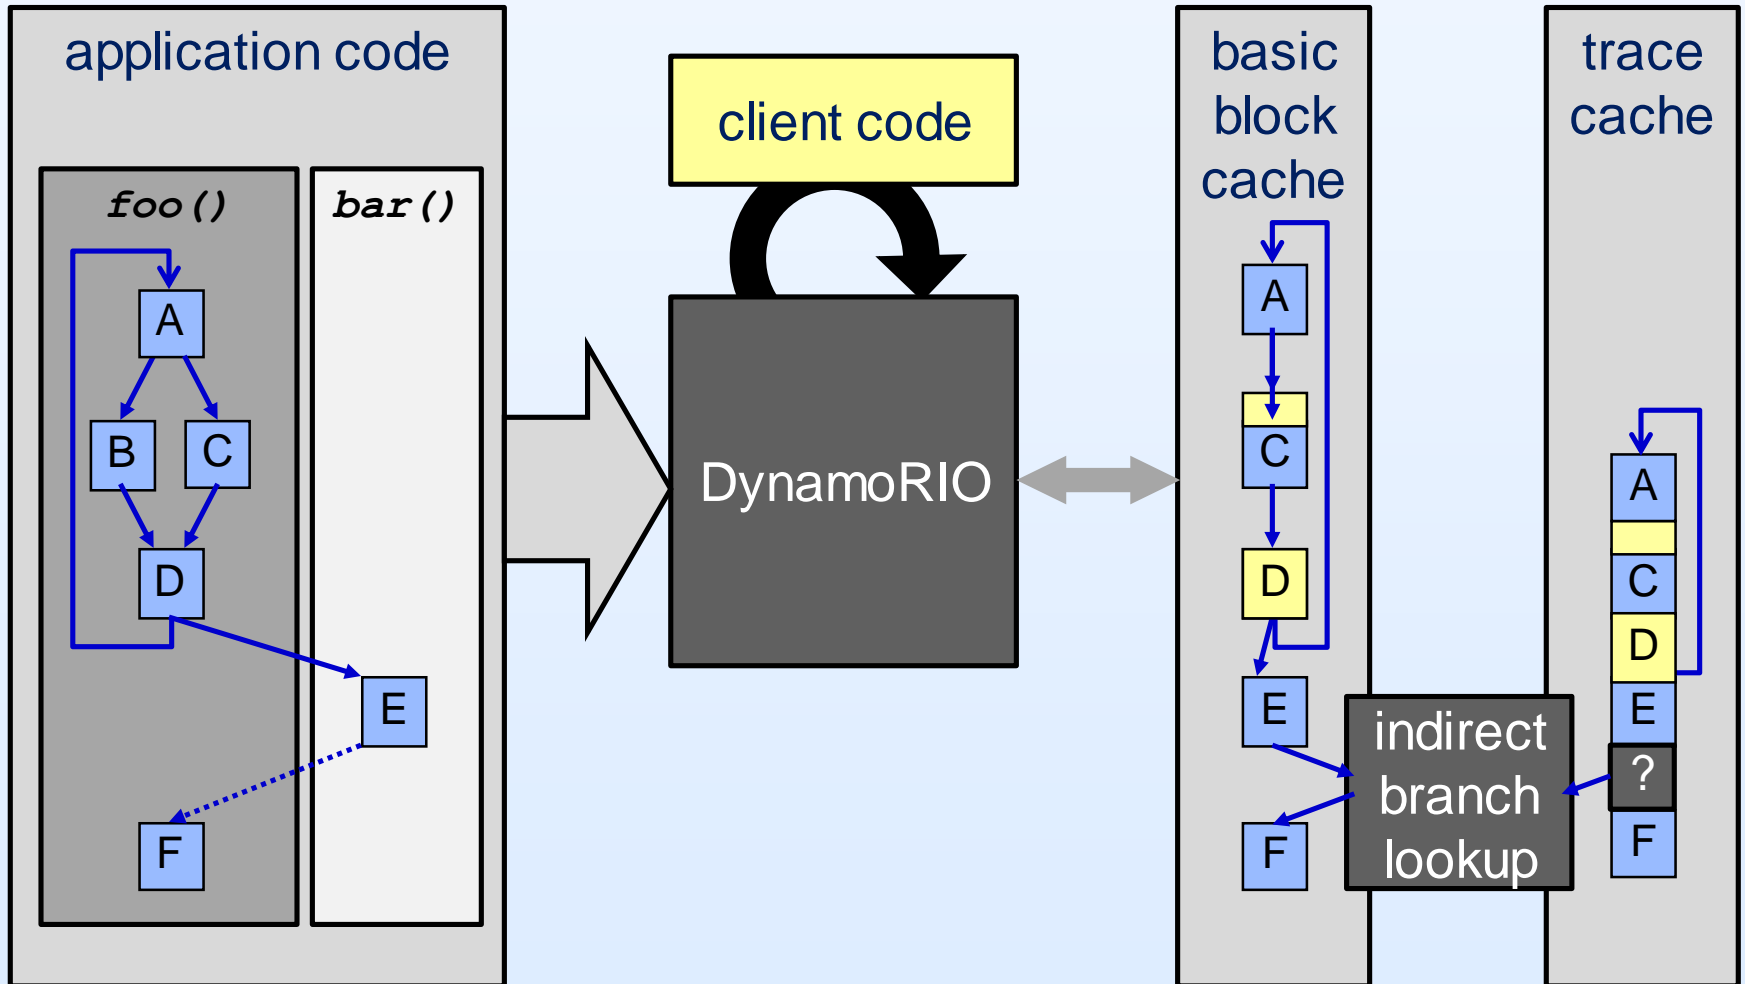
Intercepting Windows Messages



Overview Outline

- Efficient
- Transparent
- Comprehensive
- Customizable
 - Clients

DynamoRIO + Client \Rightarrow Tool



DynamoRIO

Demo

DynamoRIO Demo

The image displays two windows side-by-side. The left window is a Microsoft PowerPoint presentation titled "DynamoRIO-tutorial-feb2015.pptx". The current slide, slide 72, is titled "DynamoRIO Demo" and features the text "DynamoRIO Demo" in a large, black, sans-serif font. The right window is the DynamoRIO control interface, titled "<select an instance> - DynamoRIO". It shows a list of running instances, with "237332 R POWERPNT.EXE" selected. The status is "Running". Below the instance list, a large text area displays various performance metrics for the selected instance. At the bottom of the window, there is a status bar showing "Ready".

Performance Metrics:

Threads under DynamoRIO control	=	9
Peak threads under DynamoRIO control	=	15
Threads ever created	=	22
Callbacks	=	34679
APCs	=	0
Exceptions	=	0
System calls, pre	=	42557
System calls, post	=	7844
Native modules present	=	0
Application mmmaps	=	167
Application munmaps	=	73
Basic block fragments generated	=	507940
Trace fragments generated	=	48308
Coarse-grain units	=	0
Peak coarse-grain units	=	0
Persisted caches successfully loaded	=	0
Fcache units on live list	=	747
Peak fcache units on live list	=	747
Fcache units on free list	=	5
Peak fcache units on free list	=	5
Heap units on live list	=	243
Peak heap units on live list	=	245
Heap units on free list	=	4
Peak heap units on free list	=	5

Additional Metrics:

Instructions	=	15574688086
Floating point instrs	=	179834734
System calls	=	2559711

DynamoRIO API

8:30- 8:40	Welcome + DynamoRIO History
8:40- 9:10	Tool Demonstrations
9:10- 9:30	DynamoRIO System Overview
9:30- 9:45	DynamoRIO API Part 1
9:45-10:00	ARM Status + Demonstrations
10:00-10:30	<i>Break</i>
10:30-10:55	DynamoRIO API Part 2
10:55-11:20	How to Create Your Own Tool
11:20-11:40	Tool Internals
11:40-11:45	Q & A

Cross-Platform Approach

- A consistent interface that works across platforms
 - Windows, Linux, Android, MacOS
 - x86, AMD64, ARM, AArch64
 - Thread-private versus thread-shared
- Same client source code generally works on all combinations of platforms
- Main exception is low-level hand-tuned instruction sequences
 - Even that can be made mostly cross-platform using provided IR utilities explained later

DynamoRIO Clients

- DynamoRIO
 - “*Interpret*” the application execution
 - API: provide a cross-platform interface for clients
 - Event callbacks
 - Utilities
- Client code
 - Extend DynamoRIO with customized operations
 - Event driven: operate on different events
 - Manipulate code stream
- DynamoRIO Extensions
 - Libraries that extend DynamoRIO API
 - *drmgr, drsyms, drwrap, ...*

Events

- Application events
 - Runtime events occur during application execution
 - E.g. pre/post system calls
- DynamoRIO events
 - Code stream
 - Basic block/trace creation
 - Nudge
 - State restore
- Event callbacks registration
 - `dr_register_*_event`
 - `dr_unregister_*_event`

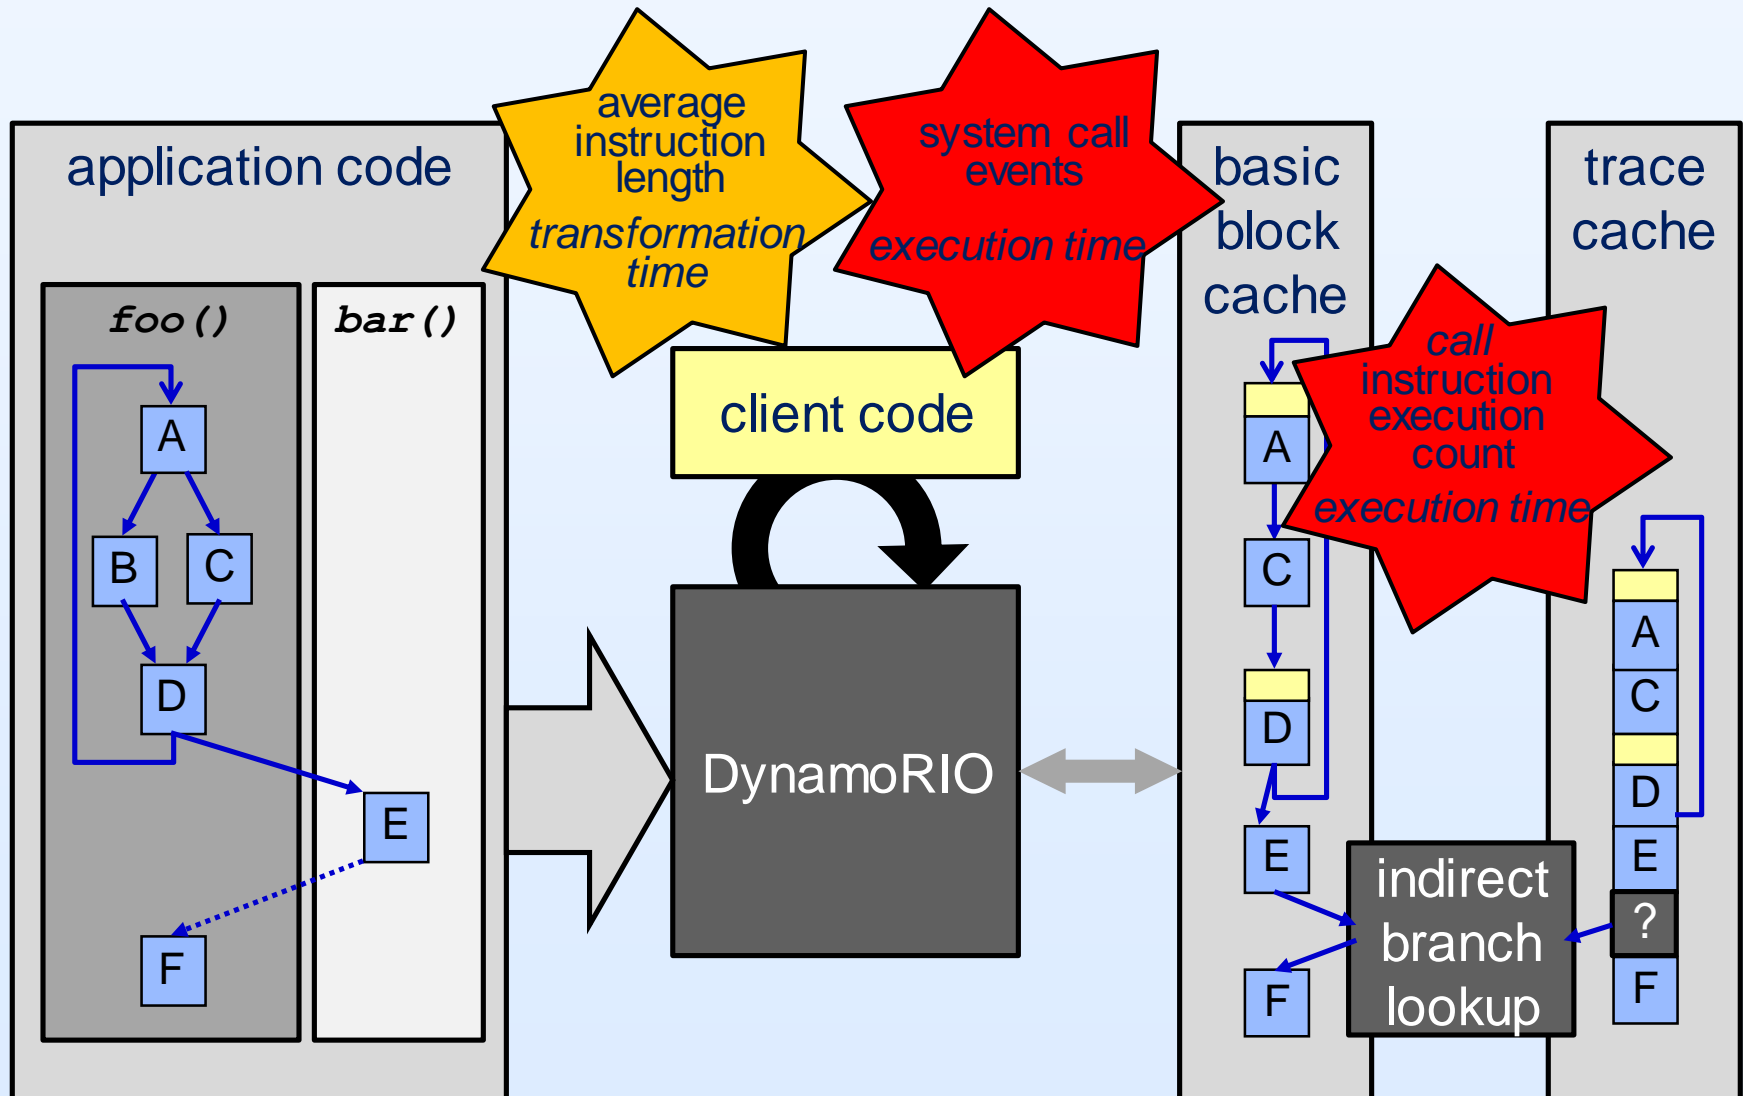
Application Events

- Process start, stop, fork
- Thread creation and deletion
- Library load and unload
- Exception (Windows)
 - Client chooses whether to deliver or suppress
- Signal (Linux)
 - Client chooses whether to deliver, suppress, bypass, or redirect
- Pre- and post- system call
 - Platform-independent system call parameter access
 - Client can modify a system call
 - Invoke additional system call as the app

DynamoRIO Events

- Code stream
 - Client has opportunity to inspect and potentially modify every single application instruction, immediately before it executes
 - Event happens at *transformation* time
 - Modifications or inserted code will operate at *execution* time
 - Entire application code stream
 - Basic block creation event: can modify the block
 - For comprehensive instrumentation tools
 - Or, focus on hot code only
 - Trace creation event: can modify the trace
 - Custom trace creation: can determine trace end condition
 - For optimization and profiling tools

Transformation Time vs Execution Time



Client Example: Dynamic Basic Block Count

```
#include "dr_api.h"
DR_EXPORT void
dr_client_main(client_id_t id, int argc, const char *argv[]) {
    dr_register_bb_event(event_bb);
}
```

```
static void bbcount() { global_count++; }
```



execution time

```
static dr_emit_flags_t
event_bb(void *drcontext, void *tag, instrlist_t *bb,
         bool for_trace, bool translating) {
    dr_insert_clean_call(drcontext, bb, instrlist_first_app(bb),
                        (void *)bbcount, false /* save fpstate */, 0);
    return DR_EMIT_DEFAULT;
}
```



transformation time

DynamRIO Extensions

- Challenges (while building Dr. Memory)
 - Multiple components
 - System call database, shadow memory management, memory allocation tracking, function wrapping, etc.
 - Manipulate the same instruction list
 - Different instrumentation passes may interfere with each other
 - Good performance
 - Avoid unnecessary save/restore
 - Instrumentation across multiple instructions (e.g., shadow xl8)
 - Code reuse
- Solution
 - DynamoRIO Extensions
 - drmgr, drwrap, drsyscall, ...

DynamoRIO Extensions

- DynamoRIO API is extended via libraries called Extensions
- Both static and shared supported
- Built and packaged with DynamoRIO
- Easy for a client to use with CMake
 - `use_DynamoRIO_extension(myclient extensionname)`

DynamoRIO Extensions

- Current Extensions:
 - *drreg*: register stealing and allocating
 - *drsyms*: symbol table and debug information lookup
 - *drcontainers*: hashtable, vector, and table
 - *drmgr*: multi-instrumentation mediation
 - *drwrap*: function wrapping and replacing
 - *drutil*: memory tracing, string loop expansion
 - *drx*: multi-process management, misc utilities
 - *drsyscall*: system call monitoring: system call names, numbers, parameter types, memory references
 - *drdecode*: standalone IA32/AMD64/ARM/Thumb/AArch64 decoding/encoding library
 - *umbra*: shadow memory framework

DynamoRIO Extensions: drmgr

- Mediation among multiple components
 - Callbacks invocation orders: priority

```
drmgr_priority_t {  
    size_t      struct_size;  
    const char  *name;      // relative order  
    const char  *before;  
    const char  *after;  
    int         priority;   // numeric order  
};
```

- Basic block instrumentation: 4 instrumentation phases
 - Application-to-application transformation
 - Application analysis
 - Instrumentation insertion
 - Instrumentation optimization

drmgr: Instrumentation Phases

- Application-To-Application
- Application Analysis
- Instrumentation Insertion
- Instrumentation Optimization

```
bool drmgr_register_bb_instrumentation_ex_event  
  (drmgr_app2app_ex_cb_t app2app_func,  
   drmgr_ilst_ex_cb_t    analysis_func,  
   drmgr_insertion_cb_t   insertion_func,  
   drmgr_ilst_ex_cb_t     instru2instru_func,  
   drmgr_priority_t        *priority)
```

drmgr: Instrumentation Phases

- Application-To-Application

```
bool drmgr_register_bb_app2app_event  
    (drmgr_xform_cb_t func, drmgr_priority_t *priority)
```

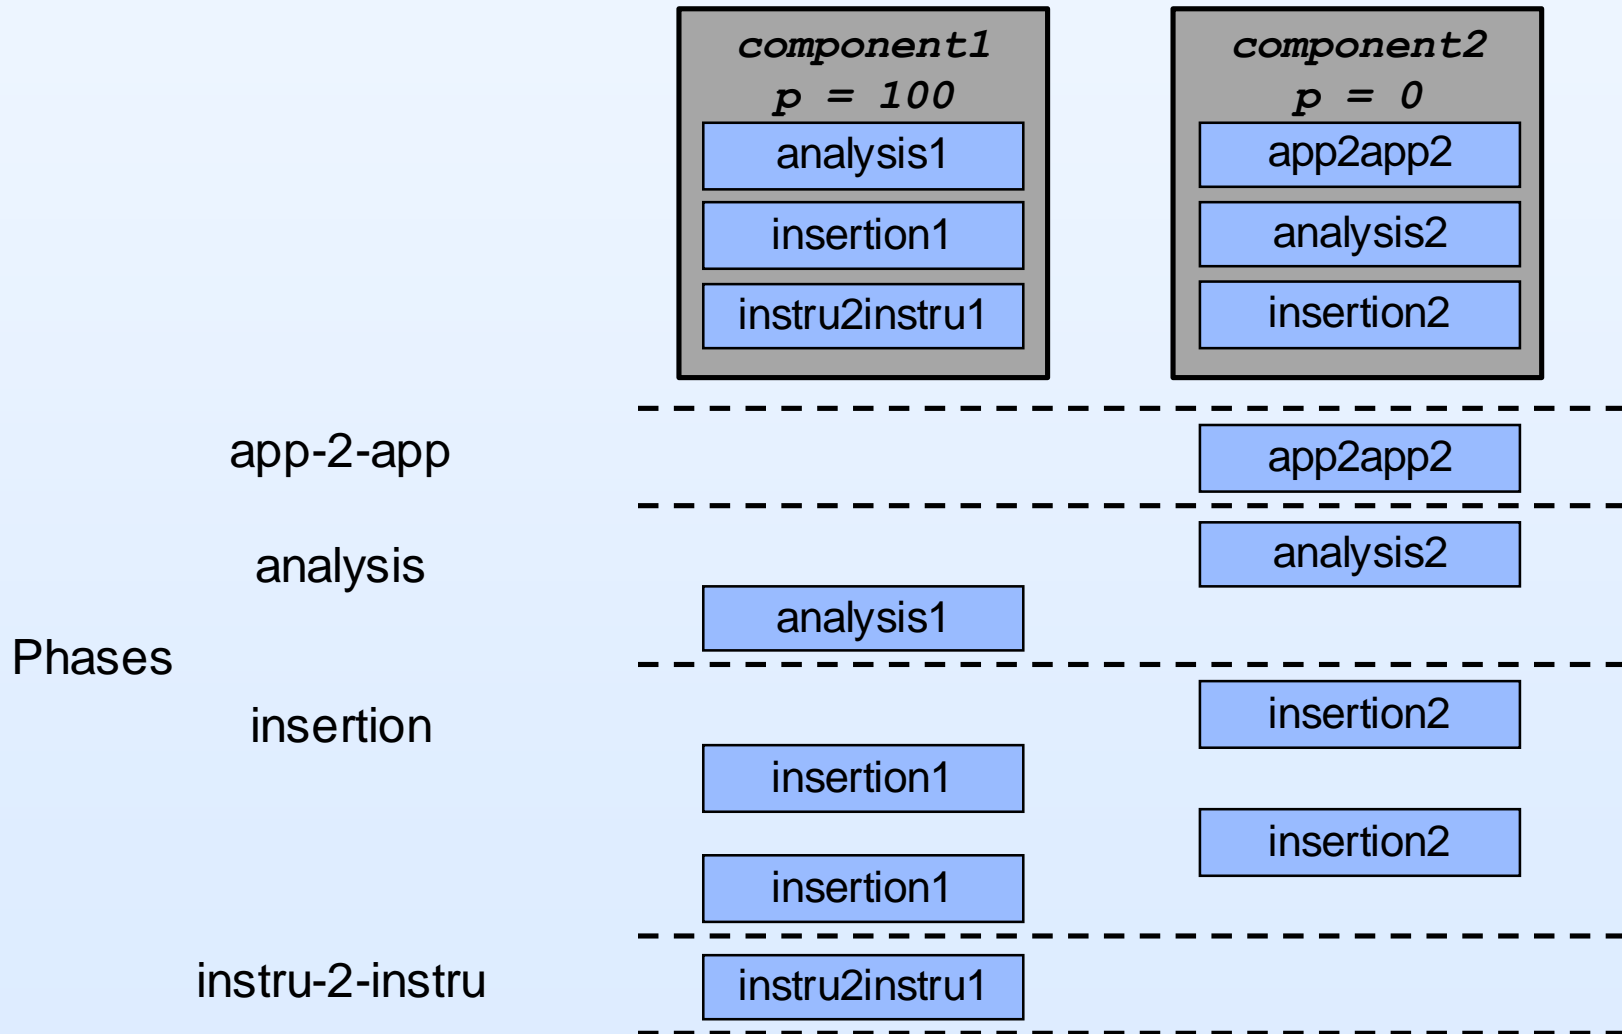
- Application Analysis
- Instrumentation Insertion

```
bool drmgr_register_bb_instrumentation_event  
    (drmgr_analysis_cb_t  analysis_func,  
     drmgr_insertion_cb_t insertion_func,  
     drmgr_priority_t     *priority);
```

- Instrumentation Optimization

```
bool drmgr_register_bb_instru2instru_event  
    (drmgr_xform_cb_t func, drmgr_priority_t *priority);
```


drmgr: Instrumentation Phases



drmgr: instruction counting

```
static dr_emit_flags_t
event_bb_analysis(void *drcontext, void *tag, instrlist_t *bb,
                  bool for_trace, bool translating, void **user_data)
{
    instr_t *instr;
    uint num_instrs;

    for (instr = instrlist_first_app(bb), num_instrs = 0;
         instr != NULL;
         instr = instr_get_next_app(instr)) {
        num_instrs++;
    }

    *user_data = (void *) (ptr_uint_t) num_instrs;

    return DR_EMIT_DEFAULT;
}
```




drmgr: instruction counting

```
static dr_emit_flags_t
event_app_instruction(void *drcontext, void *tag, instrlist_t *bb, instr_t *instr,
                     bool for_trace, bool translating, void *user_data)
{
    uint num_instrs;

    if (!drmgr_is_first_instr(drcontext, instr))
        return DR_EMIT_DEFAULT;

    num_instrs = (uint)(ptr_uint_t) user_data;
    dr_insert_clean_call(drcontext, bb, instrlist_first_app(bb),
                        (void *)inscount, false /* save fpstate */, 1,
                        OPND_CREATE_INT32(num_instrs));

    return DR_EMIT_DEFAULT;
}
```



ARM

8:30- 8:40 Welcome + DynamoRIO History

8:40- 9:10 Tool Demonstrations

9:10- 9:30 DynamoRIO System Overview

9:30- 9:45 DynamoRIO API Part 1

9:45-10:00 ARM Status + Demonstrations

10:00-10:30 Break

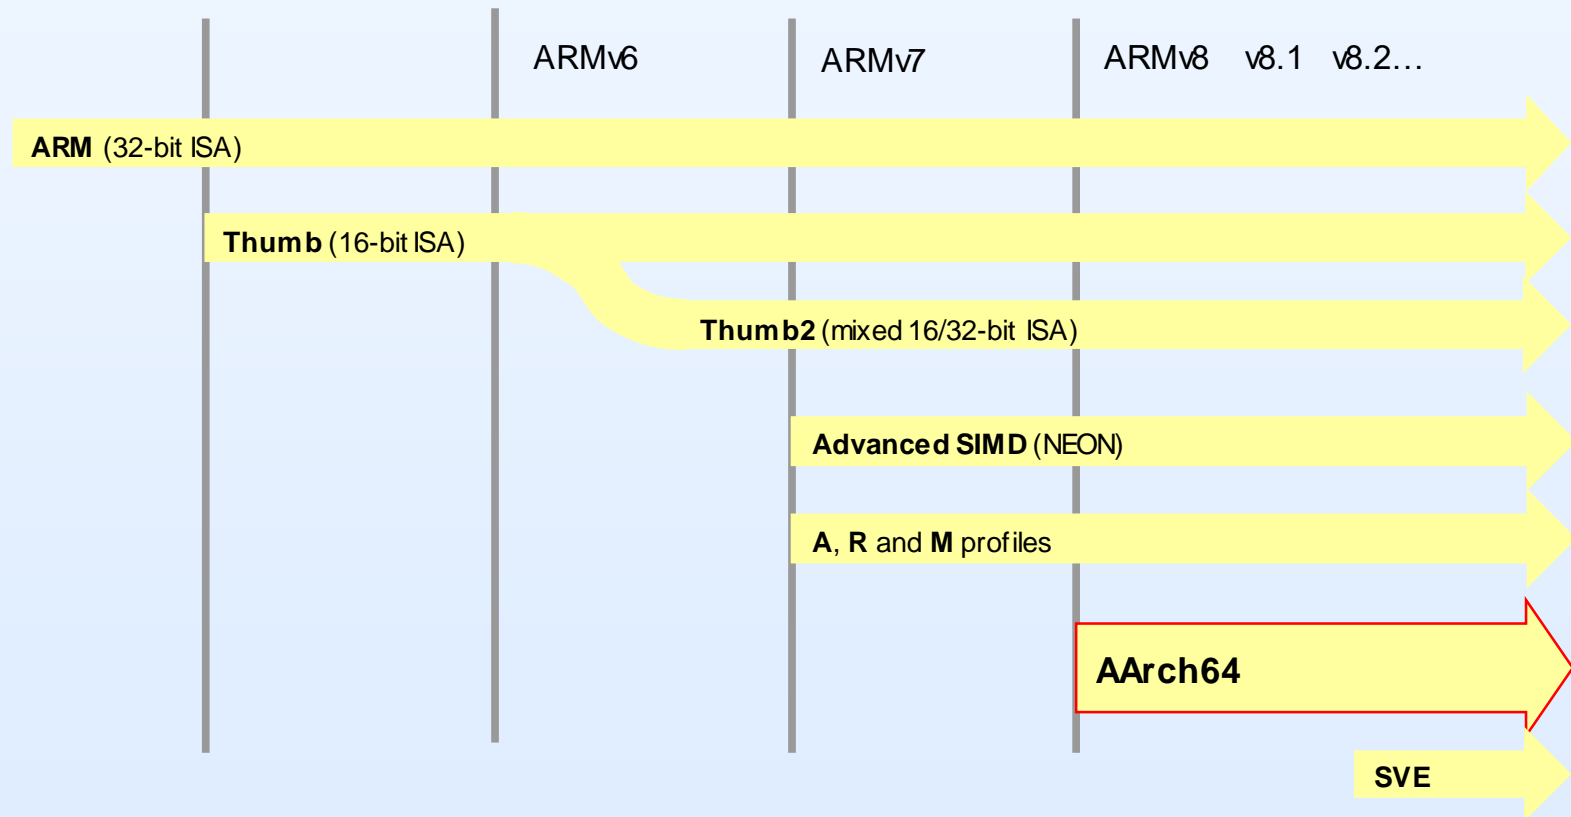
10:30-10:55 DynamoRIO API Part 2

10:55-11:20 How to Create Your Own Tool

11:20-11:40 Tool Internals

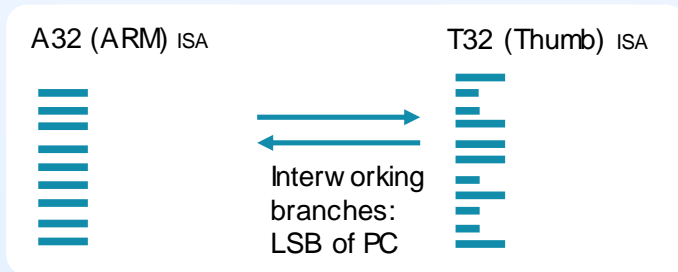
11:40-11:45 Q & A

ARM background: architecture timeline



ARM instruction sets and execution states (as of ARMv8)


AArch32 execution state



GPR registers (512 bits / 64 bytes total):
16 32-bit registers R0 to R15
R15 is PC – means lots of interesting branches!
R13 is SP
R14 is link register (but may be used as work reg)
Conventions for frame pointers etc. – not architected

VFP/NEON registers (2048 bits / 256 bytes total)
32, 64 and 128 bit registers overlay the register file

AArch64 execution state



GPR registers (1984 bits / 248 bytes total):
31 64-bit registers R0 to R30
Full register accessed as X0 to X30
Low words accessed as W0 to W30
'31' in a register operand indicates zero register (XZR or WZR)
Dedicated PC – written by branch instructions only
Dedicated SP

VFP/NEON registers (4096 bits / 512 bytes total)
V0-V31 can be each accessed as 8, 16, 32, 64 or 128 bits

ARM/AArch64 unimplemented features

- Clean call optimisation
- Trace construction
- Thread-private fragment cache
- Some samples: bbcount, cbr, countcalls,...
- AArch64 Android
- AArch64 Dr. Memory
- AArch64 encoder/decoder: FP/SIMD non-memory-access
- Load/store exclusive infinite loop on some hardware (i#1698)
- AArch64 handling of flags


Status of ARM/AArch64

- Number of tests run by "make test" on Linux
 - ARM 131
 - AArch64 139
 - x86_64 251
- AArch64 is not yet as widely tested
 - Firefox runs under DynamoRIO on some OS versions
 - Firefox crashes under DynamoRIO on other OS versions

Writing clients for ARM/AArch64

- Safe to use all client APIs which are covered in the test suites
- Code manipulation API restrictions (read ARMv8)
 - INSTR_CREATE macro covers common opcodes
 - Not yet completed for AArch64
 - Limited offset ranges
 - Direct or conditional branch, pc-relative access, shift offsets, etc.
 - Additional scratch register is required in some cases
 - Saving SP value to memory, etc.
 - Cannot store an immediate directly to memory
 - `INSTR_CREATE_str(dc, OPND_CREATE_INT(imm))` **wrong!**
 - AArch64: SP/ZR is not a general purpose register!
 - `dr_save_reg(dc, ..., opnd_create_reg(DR_REG_XSP))` **wrong!**

Preliminary Performance Comparison



AArch64 Demonstrations

(swap to live demo)

DynamoRIO API

8:30- 8:40 Welcome + DynamoRIO History

8:40- 9:10 Tool Demonstrations

9:10- 9:30 DynamoRIO System Overview

9:30- 9:45 DynamoRIO API Part 1

9:45-10:00 ARM Status + Demonstrations

10:00-10:30 Break

10:30-10:55 DynamoRIO API Part 2

10:55-11:20 How to Create Your Own Tool

11:20-11:40 Tool Internals

11:40-11:45 Q & A

DynamoRIO API: Instruction Representation

- Full IA-32/AMD64/AArch32/AArch64 instruction representation
- Instruction creation with auto-implicit-operands
- Operand iteration
- Instruction lists with iteration, insertion, removal
- Decoding at various levels of detail
- Encoding

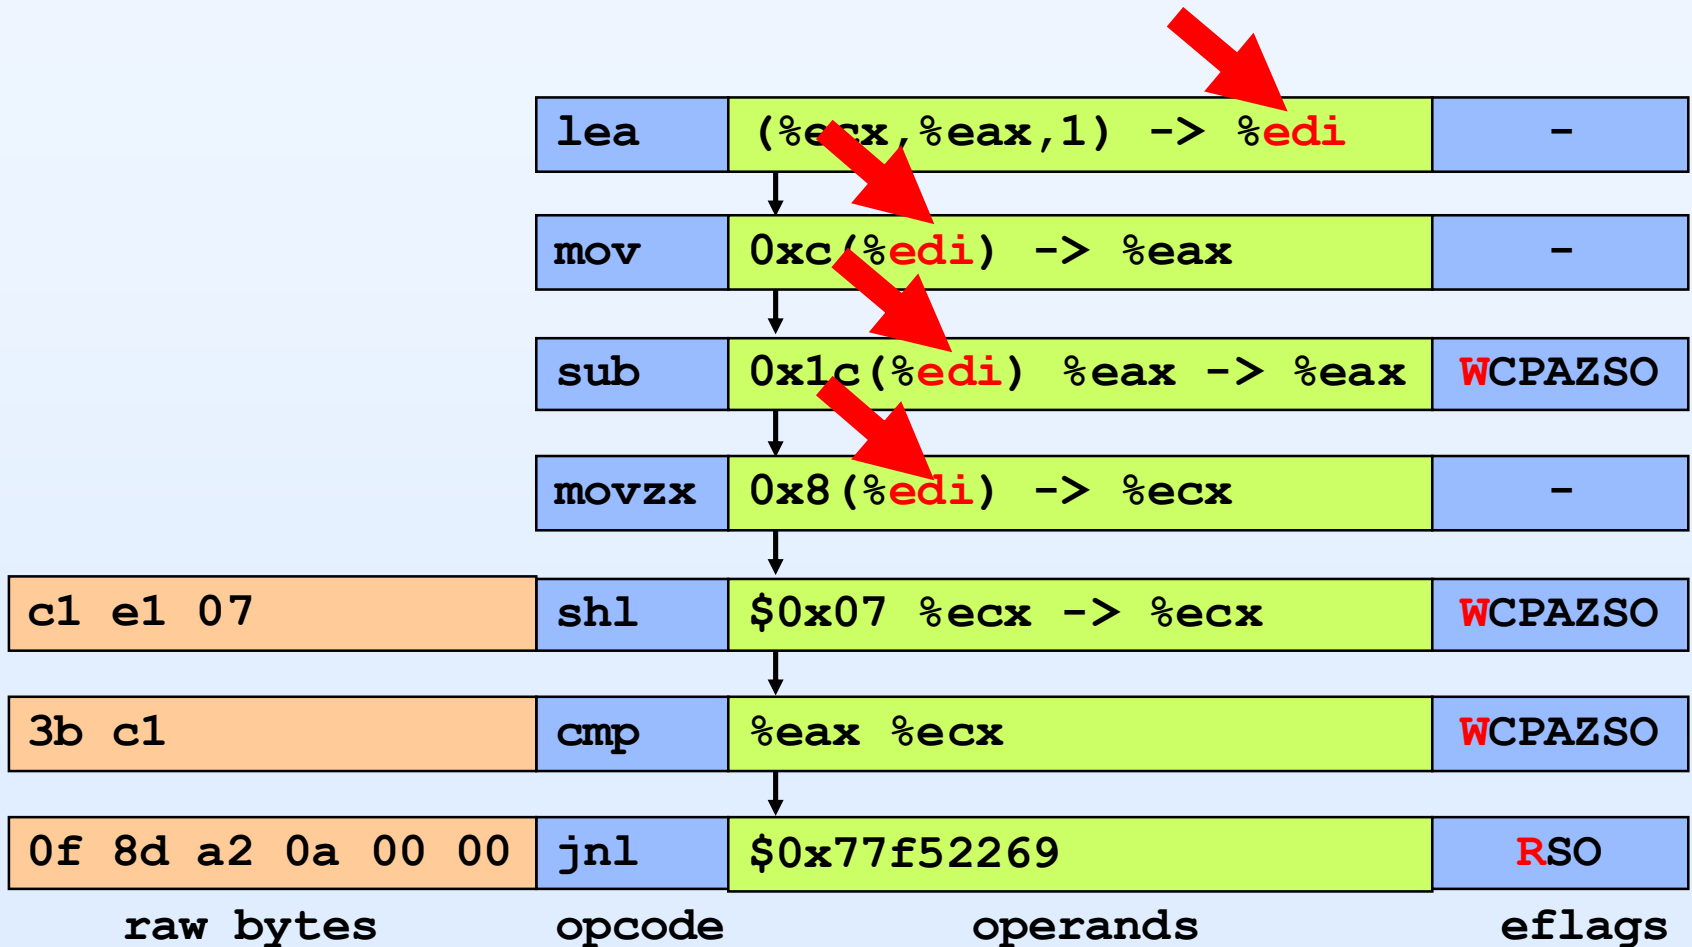
Reminder: No Intermediate Layers

- IR mirrors underlying ISA
 - Preserves optimized application code
 - Intermediate layers incur significant performance impact
 - QEMU (user-mode) **6x** slower than DR, Valgrind **4x** slower than DR
 - *This is the key to good performance*
- Still have an abstraction layer
 - Block or trace = list of instructions
 - Instruction = lists of source and destination operands
 - Tool code often still cross-platform
 - “Does this instruction read memory?”

Instruction Representation

8d 34 01	lea	(%ecx,%eax,1) -> %esi	-
8b 46 0c	mov	0xc(%esi) -> %eax	-
2b 46 1c	sub	0x1c(%esi) %eax -> %eax	WCPAZSO
0f b7 4e 08	movzx	0x8(%esi) -> %ecx	-
c1 e1 07	shl	\$0x07 %ecx -> %ecx	WCPAZSO
3b c1	cmp	%eax %ecx	WCPAZSO
0f 8d a2 0a 00 00	jnl	\$0x77f52269	RSO
raw bytes	opcode	operands	eflags

Instruction Representation



Instruction Creation

- Method 1: use the INSTR_CREATE_opcode macros that fill in implicit operands automatically:

```
instr t *instr = INSTR_CREATE dec(dcontext,  
    opnd create reg(DR_REG_EDX));
```

- Method 2: specify opcode + all operands (including implicit operands):

```
instr t *instr = instr create(dcontext);  
instr set opcode(instr, OP dec);  
instr set num opnds(dcontext, instr, 1, 1);  
instr set dst(instr, 0, opnd create reg(DR_REG_EDX));  
instr set src(instr, 0, opnd create reg(DR_REG_EDX));
```

Cross-Platform IR Support

- Instruction creation: `XINST_CREATE_*` macros
 - `XINST_CREATE_load()`, `XINST_CREATE_jump()`, etc.
- Generic instruction and operand queries
 - E.g.: `instr_writes_memory()`, `instr_reads_from_reg()`, `instr_compute_address()`, `instr_is_return()`
 - E.g.: `opnd_is_memory_reference()`, `opnd_uses_reg()`
- Generic instrumentation creation helpers
 - E.g.: `instrlist_insert_mov_immed_ptrsz()`, `instrlist_insert_mov_instr_addr()`
- ISA-specific concepts applied to all ISA's
 - E.g., predication
- *drreg* Extension for using registers without naming them

ISA Modes


- IR covers all modes
 - Abstracts away underlying mode
- ARM and Thumb
 - ARM build fully supports both modes
- AMD64 and IA-32
 - 32-bit x86 build of DynamoRIO only handles 32-bit code
 - 64-bit AMD64 build of DynamoRIO decodes/encodes both 32-bit and 64-bit code
 - Current release does support executing applications that mix the two modes on Windows
 - “X” cross-mode registers
 - DR_REG_XAX is DR_REG_EAX when compiled 32-bit, and DR_REG_RAX when compiled 64-bit

ISA Modes: Per-Thread and Per-Instruction

- When going to or from the IR, the thread mode and instruction mode determine how instructions are interpreted
- When decoding, current thread's mode is used
 - Can be changed with `dr_set_isa_mode()`
- When encoding, that instruction's mode is used
 - When created, set to mode of current thread
 - Can be changed with `instr_set_isa_mode()`

Linear Control Flow

- Both basic blocks and traces are linear
- Instruction sequences are all single-entrance, multiple-exit
- Greatly simplifies analysis algorithms



DynamoRIO API: Code Manipulation

- Processor information
- State preservation
 - Eflags, arith flags, floating-point state, MMX/SSE state
 - Spill slots, TLS, CLS
 - *drreg* Extension
- Clean calls to C code
- Dynamic instrumentation
 - Replace code in the code cache
- Branch instrumentation
 - Convenience routines

State Preservation

- Application context is the native context in the code cache: it must be explicitly preserved
 - Another key design choice for performance
- Client must spill registers before using them
 - Must also save and restore condition codes
 - *drreg* Extension makes this easy
- In C/C++ code, client must preserve floating-point state
 - Not saved by DR context switches (unless request on clean call)
 - `dr_insert_save_fpstate()`, `dr_insert_restore_fpstate()`

drreg Extension: Register Preservation

- *drreg* mediates general-purpose scratch register and condition code spilling and restoring
- Reservation model: request exclusive access
- Lazily restores to avoid superfluous adjacent spills
- Most library routines take in a scratch register
- Example:

```
/* acquire a register, spilling it if necessary */  
drreg_reserve_register(drcontext, bb, inst, NULL, &reg);  
/* pass register to library routine */  
drx_buf_insert_load_buf_ptr(drcontext, buf, bb, inst, reg);  
/* release the register, restoring from memory if necessary */  
drreg_unreserve_register(drcontext, bb, inst, reg);
```


Where to Spill?

- Absolute addressing
 - Thread-private code caches only
- Application stack
 - Not reliable or transparent
- Stolen register
 - Used by DynamoRIO on ARM and AArch64
- Segment
 - Used by DynamoRIO on AMD64 and IA-32

Thread-Local Storage

- Spill slots for registers
 - DR provides several fast slots and slower slots
 - *drreg* provides a configurable number of fast slots
- General thread-local storage
 - `drmgr_register_tls_field()`
 - `drmgr_insert_read_tls_field()`, `drmgr_insert_write_tls_field()`
 - From C code: `drmgr_get_tls_field()`, `drmgr_set_tls_field()`
 - Parallel routines for CLS fields (advanced Windows topic)

Clean Calls

```
if (instr_is_mbr(instr)) {
    app_pc address = instr_get_app_pc(instr);
    uint opcode = instr_get_opcode(instr);
    dr_insert_clean_call(drcontext, ilist, instr, (void *) at_mbr,
                        false/*don't need to save fp state*/,
                        2 /* 2 parameters */,
                        /* opcode is 1st parameter */
                        OPND_CREATE_INT32(opcode),
                        /* address is 2nd parameter */
                        OPND_CREATE_INTPTR(address));
}
```

- Saved interrupted application state can be accessed using `dr_get_mcontext()` and modified using `dr_set_mcontext()`

Clean Call Inlining

- Simple clean callees will be automatically optimized and potentially inlined
- `-opt_cleancall` runtime option controls aggressiveness
- Current requirements for inlining:
 - Leaf routine (may call PIC get-pc thunk)
 - Zero or one argument
 - Relatively short
- Compile the client with optimizations to improve clean call optimization
- Look in debug logfile for “CLEANCALL” to see results

Dynamic Instrumentation

- Thread-shared: flush all code corresponding to application address and then re-instrument when re-executed
 - Can flush from clean call, and use `dr_redirect_execution()` since cannot return to potentially flushed cache fragment
- Thread-private: can also replace particular fragment (does not affect other potential copies of the source app code)
 - `dr_replace_fragment()`

Flushing the Cache

- Immediately deleting or replacing individual code cache fragments is available for thread-private caches
 - Only removes from that thread's cache
- Two basic types of thread-shared flush:
 - Non-precise: remove all entry points but let target cache code be invalidated and freed lazily
 - Precise/synchronous:
 - Suspend the world
 - Relocate threads inside the target cache code
 - Invalidate and free the target code immediately

Flavors of Cache Flushing

- Thread-shared flush API routines:
 - `dr_unlink_flush_region()`: non-precise flush
 - `dr_flush_region()`: synchronous flush
 - `dr_delay_flush_region()`:
 - No action until a thread exits code cache on its own
 - If provide a completion callback, synchronous once triggered
 - Without a callback, non-precise

Memory Tracing

- *drutil* Extension provides utilities for memory address tracing:
 - Address acquisition
 - Converts segment-based addresses into linear addresses
 - String loop expansion
 - Turns x86 string loops into regular loops with explicit control flow for simpler instrumentation

DynamoRIO API: General Utilities

- DynamoRIO provides safe utilities for transparency support
 - Separate stack
 - Separate memory allocation
 - Separate file I/O
- Utility options
 - Use DynamoRIO-provided utilities directly
 - Use shared libraries via DynamoRIO private loader
 - Malloc, etc. redirected to DynamoRIO-provided utilities
 - Use static libraries with dependencies redirected
- Risky for client to directly invoke system calls

Third-Party Libraries

- Private loader inside DynamoRIO will load any external shared libraries a client imports from
 - Loads a duplicate copy of each library and tries to isolate from the application's copy
- On Windows, private loader does not support locating SxS libraries, so use static libc with VS2005 or VS2008

Private Libraries

- Private loader on Windows
 - Not easy to fully isolate system data structures
 - PEB and key TEB fields are isolated
 - Some libraries like ntdll.dll are shared
 - To examine application state while in client code, use `dr_switch_to_app_state()`
- Private loader on Linux
 - Isolation is simpler and more complete

DynamoRIO Heap

- Three flavors:
 - Thread-private: no synchronization; thread lifetime
 - Global: synchronized, process lifetime
 - malloc(), HeapAlloc(), etc. redirect here
 - “Non-heap”: for generated code, etc.
 - No header on allocated memory: low overhead but must pass size on free
- Leak checking
 - Debug build complains at exit if memory was not deallocated

Thread and Synchronization Support

- Thread support
 - Thread-local storage
 - Callback-local storage
 - Simple mutexes
 - Read-write locks
 - Thread-private code caches, if requested
- Sideline support
 - Create new client-only thread
 - Thread-private itimer (Linux-only)
- Suspend and resume all other threads
 - “Stop-the-world”

DynamoRIO API: General Utilities, Cont'd

- Communication
 - *Nudges*: ping from external process
 - File creation, reading, and writing
 - File descriptor isolation on Linux
- Safe read/write
 - Fault-proof read/write routines
 - Try/except facility

DynamoRIO API: General Utilities, Cont'd

- Application inspection
 - Address space querying
 - Module iterator
 - Processor identification
 - Cache line size, ISA features, etc.
 - Symbol lookup
 - Function replacing and wrapping

drsyms Extension: Symbol Table Access

- The *drsyms* Extension provides access to symbol tables and debug information
- Currently supports the following:
 - Windows PDB
 - Linux ELF + DWARF2
 - Windows PE/COFF + DWARF2
 - OSX MachO + DWARF2
- API includes:
 - Address to symbol and line information
 - Symbol to address
 - Symbol enumeration and searching
 - Symbol demangling
 - Symbol types

drwrap Extension: Func Replacing/Wrapping

- *drwrap* Extension provides function replacing and wrapping
- Use `dr_get_proc_address()` to find library exports or `drsyms` Extension to find internal functions
- Function replacing replaces with *application code*
- Function wrapping calls pre and post callbacks that execute as client code around the target application function
- Arguments, return value, and whether the function is executed can all be examined and controlled

Create Your Own Tool

8:30- 8:40	Welcome + DynamoRIO History
8:40- 9:10	Tool Demonstrations
9:10- 9:30	DynamoRIO System Overview
9:30- 9:45	DynamoRIO API Part 1
9:45-10:00	ARM Status + Demonstrations
10:00-10:30	<i>Break</i>
10:30-10:55	DynamoRIO API Part 2
10:55-11:20	How to Create Your Own Tool
11:20-11:40	Tool Internals
11:40-11:45	Q & A

How to Create Your Own Tool

- Event driven
 - Application, DynamoRIO, client, and Extensions
 - Events
- Write your DynamoRIO client
 - Part 1: Register event callbacks
 - Part 2: Implement event callbacks
 - Part 3: Instrumentation
- Config, build, and run
 - CMake

How to Create Your Own Tool

- Event driven
 - Application, DynamoRIO, client, and Extensions
 - Events
- Write your DynamoRIO client
 - Part 1: Register event callbacks
 - Part 2: Implement event callbacks
 - Part 3: Instrumentation
- Config, build, and run
 - CMake

Event Driven

- Application, DynamoRIO, clients, and Extensions
 - DynamoRIO “*interprets*” the application execution
 - A client takes actions on interested events
 - Client interacts with DynamoRIO and app via API + Extensions
 - Event callbacks
 - Utility functions
- Events
 - Application events
 - E.g., pre/post system calls
 - DynamoRIO events
 - E.g., basic block creation
 - Multiple components mediated by drmgr

Write Your DynamoRIO Client

- Part 1: Register event callbacks
 - dr_client_main
 - drmgr_register_*_event()
 - App events: pre_syscall, ...
 - DR events: bb, ...
- Part 2: Implement event callbacks
 - Strace
- Part 3: Instrumentation
 - Dynamic instruction counting

Part 1: Register Event Callbacks - strace

```
#include "dr_api.h"
#include "drmgr.h"

DR_EXPORT void
dr_client_main(client_id_t id, int argc, const char *argv[]) {
    drmgr_init();
    dr_register_exit_event(event_exit);
    drmgr_register_post_syscall_event(event_post_syscall);
}

static void event_exit() {
    drmgr_exit();
}
```

Write Your DynamoRIO Client

- Part 1: Register event callbacks

- dr_client_main
- drmgr_register_*_event()
 - App events: pre_syscall, ...
 - DR events: bb, ...

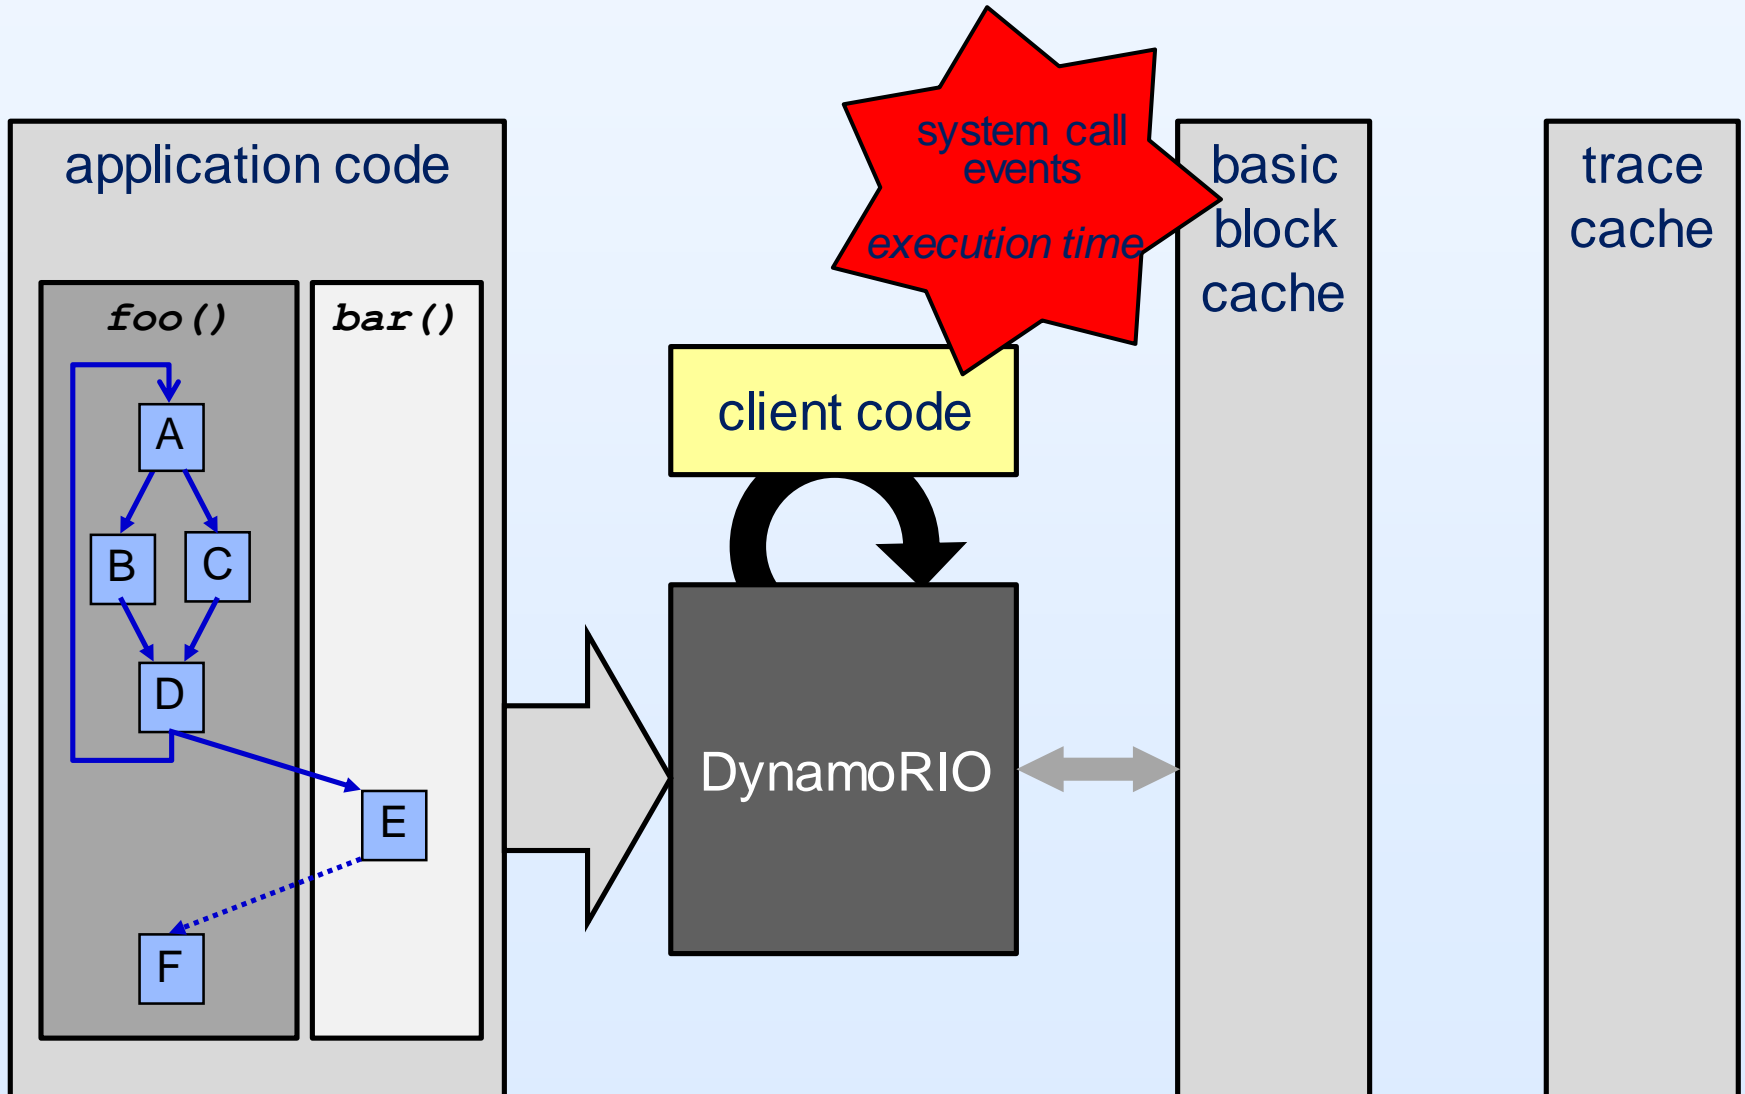
- Part 2: Implement event callbacks

- Strace

- Part 3: Instrumentation

- Dynamic instruction counting


Part 2: Implement Event Callbacks - strace



Part 2: Implement Event Callbacks - strace

```
DR_EXPORT void
dr_client_main(client_id_t id, int argc, const char *argv[]) {
    ...
    drmgr_register_post_syscall_event(event_post_syscall);
}
```


```
static void event_post_syscall(void *drcontext, int sysnum) {
    reg_t result = dr_syscall_get_result(drcontext);
    dr_printf("syscall %d: %d\n", sysnum, result);
}
```



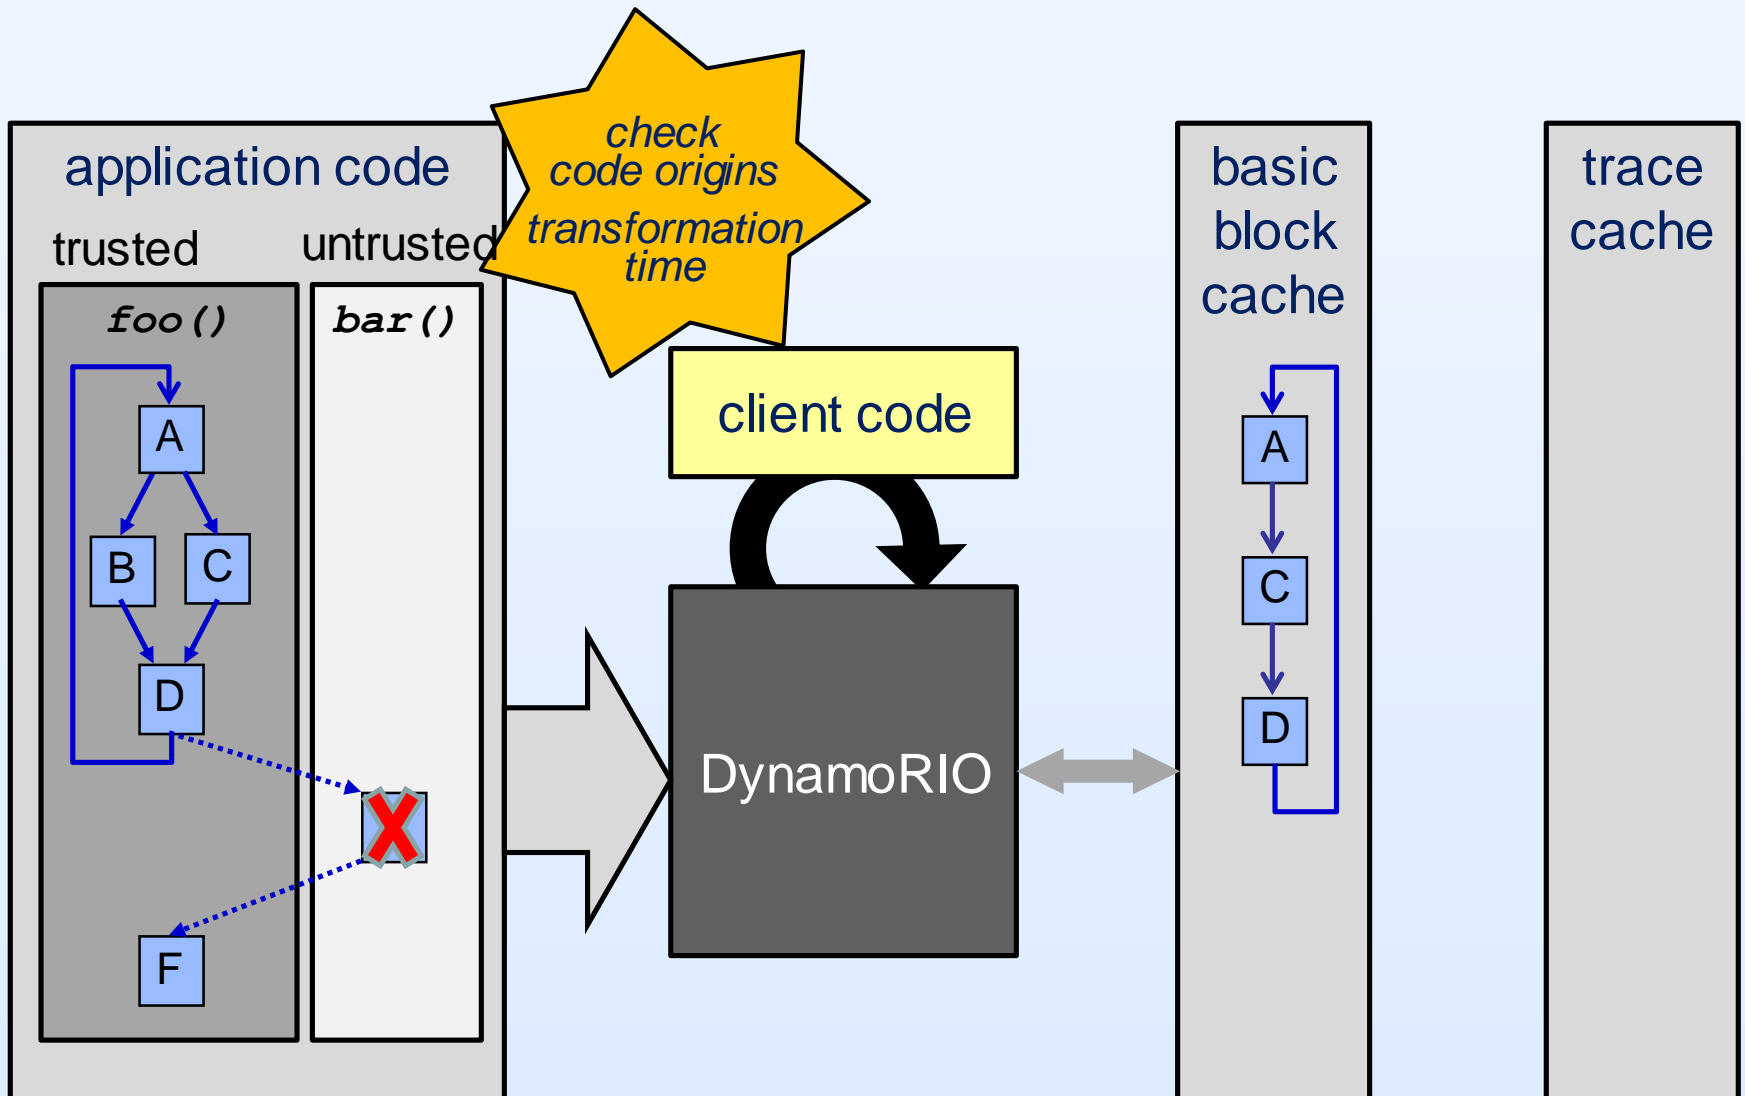
Part 2: Implement Event Callbacks - strace

```
DR_EXPORT void
dr_client_main(client_id_t id, int argc, const char *argv[]) {
    ...
    drmgr_register_pre_syscall_event(event_pre_syscall);
}

static bool event_pre_syscall(void *drcontext, int sysnum) {
    if (sysnum == SYS_write &&
        dr_syscall_get_param(drcontext, 0) == (reg_t) STDERR){
        dr_syscall_set_result(drcontext, 0);
        dr_printf("skip syscall %d\n", sysnum);
        return false;
    }
    return true;
}
```



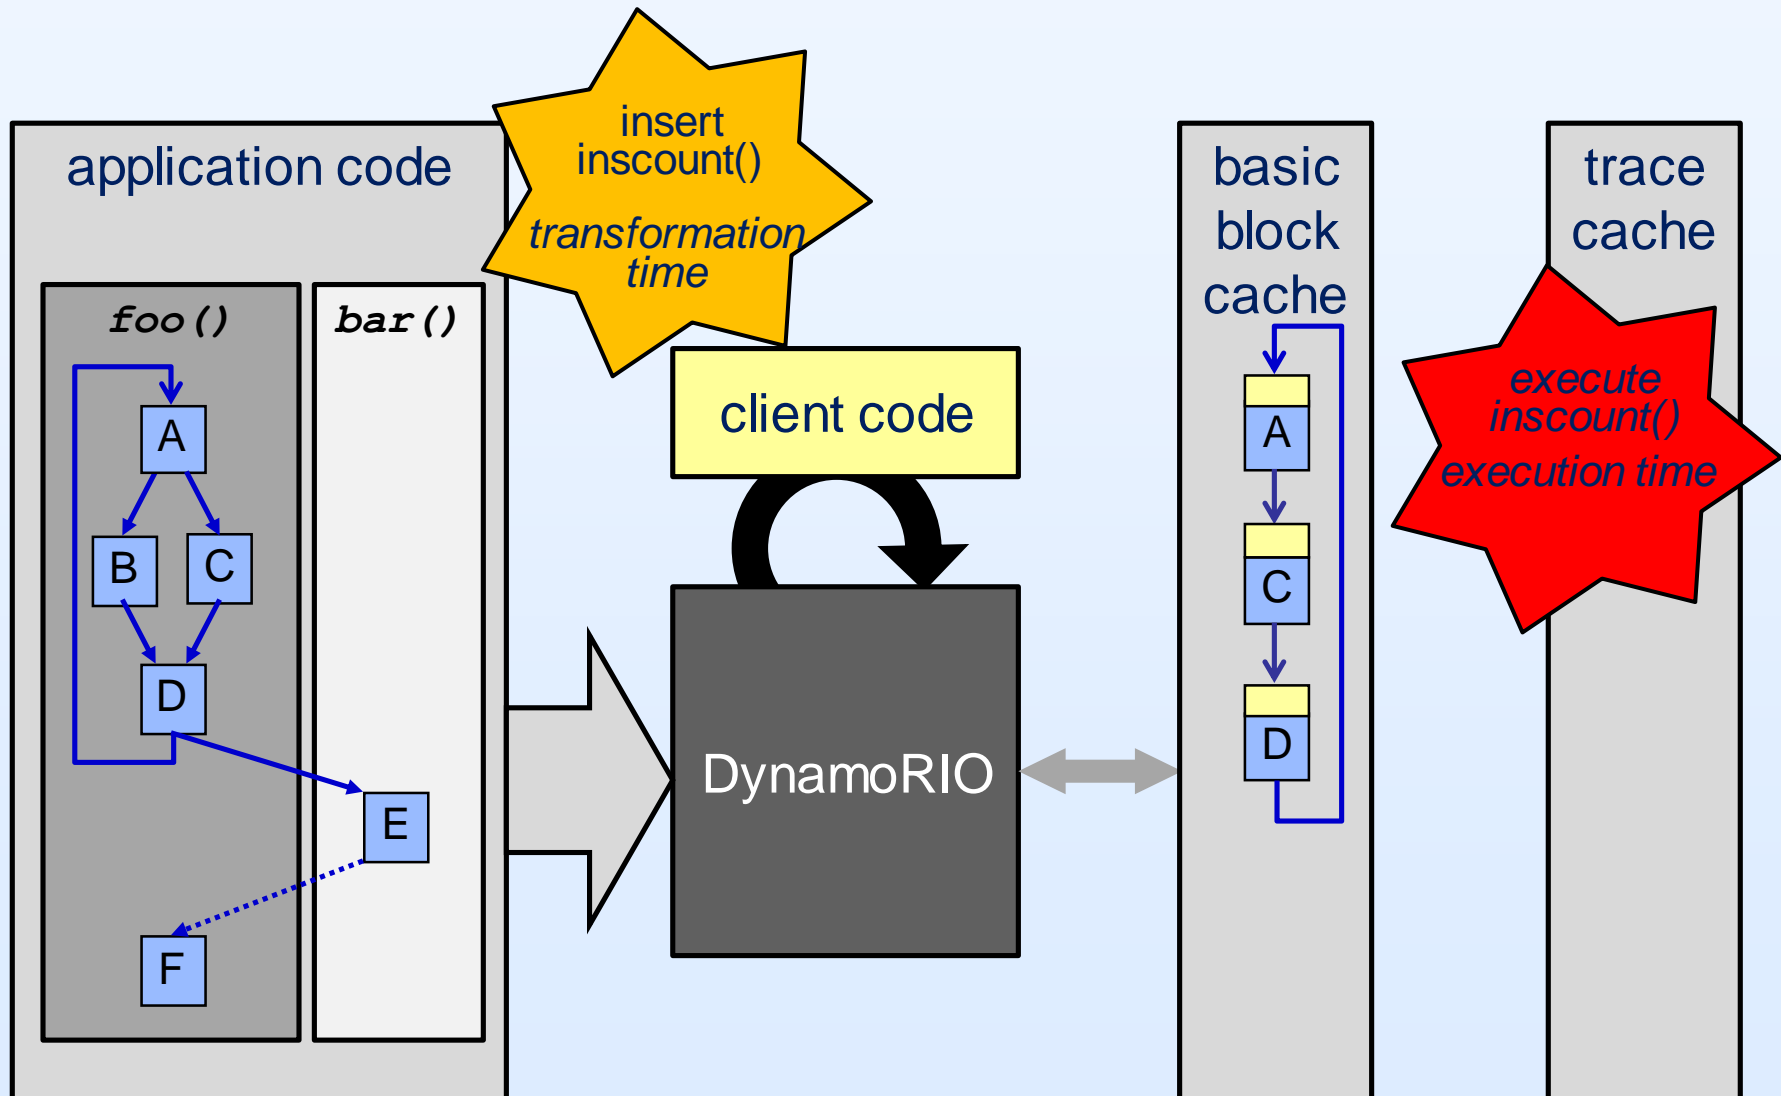
Part 2: Implement Event Callbacks - security



Write Your DynamoRIO Client

- Part 1: Register event callbacks
 - dr_client_main
 - drmgr_register_*_event()
 - App events: pre_syscall, ...
 - DR events: bb, ...
- Part 2: Implement event callbacks
 - Strace
- Part 3: Instrumentation
 - Dynamic instruction counting

Part 3: Instrumentation: instruction counting



Part 3: Instrumentation: instruction counting

```
DR_EXPORT void
dr_client_main(client_id_t id, int argc, const char *argv[]) {
    ...
    drmgr_register_bb_instrumentation_event
        (event_bb_analysis, event_app_instruction, NULL);
}

static uint global_count;
static void inscount(uint num_instrs) { global_count += num_instrs; }
```

Part 3: Instrumentation: instruction counting

```
static dr_emit_flags_t
event_bb_analysis(void *drcontext, void *tag, instrlist_t *bb,
                  bool for_trace, bool translating, void **user_data)
{
    instr_t *instr;
    uint num_instrs;

    for (instr = instrlist_first_app(bb), num_instrs = 0;
         instr != NULL;
         instr = instr_get_next_app(instr)) {
        num_instrs++;
    }

    *user_data = (void *) (ptr_uint_t) num_instrs;

    return DR_EMIT_DEFAULT;
}
```


Part 3: Instrumentation: instruction counting

```
static dr_emit_flags_t
event_app_instruction(void *drcontext, void *tag, instrlist_t *bb, instr_t *instr,
                    bool for_trace, bool translating, void *user_data)
{
    uint num_instrs;

    if (!drmgr_is_first_instr(drcontext, instr))
        return DR_EMIT_DEFAULT;

    num_instrs = (uint)(ptr_uint_t) user_data;
    dr_insert_clean_call(drcontext, bb, instrlist_first_app(bb),
                        (void *)inscount, false /* save fpstate */, 1,
                        OPND_CREATE_INT32(num_instrs));

    return DR_EMIT_DEFAULT;
}
```

Clean Calls

- C-code Callout
 - `dr_insert_clean_call{_ex}`
 - `dr_insert_{call,mbr,ubr,cbr}_instrumentation{_ex}`
 - Inserted clean callee is called from code cache on every execution of the fragment
- Full context switch
 - Save/restore full app state (expensive!)
- Automatic Optimization & Inlining
 - Clean callees are analyzed, optimized & potentially inlined
 - Optimization
 - partial context switch
 - Inlining if simple enough
 - `-opt_cleancall` runtime option controls aggressiveness

Counter Incrementing


- Instruction counting example uses a clean call for illustration purposes only
 - Additional actions can be added to the callee
 - If only a counter increment is needed, use `drx_insert_counter_update()` instead of the clean call
- `drx_insert_counter_update()`
 - Convenience routine for inserting inlined counter updates
 - Supports 32-bit and 64-bit counters
 - Supports atomic updates
 - Automatically saves the condition codes if necessary
 - Merges adjacent counter updates

Part 3: Instrumentation: instruction counting

```
TAG 0x7fb79f11a0de
+0  test  [7fb79f32ccc0h], 2
+7  jnz   7fb79f11a163h
END 0x7fb79f11a0de
```



```
TAG 0x7fb79f11a0de
+0  add   [722020a8h], 2
+10 test  [7fb79f32ccc0h], 2
+17 jnz   7fb79f11a163h
END 0x7fb79f11a0de
```



```
TAG 0x7fb79f11a0de
+0  mov    gs:[0h], rcx
+9  mov    rcx, gs:[20h]
+18 mov    [rcx + 2a8h], rdi
+25 mov    edi, 2
+30 mov    edi, edi
+32 add    [722020a8h], rdi
+39 mov    rdi, [rcx + 2a8h]
+46 mov    rcx, gs:[0h]
+55 test   [7fb79f32ccc0h], 2
+62 jnz    7fb79f11a163h
END 0x7fb79f11a0de
```

How to Create Your Own Tool

- Event driven
 - Application, DynamoRIO, client, and Extensions
 - Events
- Write your DynamoRIO client
 - Part 1: Register event callbacks
 - Part 2: Implement event callbacks
 - Part 3: Instrumentation
- Config, build, and run
 - CMake

Config, Build, and Run

- CMake
 - <http://www.cmake.org/>
 - Generates build files for native compiler of choice
 - Makefiles for UNIX, nmake, etc.
 - Visual Studio project files
- CMakeLists.txt

```
add_library(myclient SHARED myclient.c)
find_package(DynamoRIO)
if (NOT DynamoRIO_FOUND)
    message(FATAL_ERROR "DynamoRIO package
        required to build")
endif(NOT DynamoRIO_FOUND)
configure_DynamoRIO_client(myclient)
use_DynamoRIO_extension(myclient drmgr)
```

Config, Build, and Run

- Config
 - `cmake /path/to/your/client/`
 - `ccmake`, `cmake-gui`
- Build
 - `make`
 - `cmake --build .`
- Run
 - Method 1 (one step)
 - `drrun -c <client> <client options> -- <app cmdline>`
 - Method 2 (two steps, for better child process blacklisting, etc.)
 - `drconfig -reg <appname> -c <client> <client options>`
 - `drinject <app cmdline>`

Runtime Options

- Pass options to drconfig/drrun
- A large number of options; the most relevant are:
 - -t <tool>
 - -c <client lib> <client options>
 - -thread_private
 - -follow_children
 - -opt_cleancall
 - -tracedump_text and -tracedump_binary
 - -prof_pcs
 - -code_api
 - -help

Non-Standard Deployment

- drdecode
 - Static IA-32/AMD64/ARM/AArch64 decoding/encoding/instruction manipulation library
- Standalone API
 - Use DynamoRIO as a library of IA-32/AMD64/ARM/AArch64 manipulation routines plus cross-platform file i/o, locks, etc.
- Start/Stop API
 - Can instrument source code with where DynamoRIO should control the application

Debugging Your DynamoRIO Client

- Run with the debug build of DynamoRIO
 - `drrun -debug ...`
 - `-stderr_mask 0xN`
 - `-ignore_assert_list '*'`
 - `-disable_traces`
- Run with logging
 - `drrun -debug -loglevel 3`
- Run under a debugger
 - `drrun -debug -loglevel 3 -msgbox_mask 0xf ...`
 - `-no_hide` (windows)
 - Attach debugger to the process

Using a Debugger on a DynamoRIO Client

- Debug client code
 - Load debug symbol for your client
 - `add-symbol-file '/path/to/your/client/libname.so' 0xxxxxxx`
 - Normal debugging
- Debug instrumented code
 - Printf debugging
 - `{opnd, instrlist, instr}_disassemble`
 - Breakpoint
 - `build_basic_block_fragment`
 - `build_bb_ilist`
 - `enter_fcache`
 - Watchpoint (data breakpoint)

Obtaining Help

- Read the documentation
 - <http://dynamorio.org/docs/>
- Look at the provided tools and sample clients
 - In the documentation
 - In the release package: samples/
- Ask on the DynamoRIO Users discussion forum/ mailing list
 - <http://groups.google.com/group/dynamorio-users>

Reporting Bugs

- Search the issue tracker linked from <http://dynamorio.org> first
 - <https://github.com/DynamoRIO/dynamorio/issues>
- File a new issue if not found
- Follow conventions on wiki
 - <https://github.com/DynamoRIO/dynamorio/wiki/Bug-Reporting>
 - CRASH, APP CRASH, HANG, ASSERT
- Example titles:
 - CRASH (1.3.1 calc.exe)
vm_area_add_fragment:vmareas.c(4466)
 - ASSERT (1.3.0 suite/tests/common/segfault)
study_hashtable:fragment.c:1745 ASSERT_NOT_REACHED

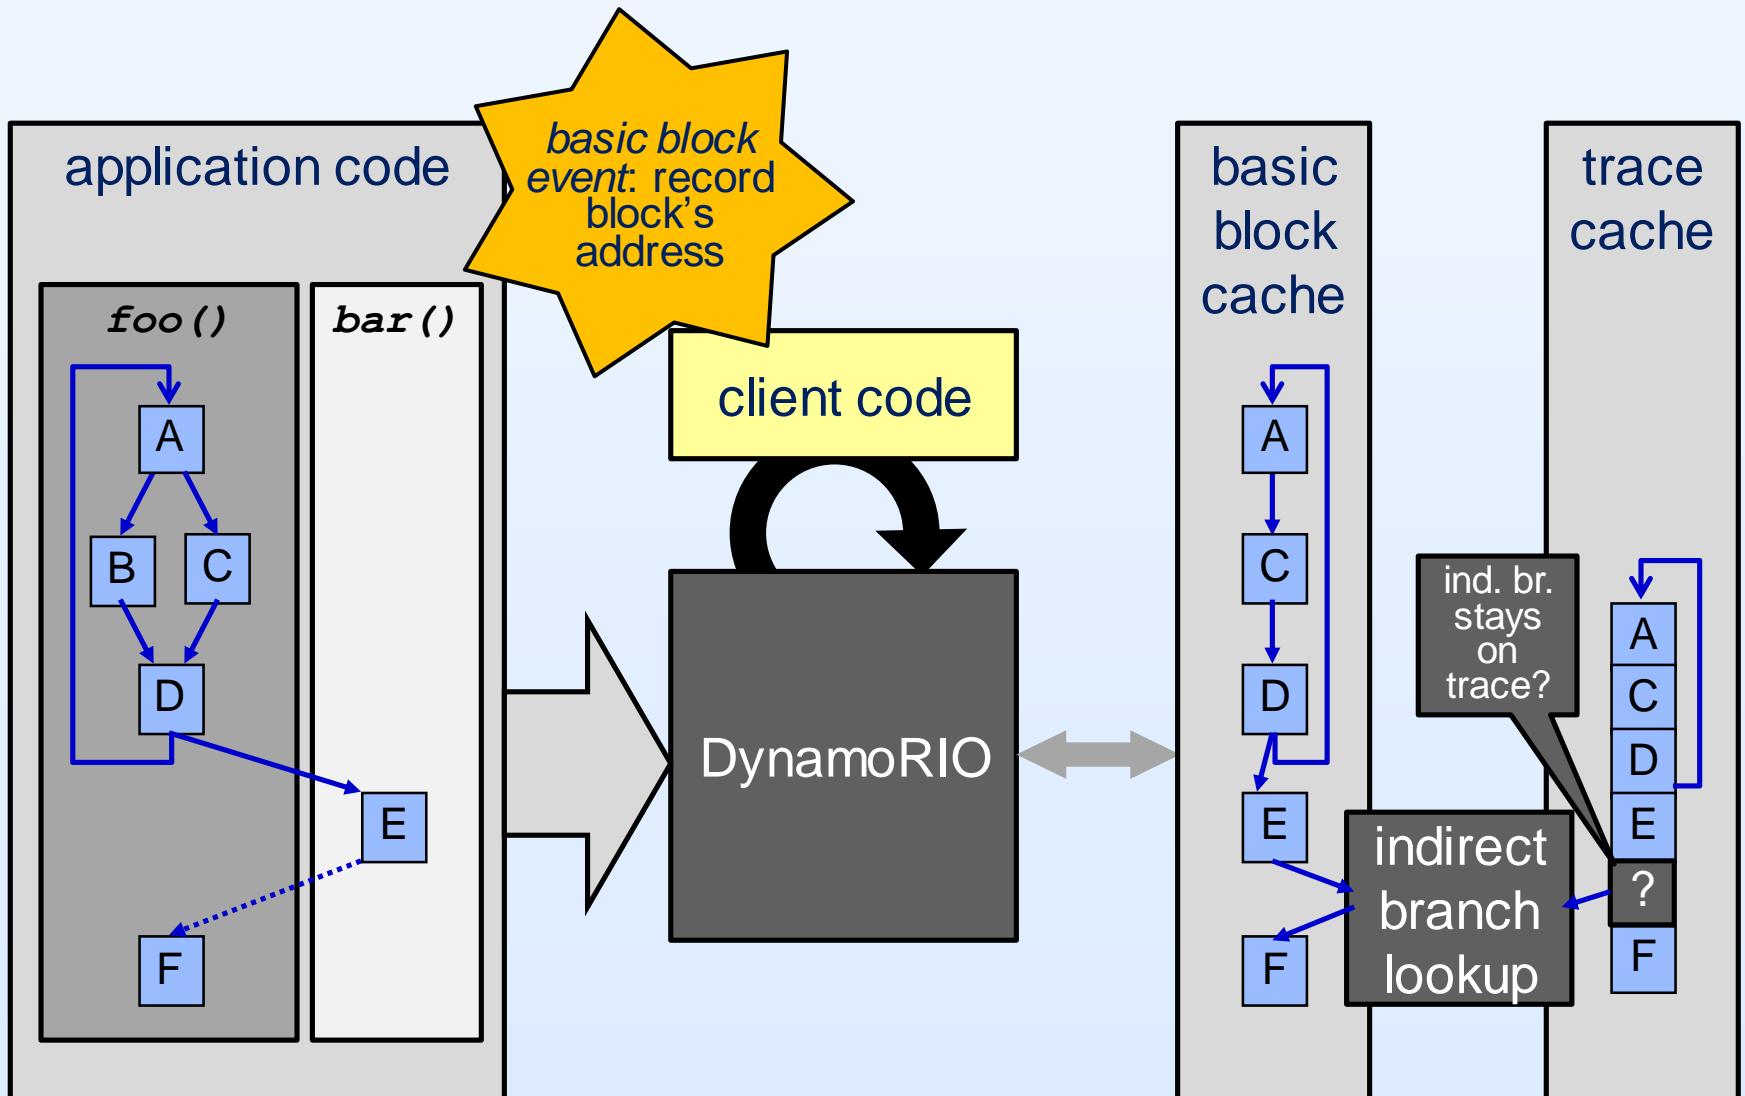
Tool Internals

8:30- 8:40	Welcome + DynamoRIO History
8:40- 9:10	Tool Demonstrations
9:10- 9:30	DynamoRIO System Overview
9:30- 9:45	DynamoRIO API Part 1
9:45-10:00	ARM Status + Demonstrations
10:00-10:30	<i>Break</i>
10:30-10:55	DynamoRIO API Part 2
10:55-11:20	How to Create Your Own Tool
11:20-11:40	Tool Internals
11:40-11:45	Q & A

Dr. Cov

- Collects code coverage information
- Records which basic blocks have been executed
- Operates on unmodified binaries
- Results are post-processed into the standard lcov format
- Usage:
 - `drrun -t drcov -- <app>`
 - `tools/bin32/drcov2lcov -input drcov*.log`
 - Now pass the resulting data file to lcov's genhtml script

Dr. Cov Implementation

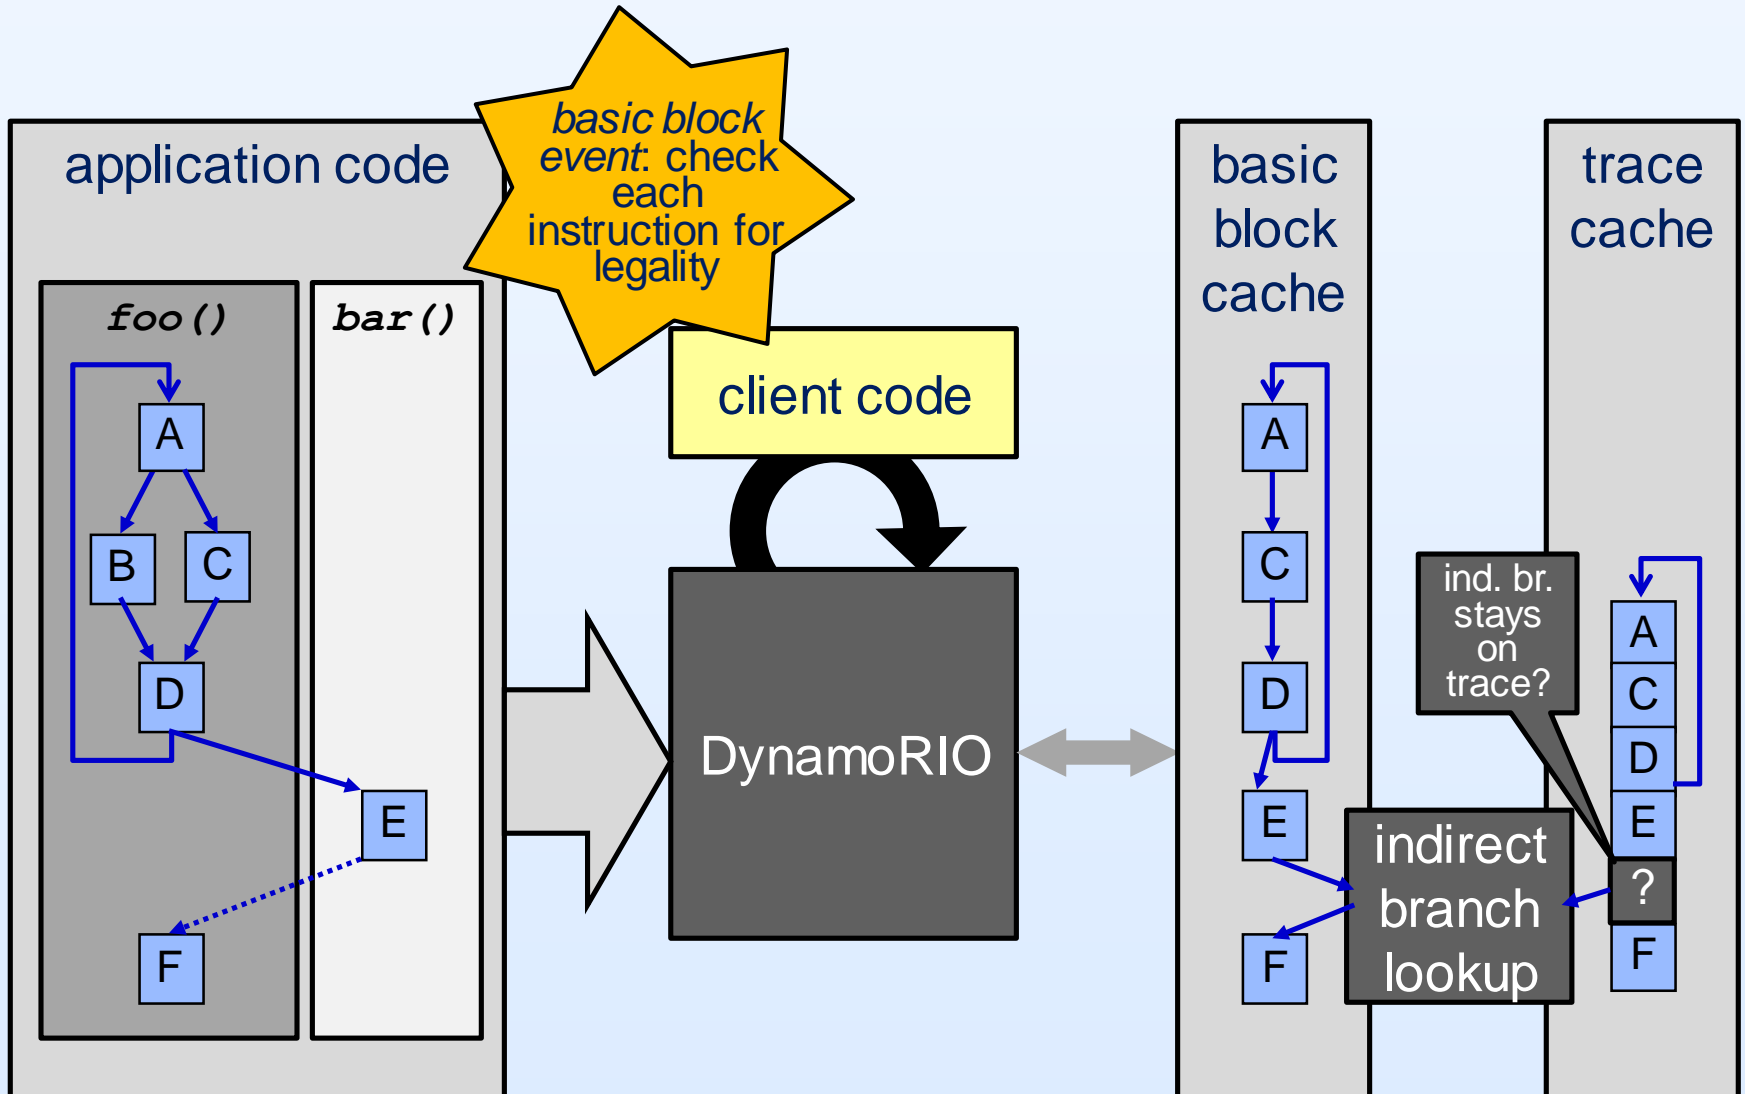


Dr. CPUSim: Legacy CPU Checker

- A tool for testing that applications will run correctly on legacy processors
- Looks for instructions that are illegal on the specified CPU model
- Fools `cpuid` checks
- Usage:
 - `drrun -t drcpusim -cpu PentiumPro -- <app>`
- Example:

```
% bin32/drrun -t drcpusim -cpu Pentium -continue -- hello.exe
<Invalid Pentium instruction "movd" @ hello.exe+0x85c3.
Continuing.>
```

Dr. CPUSim Implementation

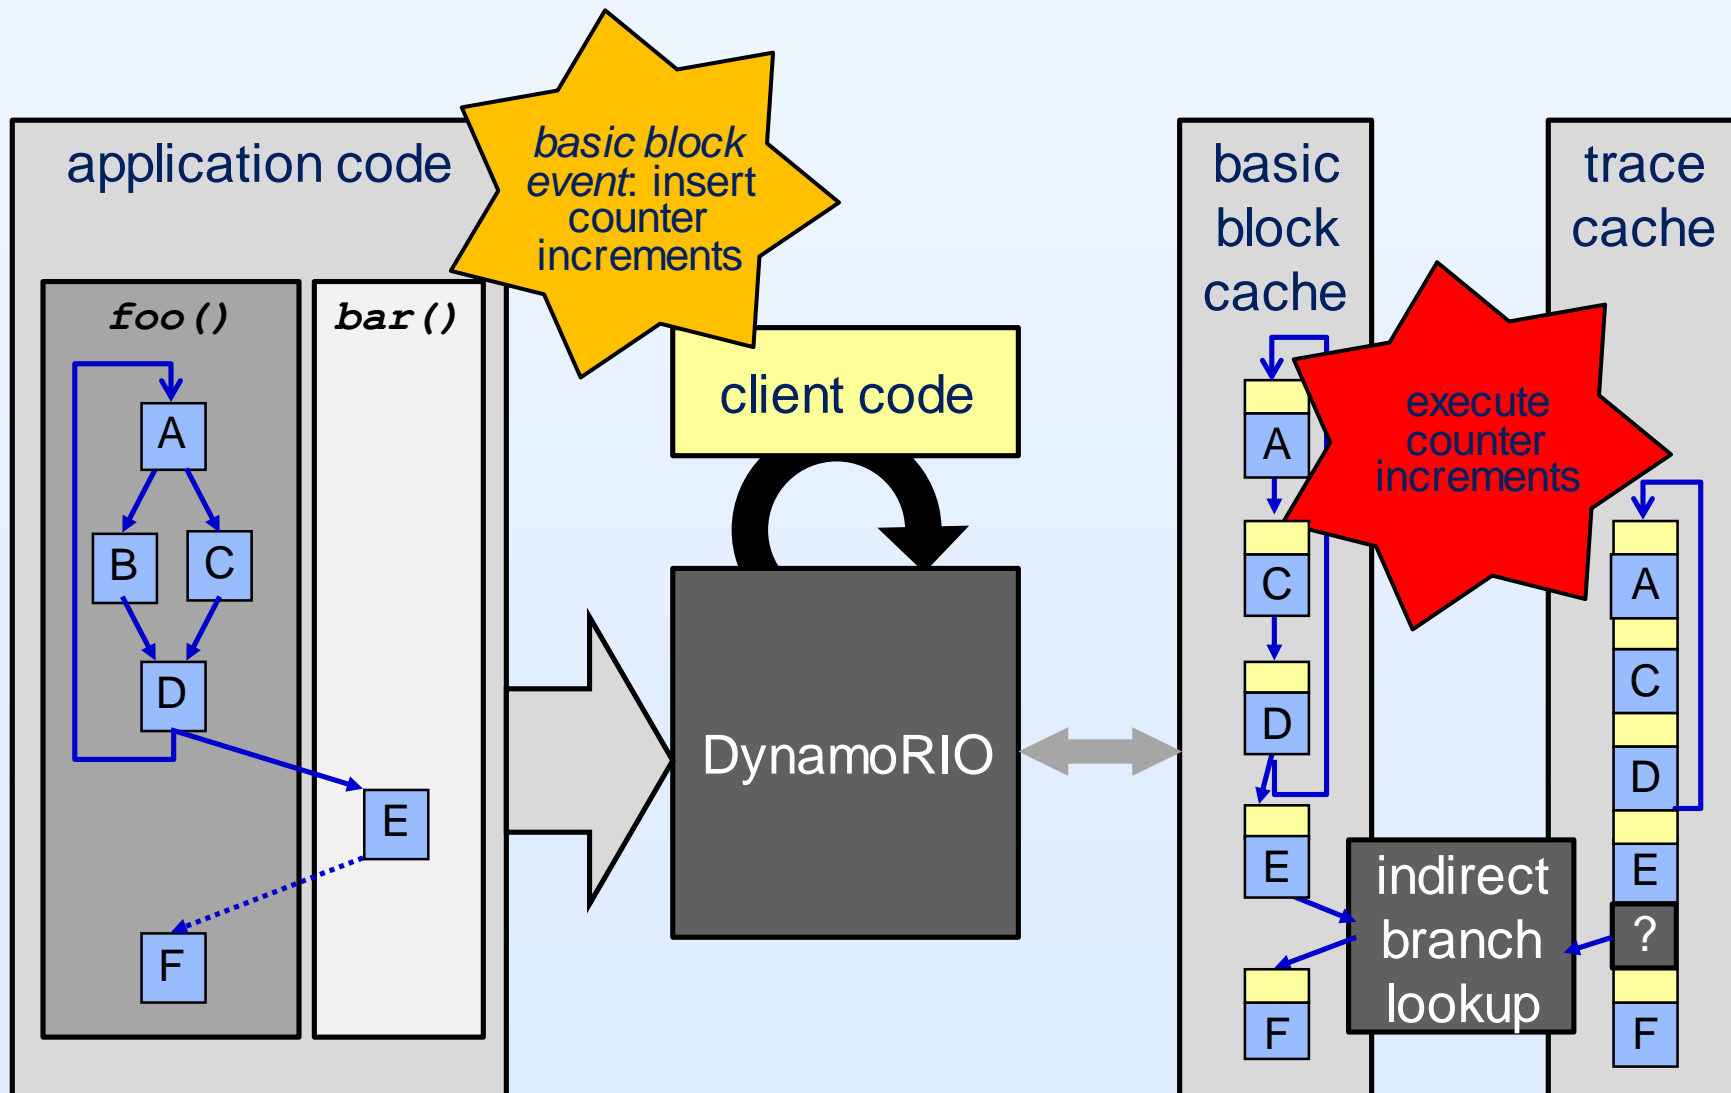


Opcode Mix Tool

- Monitors every instruction executed
- Counts the dynamic execution of each ISA opcode for the target application workload
- Usage:
 - `drrun -c samples/bin32/opcodes.dll -- <app>`
- Example:

```
% bin64/drrun -c samples/bin64/libopcodes.so -- gzip
lib64/libdrinjectlib.a
Client opcodes is running
Top 15 opcode execution counts in 64-bit AMD64 mode:
3932460 : shl
4607686 : shr
6736566 : cmovnb
<..>
```

Opcode Mix Tool Implementation



Dr. Strace: System Call Tracer

- Monitors all system calls executed by a target application
- Prints a trace of system calls and their arguments
- Operates on Windows, Linux, and Android
- Usage:

- `drrun -t drstrace -- <app>`

- Example:

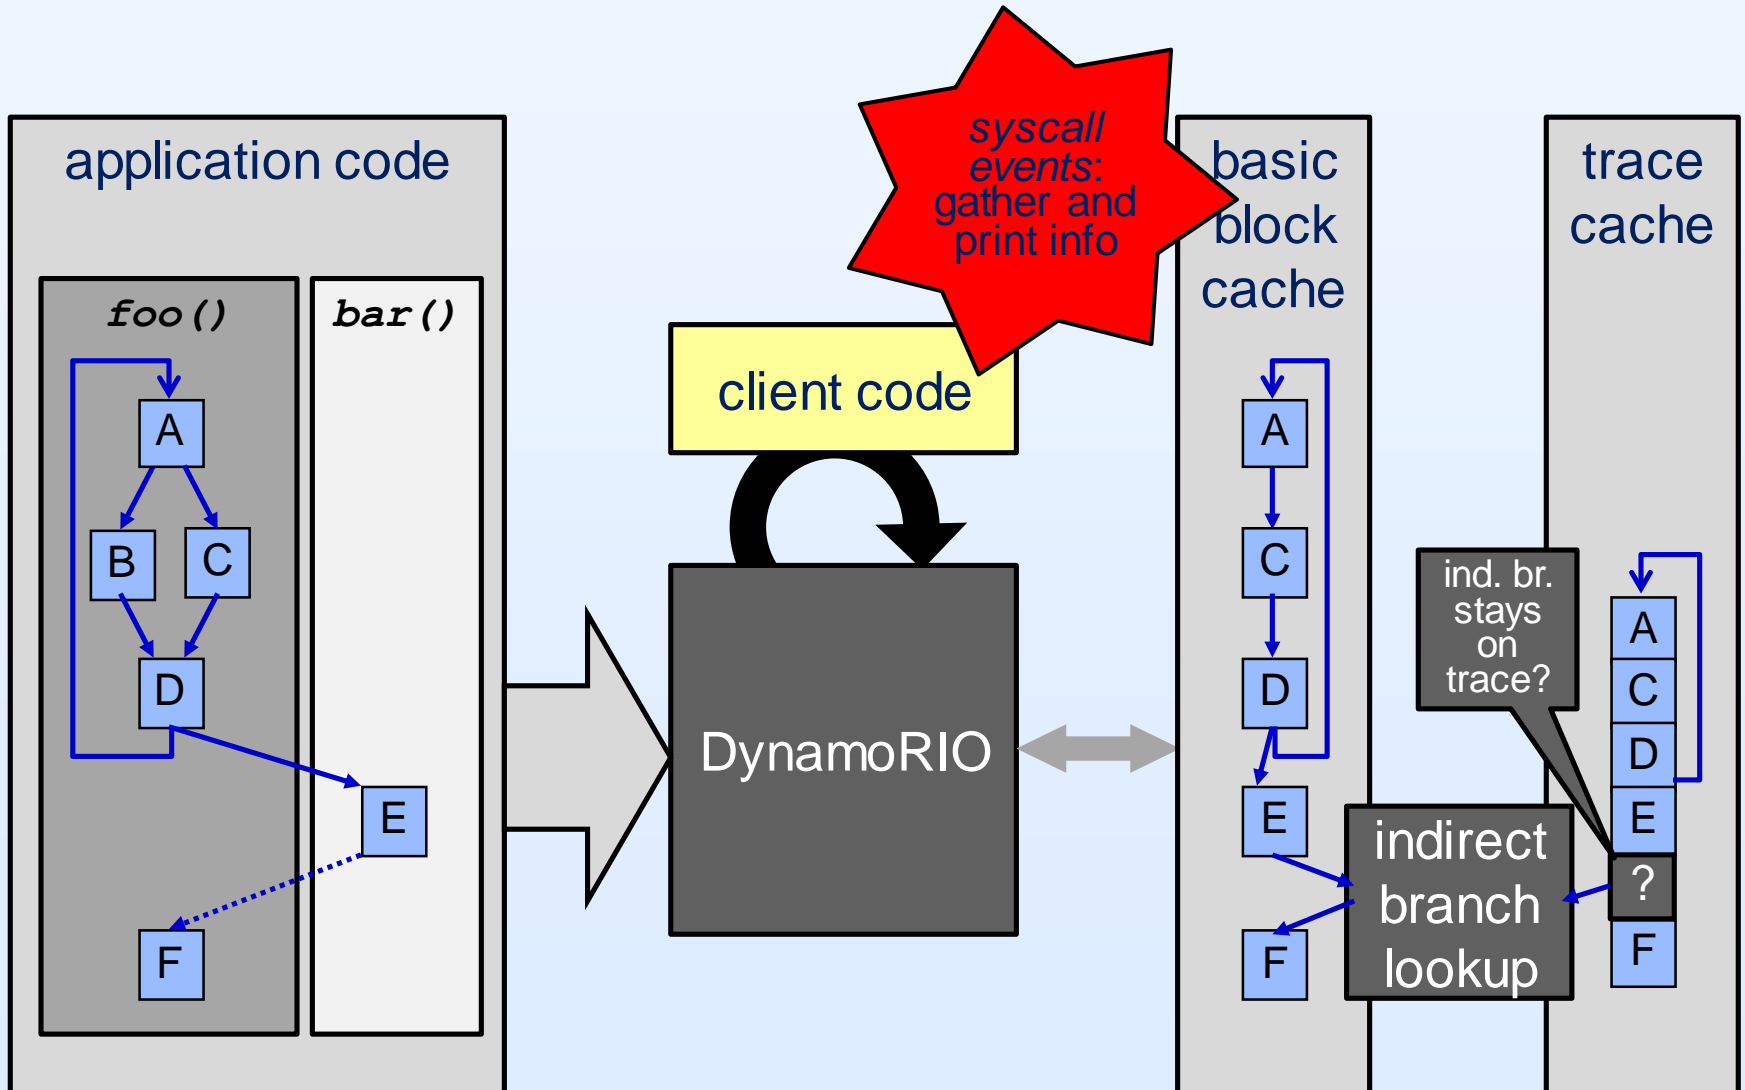
NtOpenKeyEx

```
arg 0: 0x001fcd0c (type=HANDLE*, size=0x4)
arg 1: 0x109 (type=unsigned int, size=0x4)
arg 2: len=0x18, root=0x3c, name=150/152 "SOFTWARE\Microsoft\Windows
      NT\CurrentVersion\LanguagePack\SurrogateFallback", att=0x40,
      sd=0x00000000, sqos=0x00000000 (type=OBJECT_ATTRIBUTES*, size=0x4)
arg 3: REG_OPTION_RESERVED or REG_OPTION_NON_VOLATILE (type=named constant,
      value=0x0, size=0x4)
```

succeeded =>

```
arg 0: 0x001fcd0c => 0x134 (type=HANDLE*, size=0x4)
retval: 0x0 (type=NTSTATUS, size=0x4)
```

Dr. Strace Implementation



Dr. Syscall System Call Database

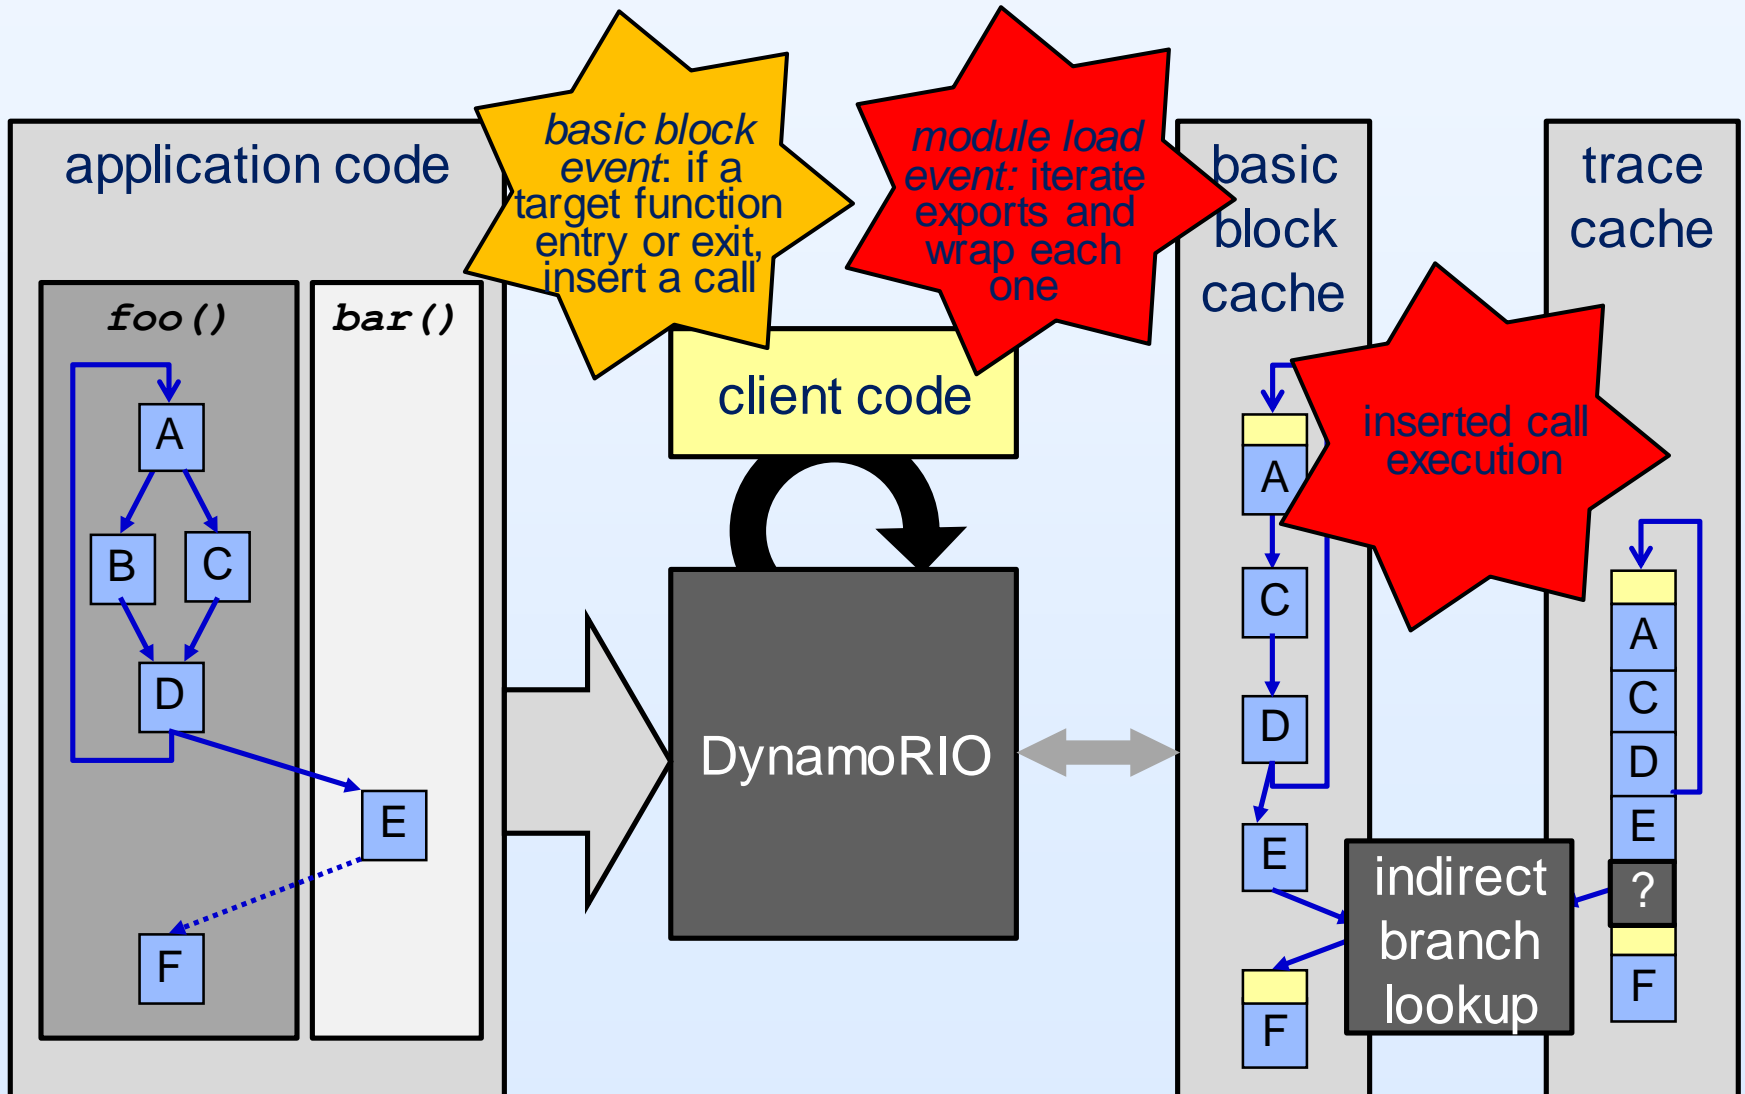
- *drsyscall* Extension
- Windows, Linux, and MacOS system call data
- Library interface to convert numbers to names or vice versa
- Static argument type iteration
- Dynamic argument value iteration
- Dynamic success querying
- Dynamic iteration over each memory region read or written

Dr. Ltrace: Library Tracer

- Intercepts calls to exported library functions
- Prints a trace of each function called and some argument information
- Usage:
 - `drrun -t drltrace -- <app>`
- Example:

```
~~~~ ntdll.dll!RtlEnterCriticalSection(0x006fc758, 0x0018fe50)
~~~~ KERNEL32.dll!WriteFile(0x00000268, 0x0018e9b0)
~~~~ KERNELBASE.dll!WriteFile(0x00000268, 0x0018e9b0)
Hello, world!
~~~~ ntdll.dll!RtlLeaveCriticalSection(0x006fc758, 0x0018fe50)
```


Dr. Ltrace Implementation

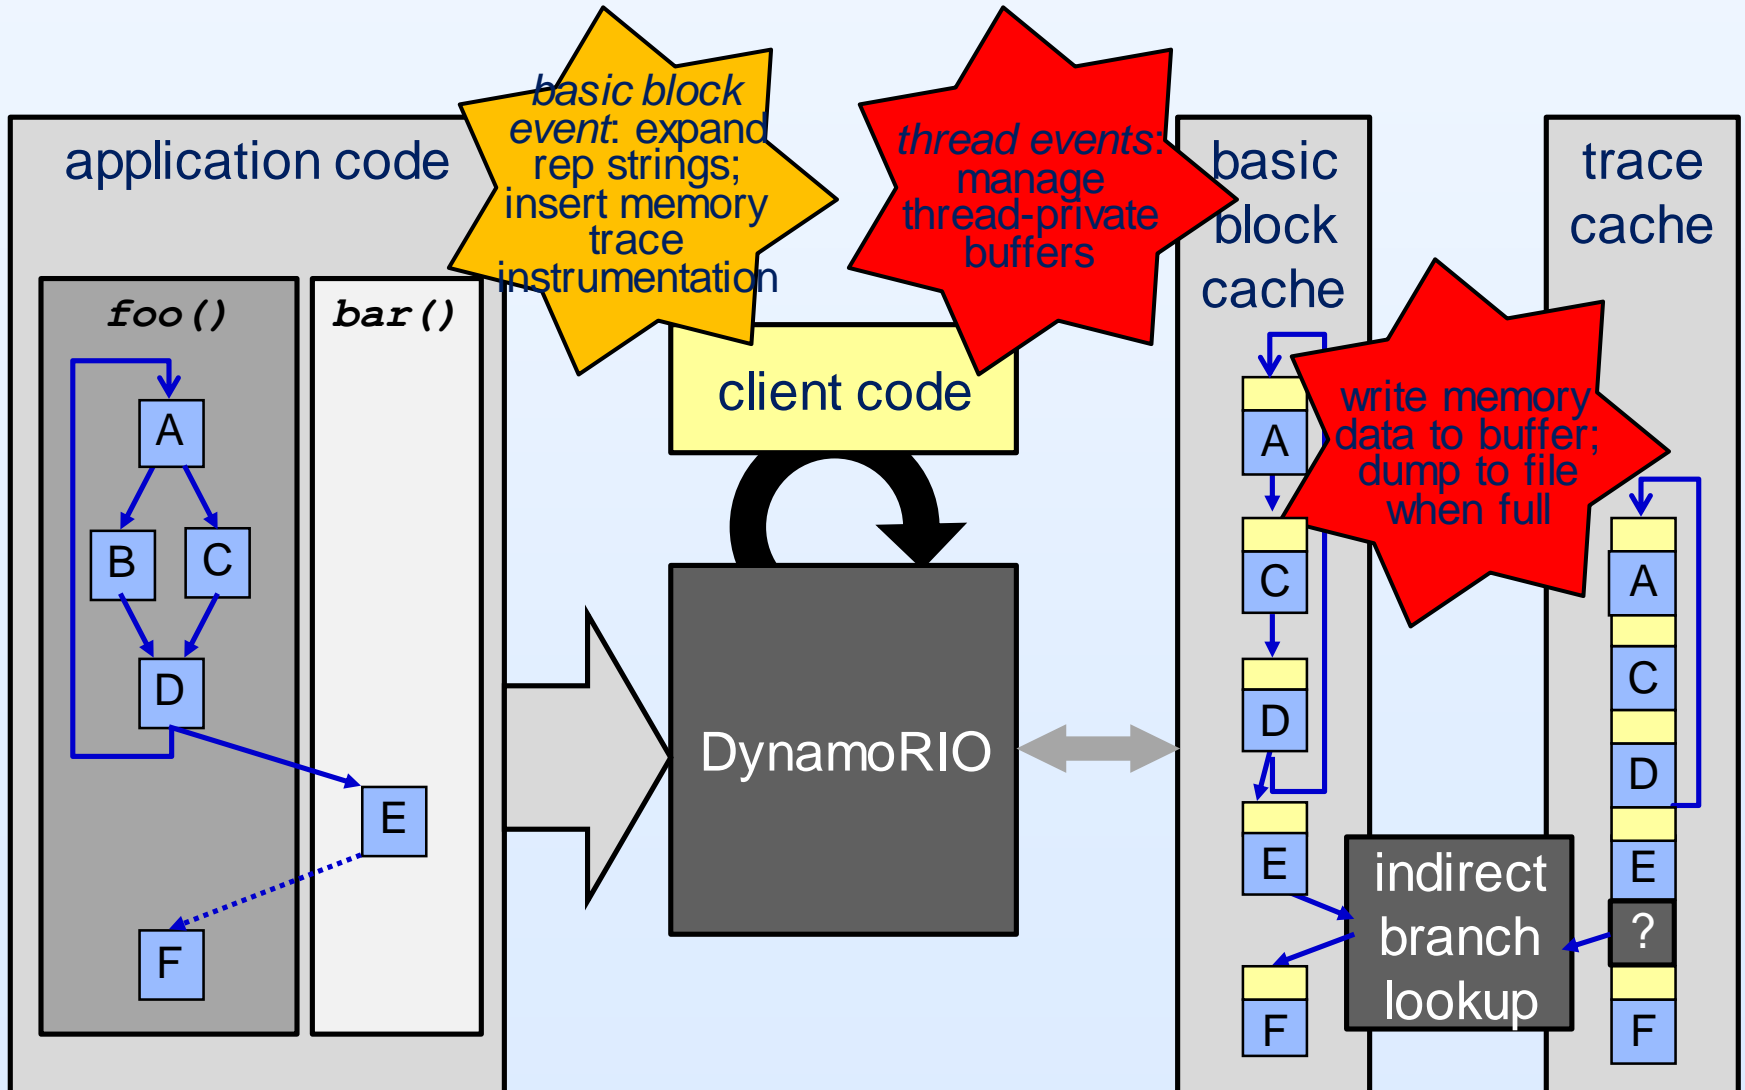


Memory Trace Tool

- Instruments every load and store
- Writes the trace of memory accesses to a file
- Usage:
 - `drrun -c samples/bin64/libmemtrace_x86_text.so -- <app>`
- Results:

Format: <instr address>,<(r)ead/(w)rite>,<data size>,<data address>
0x00007fb31a5462d3,w,8,0x00007ffc4febfb18
0x00007fb31a549a8f,r,8,0x00007fb31a767e70
0x00007fb3199a54fc,r,4,0x00007ffc4febf438
0x00007fb3199a5526,r,1,0x0000000000400e98
0x00007fb31a549a96,w,8,0x00007fb31a767c98
0x00007fb3199af4a0,r,8,0x00007ffc4febf9f8
0x0000000000400db4,r,8,0x00007ffc4febfa48

Memory Trace Tool Implementation

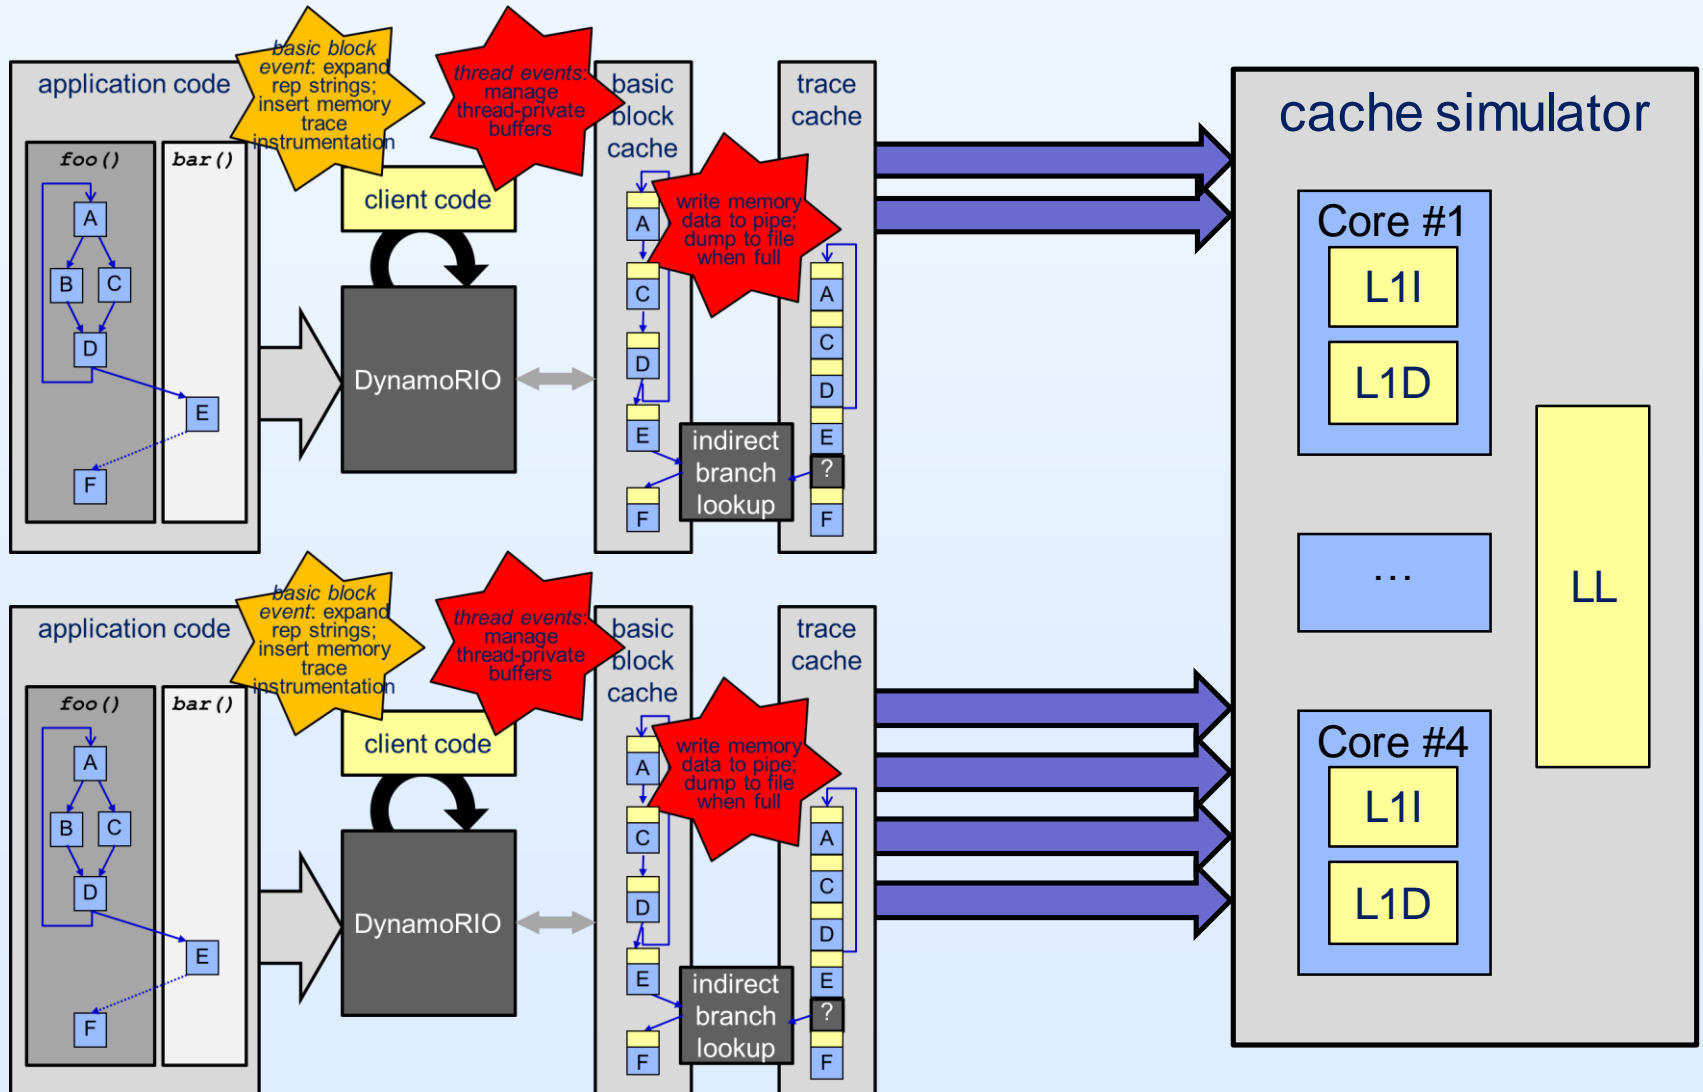


Dr. CacheSim

- Multi-process, multi-thread online cache simulator
- Each thread collects and sends its own memory access trace to a single simulator
- Usage:
 - `drrun -t drcachesim -- <app>`
- Results:

```
% bin64/drrun -t drcachesim -- ./multi_threaded_pi_estimator
Estimation of pi is 3.142425985001098
---- <application exited with code 0> ----
Core #0 (1 thread(s))
  L1I stats:
    Hits:                243,293
    Misses:               1,077
    Miss rate:            0.44%
  L1D stats:
    Hits:                67,738
    Misses:               3,471
    Miss rate:            4.87%
Core #1 (1 thread(s))
  L1I stats:
    Hits:                8,589
    ...
```

Dr. CacheSim Implementation

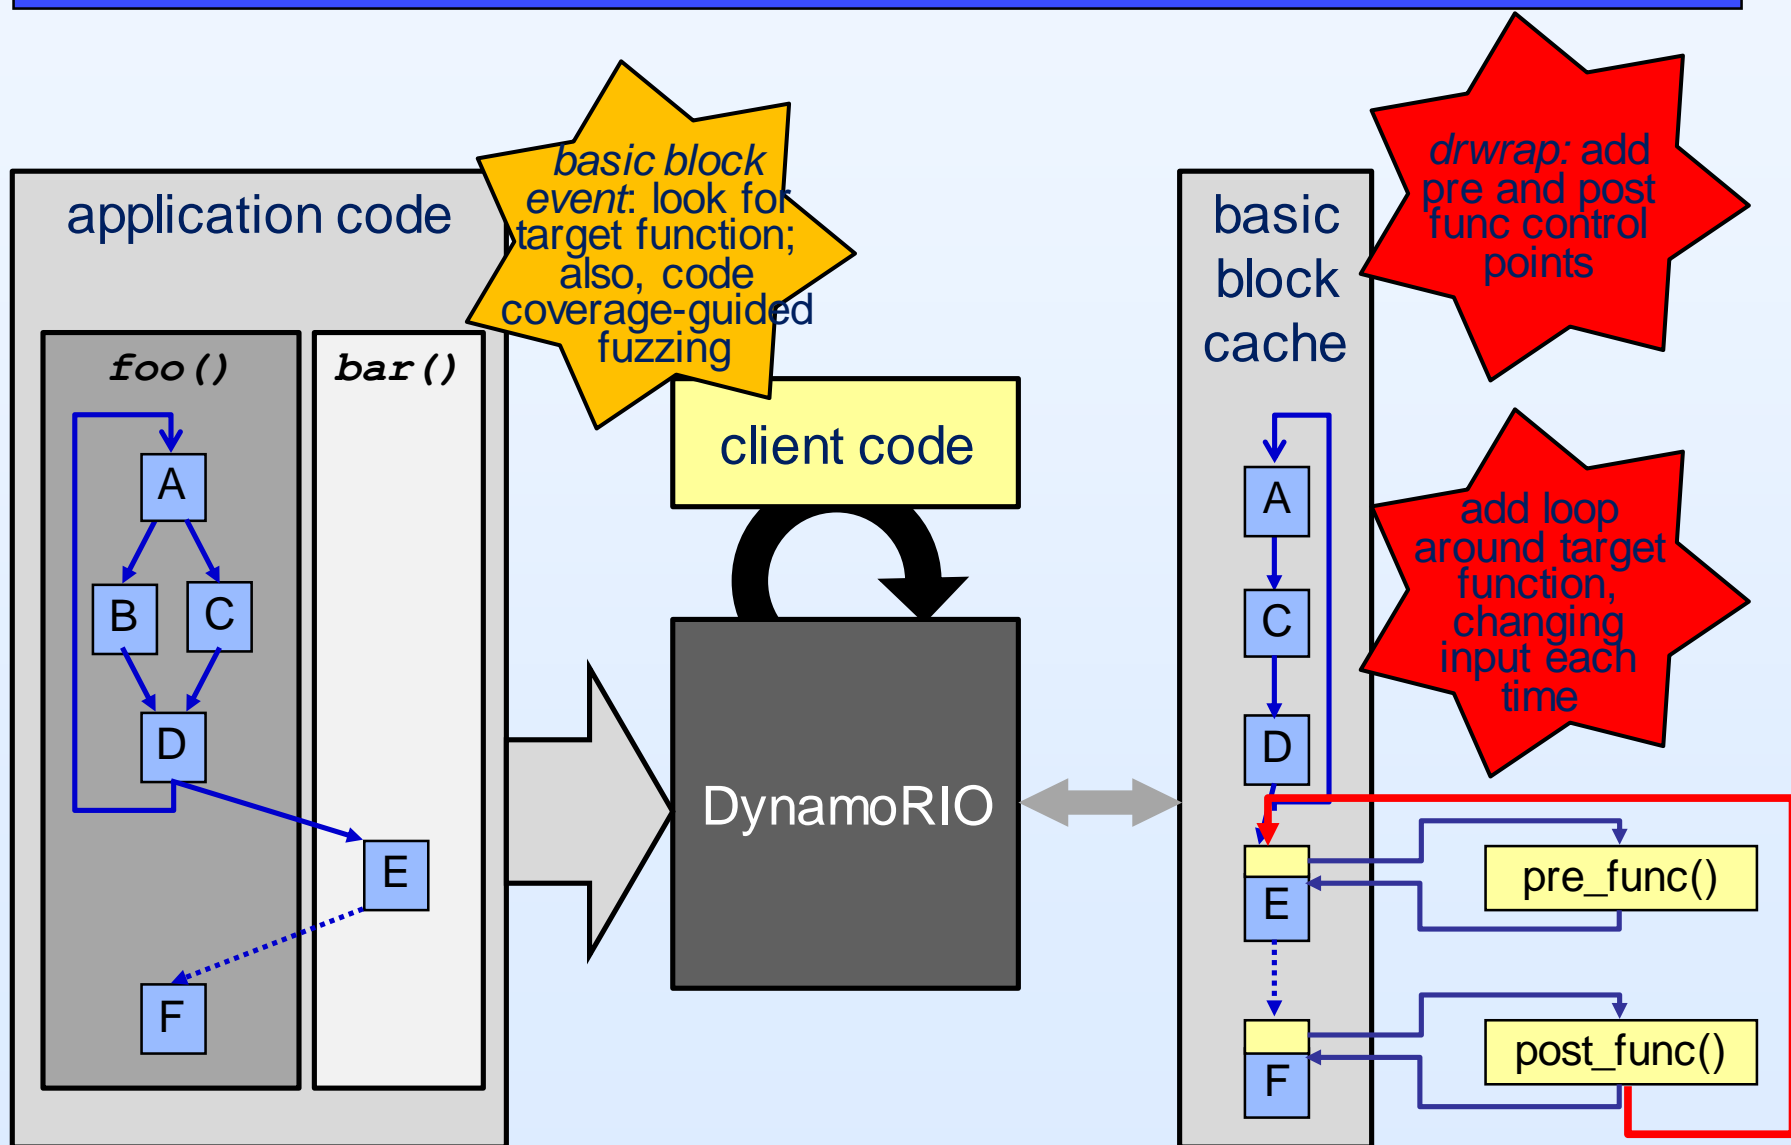


Dr. Fuzz

- In-process function-level fuzzing
- Repeatedly runs a target function, varying the inputs
- Basic usage:
 - `drrun -t drmemory -fuzz_function <func> -- <app>`
- Results:

```
% drrun -t drmemory -batch -fuzz_function some_function -fuzz_num_iters 10 -- ./fuzzme.exe
~~Dr.M~~ Dr. Memory version 1.11.0
~~Dr.M~~ Running "./fuzzme.exe"
in some_function
hello!
in some_function
~~Dr.M~~
~~Dr.M~~ Error #1: UNADDRESSABLE ACCESS beyond heap bounds: reading 0x00f58efc-0x00f58f00 4 byte(s)
~~Dr.M~~ # 0 some_function [c:\derek\cgo2017\demos\fuzzme.c:8]
~~Dr.M~~ # 1 main [c:\derek\cgo2017\demos\fuzzme.c:18]
in some_function
in some_function
hello!
...
```

Dr. Fuzz Implementation



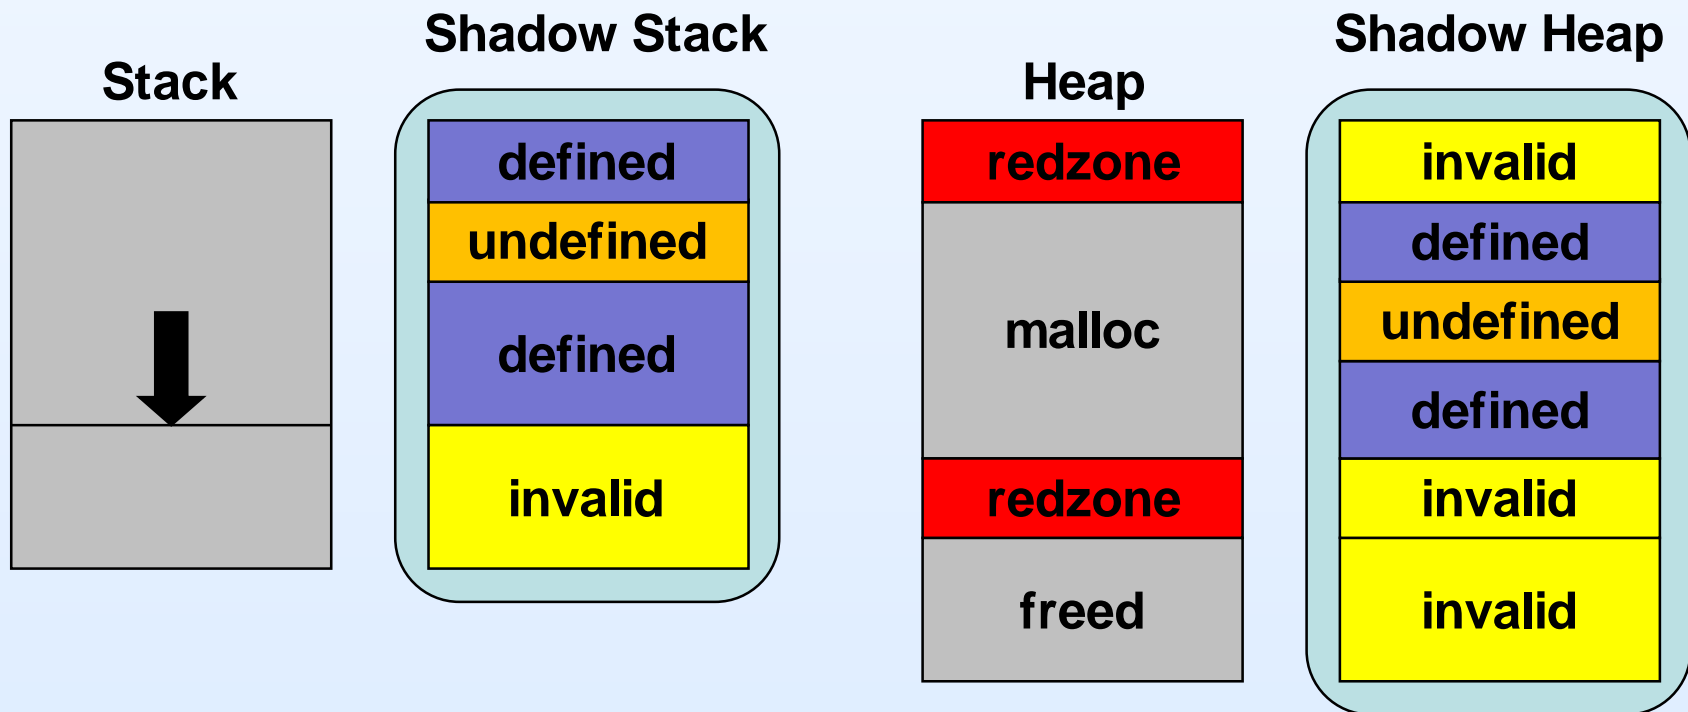
Dr. Memory: Memory Debugger

- Detects reads of uninitialized memory
- Detects heap errors
 - Out-of-bounds accesses (underflow, overflow)
 - Access to freed memory
 - Invalid frees
 - Memory leaks
- Detects other accesses to invalid memory
 - Stack tracking
- Detects Windows handle leaks
- Operates at runtime on unmodified Windows, Linux, MacOS, and Android binaries

Dr. Memory Instrumentation

- Monitor all memory accesses, stack adjustments, and heap allocations
- Shadow each byte of app memory
- Each byte's shadow stores one of 3 values:
 - Unaddressable
 - Uninitialized
 - Defined

Dr. Memory Shadow Memory



Monitoring Stack Changes

- As stack is extended and contracts again, must update stack shadow as unaddressable vs uninitialized
- Push, pop, or any write to stack pointer
- Try to distinguish large alloc/dealloc from stack swap

Partial-Word Defines But Whole-Word Transfers

- Sub-dword variables are moved around as whole dwords
- Cannot raise error when a move reads uninitialized bits
- Must propagate on moves and thus must shadow registers
 - Propagate shadow values by mirroring app data flow
- Check system call reads and propagate system call writes
 - Else, false negatives (reads) or positives (writes)
- Raise errors instead of propagating at certain points
 - Report errors only on “significant” reads

Shadowing Registers

- Use multiple TLS slots
 - `dr_raw_tls_calloc()`
 - Alternative: steal register
- Can read and write w/o spilling
- Bring into spilled register to combine w/ other args
 - Defined=0, uninitialized=1
 - Combine via bitwise or

Types Of Instrumentation

- Clean call
 - Simplest, but expensive in both time and space: full context switch from application state to tool state with separate stack to execute C code
- Shared clean call (now auto-provided by DR)
 - Saves space
- Lean procedure
 - Shared routine with smaller context switch than full clean call
 - Jump-and-link rather than swapping stack
 - Array of routines, one per pair of dead registers
- Inlined
 - Smallest context switch, but should limit to small sequences of instrumentation

Non-Code-Cache Code

- Use `dr_nonheap_alloc()` to allocate space to store code
- Generate code using DR's IR and emit to target space
- Mark read-only once emitted via `dr_memory_protect()`

Jump-and-Link

- Rather than using call+return, avoid stack swap cost by using jump-and-link
 - Store return address in a register or TLS slot
 - Direct jump to target
 - Indirect jump back to source

```
PRE(bb, inst, INSTR_CREATE_mov_st(drcontext,  
    spill_slot_opnd(drcontext, SPILL_SLOT_2),  
    opnd_create_instr(appinst)));  
PRE(bb, inst, INSTR_CREATE_jmp(drcontext,  
    opnd_create_pc(shared_slowpath_region)));  
...  
PRE(ilist, NULL, INSTR_CREATE_jmp_ind(drcontext,  
    spill_slot_opnd(SPILL_SLOT_2)));
```


Using Faults For Faster Common Case Code

- Instead of explicitly checking for rare cases, use faults to handle them and keep common case code path fast
- Signal and exception event and restore state extended event all provide pre- and post-translation contexts and containing fragment information
- Client can return failure for extended restore state event
 - When can support re-execution of faulting cache instr, but not re-start translation for relocation

Suspending The World

- Use case: Dr. Memory leak check
 - GC-like memory scan
- Use `dr_suspend_all_other_threads()` and `dr_resume_all_other_threads()`
- Cannot hold locks while suspending

Tool Packaging

- DynamoRIO is redistributable, so you can include a copy with your tool
- drrun supports the `-t` option via a tool configuration file
 - `drrun -t drcov -- <app cmdline>`
- Custom front end to configure and launch app
 - We provide several libraries for building tool front ends:
 - `drconfiglib`
 - `drinjectlib`
 - `drfrontendlib`

Q & A / Feedback

8:30- 8:40 Welcome + DynamoRIO History

8:40- 9:10 Tool Demonstrations

9:10- 9:30 DynamoRIO System Overview

9:30- 9:45 DynamoRIO API Part 1

9:45-10:00 ARM Status + Demonstrations

10:00-10:30 Break

10:30-10:55 DynamoRIO API Part 2

10:55-11:20 How to Create Your Own Tool

11:20-11:40 Tool Internals

11:40-11:45 Q & A

Seeking Contributors

- Many potential projects
 - Contribute new tools
 - Finish feature to attach to a running process
 - Debugger integration
 - Transparent debugging and debugger extensions
 - ARM performance
 - Optimizing assembly code and generated code
 - Android
 - AArch64 support
 - Support running larger apps
 - MacOS
 - Private loader and 64-bit support need owners
 - Search “GoodContrib” label in issue tracker

Optional Slides: Versus Pin

DynamoRIO versus Pin

- DynamoRIO is open-source while Pin is closed
 - If something goes wrong, can look under the hood
 - Can adapt to particular usage scenario
 - E.g., DynamoRIO as static library linked into application
- DynamoRIO supports both Intel (x86, AMD64) and ARM (ARM, AArch64) platforms while Pin supports only Intel


DynamoRIO versus Pin: Interface

- Basic interface is fundamentally different
- Pin = insert callout/trampoline only
 - Not so different from tools that modify the original code: Dyninst, Vulcan, Detours
 - Uses code cache only for transparency
- DynamoRIO = arbitrary code stream modifications
 - Only feasible with a code cache
 - Takes full advantage of power of code cache
 - General IA-32/AMD64/ARM/AArch64 decode/encode/IR support


DynamoRIO versus Pin: Consequences

- Pin = insert callout/trampoline only
 - Pin tries to inline and optimize
 - Client has little control or guarantee over final performance
- DynamoRIO = arbitrary code stream modifications
 - Client has full control over all inserted instrumentation
 - Result can be significant performance difference
 - PiPA Memory Profiler+ Cache Simulator:
3.27x speedup w/ DynamoRIO vs 2.6x w/ Pin
 - DynamoRIO also performs callout (“clean call”) optimization and inlining just like Pin for less performance-focused clients


Base Performance Comparison (No Tool)




Base Performance Comparison (No Tool)



Base Memory Comparison (No Tool)



Base Memory Comparison (No Tool)



BBCount Pin Tool

```
static int bbcount;

VOID PIN_FAST_ANALYSIS_CALL docount() { bbcount++; }

VOID Trace(TRACE trace, VOID *v) {
    for (BBL bbl = TRACE_BblHead(trace); BBL_Valid(bbl); bbl = BBL_Next(bbl)) {
        BBL_InsertCall(bbl, IPOINT_ANYWHERE, AFUNPTR(docount),
                       IARG_FAST_ANALYSIS_CALL, IARG_END);
    }
}

int main(int argc, CHAR *argv[]) {
    PIN_InitSymbols();
    PIN_Init(argc, argv);
    TRACE_AddInstrumentFunction(Trace, 0);
    PIN_StartProgram();
    return 0;
}
```

Simple BBCount DynamoRIO Tool


```
static int bbcount;

static void docount() { bbcount++; }


static dr_emit_flags_t
event_basic_block(void *drcontext, void *tag, instrlist_t *bb,
                  bool for_trace, bool translating) {
    dr_insert_clean_call(drcontext, bb, instrlist_first(bb), docount, false, 0);
    return DR_EMIT_DEFAULT;
}

DR_EXPORT void dr_init(client_id_t id) {
    dr_register_bb_event(event_basic_block);
}
```

BBCount Performance Comparison: Simple Tool



BBCount Performance Comparison: Simple Tool



Optimized BBCount DR Tool: Raw Version

```
static int global_count;

static dr_emit_flags_t
event_basic_block(void *drcontext, void *tag, instrlist_t *bb,
                  bool for_trace, bool translating) {
    instr_t *instr, *first = instrlist_first(bb);
    uint flags;
    /* Our inc can go anywhere, so find a spot where flags are dead.
     * Technically this can be unsafe if app reads flags on fault =>
     * stop at instr that can fault, or supply runtime op */
    for (instr = first; instr != NULL; instr = instr_get_next(instr)) {
        flags = instr_get_arith_flags(instr);
        /* OP_inc doesn't write CF but not worth distinguishing */
        if (TESTALL(EFLAGS_WRITE_6, flags) && !TESTANY(EFLAGS_READ_6, flags))
            break;
    }
    if (instr == NULL)
        dr_save_arith_flags(drcontext, bb, first, SPILL_SLOT_1);
    instrlist_tool_preinsert(bb, (instr == NULL) ? first : instr,
        INSTR_CREATE_inc(drcontext, OPND_CREATE_ABSMEM((byte *)&global_count, OPSZ_4)));
    if (instr == NULL)
        dr_restore_arith_flags(drcontext, bb, first, SPILL_SLOT_1);
    return DR_EMIT_DEFAULT;
}

DR_EXPORT void dr_init(client_id_t id) {
    dr_register_bb_event(event_basic_block);
}
```


Optimized BBCount DR Tool: Leverage API

```
static int global_count;


static dr_emit_flags_t
event_basic_block(void *drcontext, void *tag, instrlist_t *bb,
                  bool for_trace, bool translating) {
    instr_t *first = instrlist_first(bb);
    drx_insert_counter_update(drcontext, bb, first, SPILL_SLOT_MAX+1,
                             (void *)&global_count, 1, 0);
    return DR_EMIT_DEFAULT;
}

DR_EXPORT void dr_init(client_id_t id) {
    dr_register_bb_event(event_basic_block);
}
```

BBCount Performance Comparison: Opt Tool



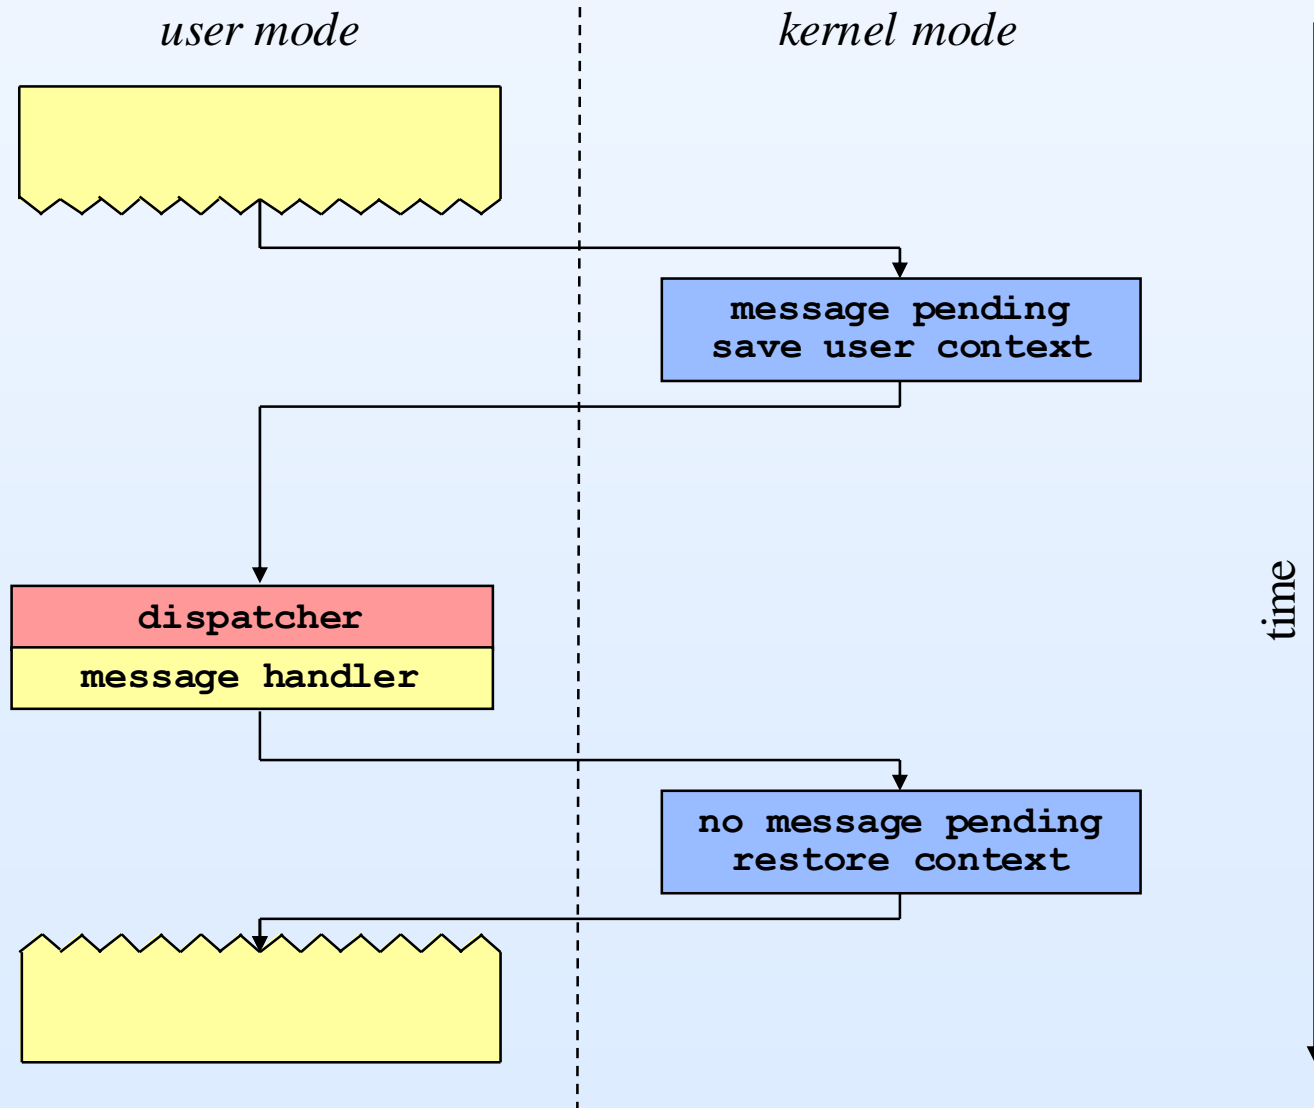
BBCount Performance Comparison: Opt Tool



Optional Slides:

Advanced API

Callback-Local Storage (CLS)




Callback-Local Storage (CLS)

- Windows callbacks interrupt execution to process an event and later resume the suspended context
- TLS data from the suspended context will be overwritten during callback execution
- CLS data is saved at the interruption point and restored at the resumption point
- Whenever keeping persistent data specific to one context rather than overall execution, use CLS instead of TLS
 - Usually only needed when storing data specific to a system call in pre-syscall event and reading it back in post-syscall event
- Can be used for Linux signals as well
- Provided by the drmgr Extension

DynamoRIO API: Translation

- Translation refers to the mapping of a code cache machine state (program counter, registers, and memory) to its corresponding application state
 - The program counter always needs to be translated
 - Registers and memory may also need to be translated depending on the transformations applied when copying into the code cache

Translation Case 1: Fault



- Exception and signal handlers are passed machine context of the faulting instruction.
- For transparency, that context must be translated from the code cache to the original code location
- Translated location should be where the application would have had the fault or where execution should be resumed

Translation Case 2: Relocation

- If one application thread suspends another, or DynamoRIO suspends all threads for a synchronous cache flush:
 - Need suspended target thread in a safe spot
 - Not always practical to wait for it to arrive at a safe spot (if in a system call, e.g.)
- DynamoRIO forcibly relocates the thread
 - Must translate its state to the proper application state at which to resume execution

Translation Approaches

- Two approaches to program counter translation:
 - Store mappings generated during fragment building
 - High memory overhead (> 20% for some applications, because it prevents internal storage optimizations) even with highly optimized difference-based encoding. Costly for something rarely used.
 - Re-create mapping on-demand from original application code
 - Cache consistency guarantees mean the corresponding application code is unchanged
 - Requires idempotent code transformations
- DynamoRIO supports both approaches
 - The engine mostly uses the on-demand approach, but stored mappings are occasionally needed

Instruction Translation Field

- Each instruction contains a translation field
- Holds the application address that the instruction corresponds to
- Set via `instr_set_translation()`

Context Translation Via Re-Creation

```
A1: mov  %ebx, %ecx  
A2: add  %eax, (%ecx)  
A3: cmp  $4, (%eax)  
A4: jle  710349fb
```

```
C1: mov  %ebx, %ecx  
C2: add  %eax, (%ecx)  
C3: cmp  $4, (%ax)  
C4: jle  <stub0>  
C5: jmp  <stub1>
```



```
D1: (A1) mov  %ebx, %ecx  
D2: (A2) add  %eax, (%ecx)  
D3: (A3) cmp  $4, (%eax)  
D4: (A4) jle  <stub0>  
D5: (A4) jmp  <stub1>
```

Application vs. Meta Instructions

- By default, instructions are treated as application instructions
 - Must have translations: `instr_set_translation()`, `INSTR_XL8()`
 - Control-flow-changing app instructions are modified to retain DynamoRIO control and result in cache populating
- *Meta* instructions are added instrumentation code
 - Not treated as part of the application (e.g., calls run natively)
 - Usually cannot fault, so translations not needed
 - Created via `instr_set_meta()` Or `instrlist_meta_append()`
- Meta instructions can reference application memory, or deliberately fault
 - A meta instruction that might fault must contain a translation
 - The client should handle any such fault

Client Translation Support

- Instruction lists passed to clients are annotated with translation information
 - Read via `instr_get_translation()`
 - Clients are free to delete instructions, change instructions and their translations, and add new tool and app instructions (see `dr_register_bb_event()` for restrictions)
 - An idempotent client that restricts itself to deleting app instructions and adding non-faulting meta instructions can ignore translation concerns
 - DynamoRIO takes care of instructions added by API routines (`insert_clean_call()`, etc.)
- Clients can choose between storing or regenerating translations on a fragment by fragment basis.

Client Regenerated Translations

- Client returns `DR_EMIT_DEFAULT` from its bb or trace event callback
- Client bb & trace event callbacks are re-called when translations are needed with `translating==true`
- Client must exactly duplicate transformations performed when the block was generated
- Client must set translation field for all added app instructions and all meta instructions that might fault
 - This is true even if `translating==false` since DynamoRIO may decide it needs to store translations anyway

Client Stored Translations

- Client returns `DR_EMIT_STORE_TRANSLATIONS` from its bb or trace event callback
- Client must set translation field for all added app instructions and all meta instructions that might fault
- Client bb or trace hook will not be re-called with `translating==true`

Register State Translation

- Translation may be needed at a point where some registers are spilled to memory
 - During indirect branch or RIP-relative mangling, e.g.
- DynamoRIO walks fragment up to translation point, tracking register spills and restores
 - Special handling for stack pointer around indirect calls and returns
- DynamoRIO tracks client spills and restores *implicitly* added by API routines
 - Clean calls, etc.
 - Explicit spill/restore (e.g., `dr_save_reg()`) client's responsibility

Client Register State Translation

- If a client adds its own register spilling/restoring code or changes register mappings it must register for the restore state event to correct the context
- The same event can also be used to fix up the application's view of memory
- DynamoRIO does not internally store this kind of translation information ahead of time when the fragment is built
 - The client must maintain its own data structures

Application Predication

- Conditionally executed instructions do exist for x86
 - `cmovcc`, `fcmovcc`, `bsr`, `bsf`
 - Previously modeled by listing the destination as a source
- Several shades of predication
 - Complete conditional execution
 - Sources always read, destination written conditionally
 - Sources always read, destination written conditionally, but condition codes are always written (`bsr` and `bsf`, e.g.)

Application Predication Handling

- We added a predicate field to each instruction
- We added qualifiers to `instr_writes_to_reg()`, `instr_reads_from_reg()`, and other common analysis funcs
 - Support querying “might write” versus “definitely writes”: liveness and other analyses differ in what they want to know
- We back-ported this to x86 and removed the dst-as-src
 - Compatibility break
- Other attributes (e.g., number of dests) remain unchanged
 - User must handle predication there

Optional Slides: ARM-Specific

IT Blocks


- Thumb allows predication only inside IT blocks, where an IT instruction specifies the predicate for 1-4 subsequent instrs
- Arithmetic instrs have different semantics inside IT vs outside
 - Outside they write condition codes while inside they do not
- Example:

```
cmp      r5, r2
ite      ls
movls    r3, #0
movhi    r3, #1
```


IT Block Decoding

- Complicates general decoding and encoding
 - Decode: if see IT instr, store the PC and only assume in IT block on subsequent decode calls if the PC matches
 - Encode: only support IT when encoding an entire list of instrs
- Complicates instrumentation

```
cmp    r5, r2
ite    ls
movls  r3, #0
movhi  r3, #1
```



```
cmp    r5, r2
ite    ls
movls  r3, #0
strhi  r3, <TLS-slot>
movs   r3, #1
```

IT Block Instrumentation

- Upon decoding, predicates are set for instrs inside IT block
- Tools that want to insert instrumentation in the IT block should call an API routine that removes the IT instrs
 - Inserted instrumentation can then ignore IT blocks
 - Mangling step will re-insert IT instrs for each predicated Thumb instruction
 - Non-default: some tools want to see original opcode mix
- For general instruction generation outside of app blocks, tool must insert IT instrs and set predicates on its own for successful encoding
- Open to suggestions for other ways of handling IT blocks

Optional Slides:

Dr. Memory

Inter-Instruction Storage

- Spill slots provided by DR are only guaranteed to be live during a single app instr
 - In practice, live until next selfmod instr
- Allocate own TLS for spill slots
 - `dr_raw_tls_calloc()`
 - *drreg* does this for you
- Steal registers across whole bb
 - Restore before each app read
 - Update spill slot after each app write
 - Restore on fault
 - *drreg* does all of this for you too

Address Space Iteration

- Repeated calls to `dr_query_memory_ex()`
- Check `dr_memory_is_in_client()` and `dr_memory_is_dr_internal()`
- Heap walk
 - API on Windows
- Initial structures on Windows
 - TEB, TLS, etc.
 - PEB, ProcessParameters, etc.

Delayed Fragment Deletion

- Due to non-precise flushing we can have a flushed bb made inaccessible but not actually freed for some time
- When keeping state per bb, if a duplicate bb is seen, replace the state and increment a counter `ignore_next_delete`
- On a deletion event, decrement and ignore unless below 0
- Can't tell apart from duplication due to thread-private copies: but this mechanism handles that if saved info is deterministic and identical for each copy

Callstack Walking

- Use case: error reporting
- Technique:
 - Start with xbp as frame pointer (fp)
 - Look for <fp,retaddr> pairs where retaddr = inside a module
- Interesting issues:
 - When scanning for frame pointer (in frameless func, or at bottom of stack), querying whether in a module dominates performance
 - msvcr80!malloc pushes ebx and then ebp, requiring special handling
 - When displaying, use retaddr-1 for symbol lookup
 - More sophisticated techniques needed in presence of FPO

Callstack Walking

- Record on every allocation so must be very fast
 - Cannot afford to do symbol lookup, so just check whether a potential return address is inside a module
- Problem: stale return addresses
- Solutions:
 - For full mode, use definedness info
 - For other modes, zero out the return address upon returning from a function

Using Nudges

- Daemon apps do not exit
- Request results mid-run
- Cross-platform
 - Signal on Linux
 - Remote thread on Windows